

Language Models


Boris Zubarev

 @bobazooba

Word Embeddings


w2v, glove, etc

- Just key—value storage at inference
- Don't change from relationships with other words in the current text




Language Model

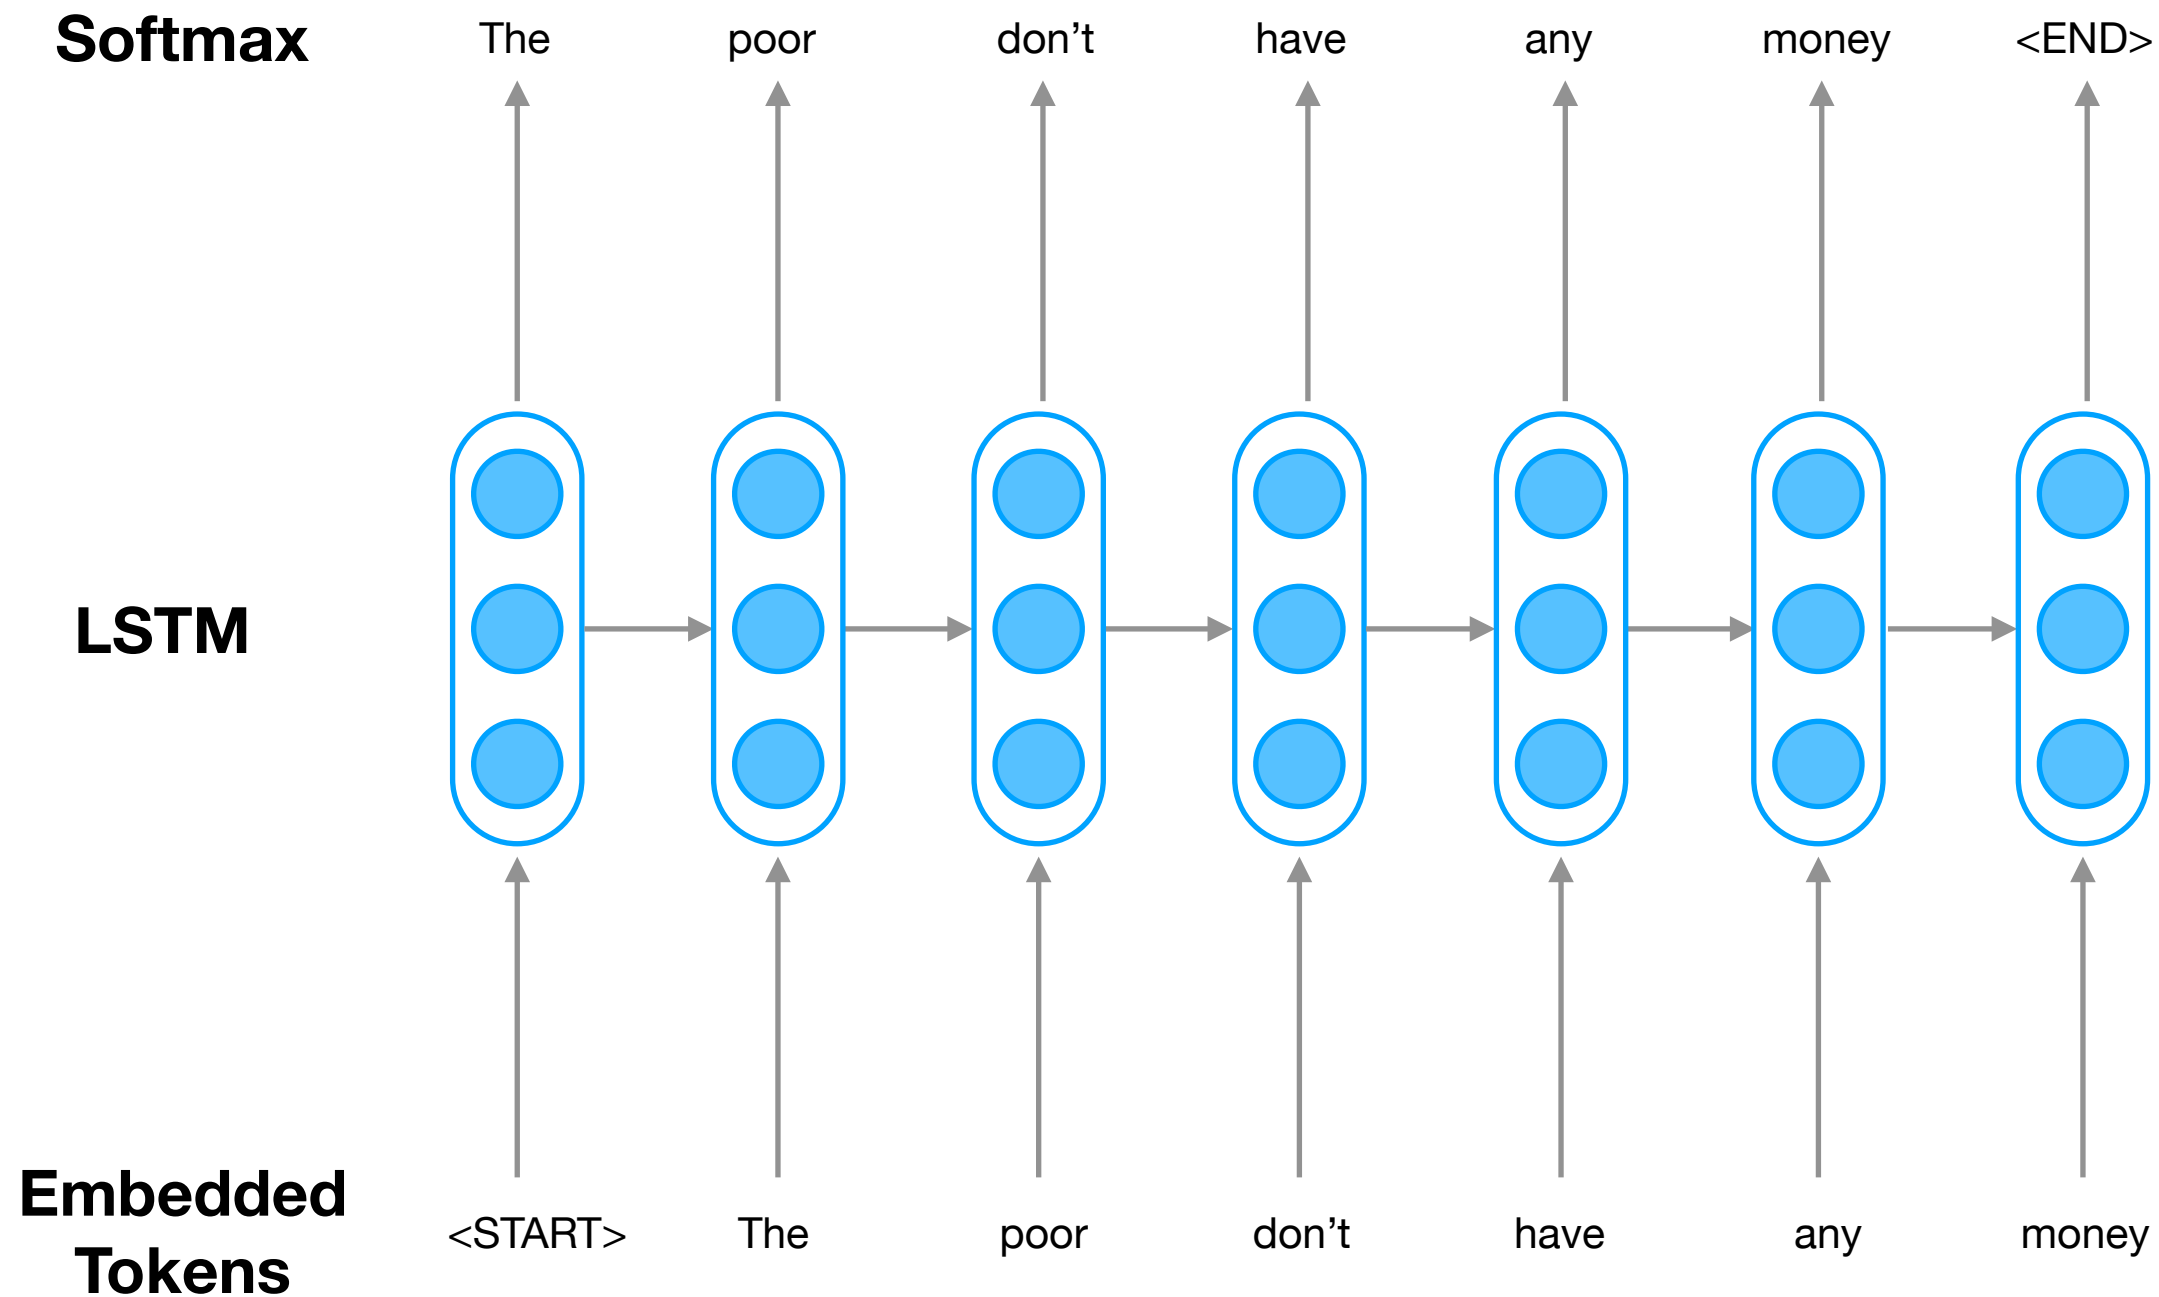
LSTM



Language Model



Language Model



Language Model

Inference

LSTM

**Embedded
Tokens**

<START>


Language Model

Inference

Softmax

LSTM


Embedded
Tokens



<START>


Language Model

Inference




Language Model

Inference




Language Model

Inference




Language Model

Inference




Language Model

Inference




Language Model

Inference



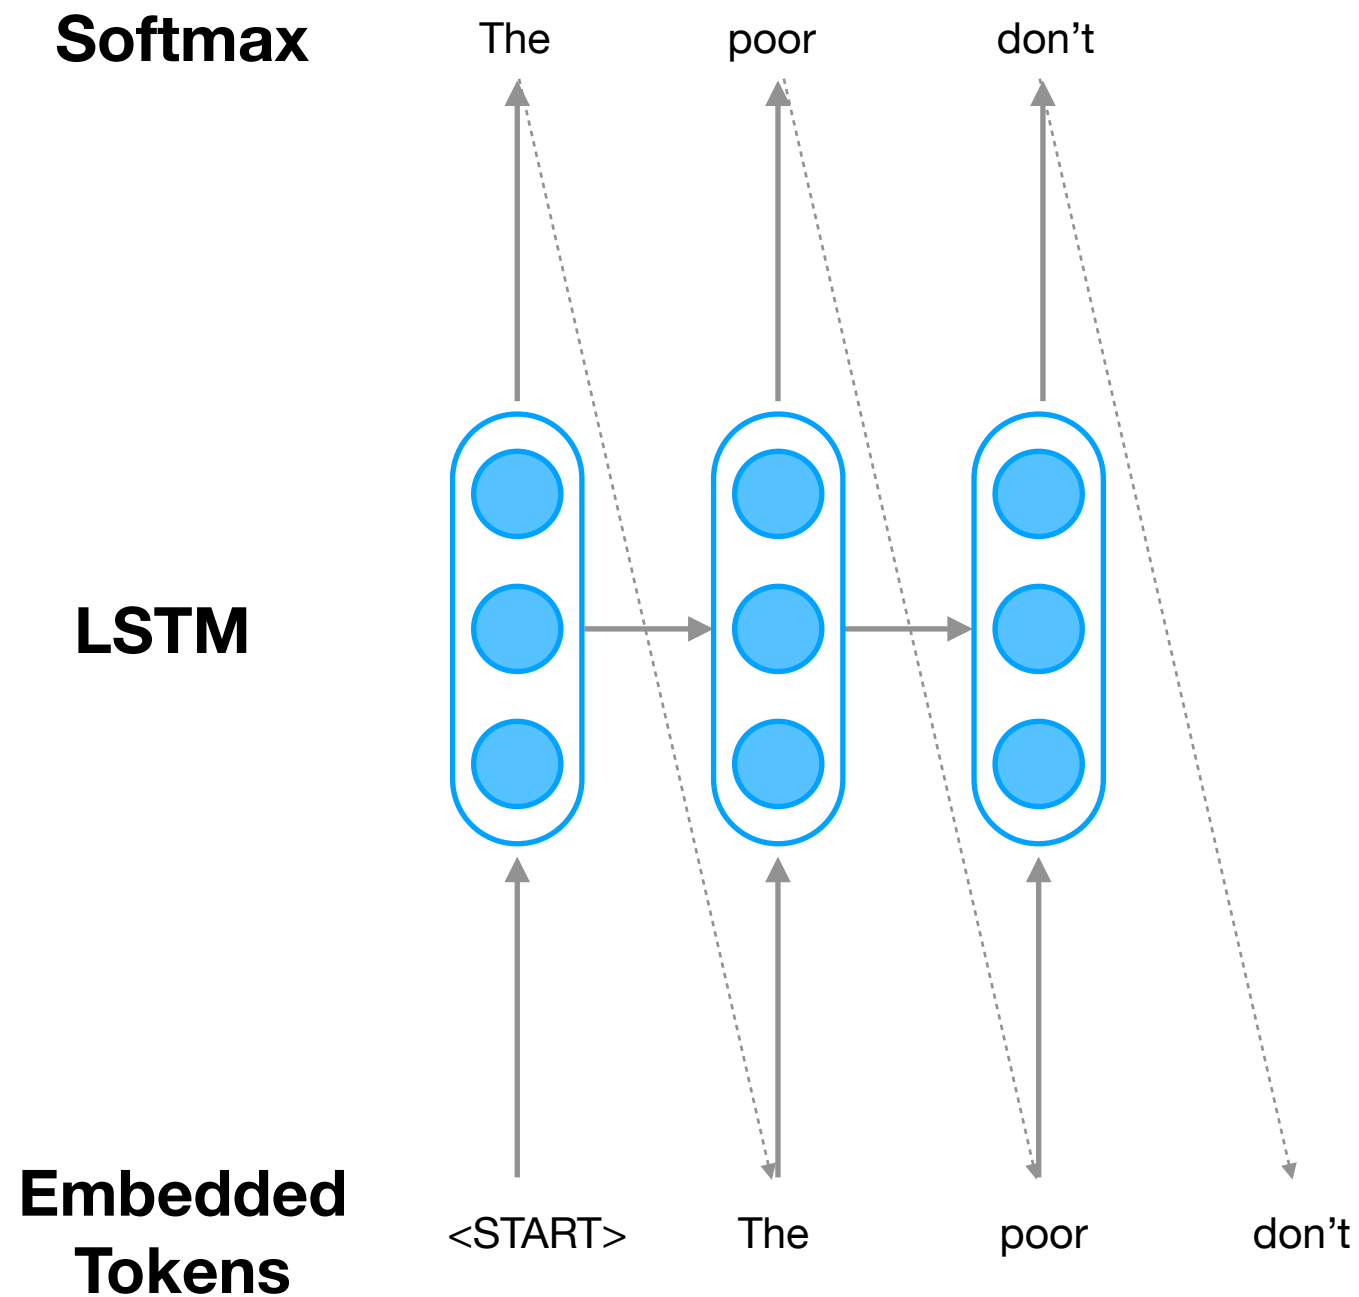
Language Model

Inference



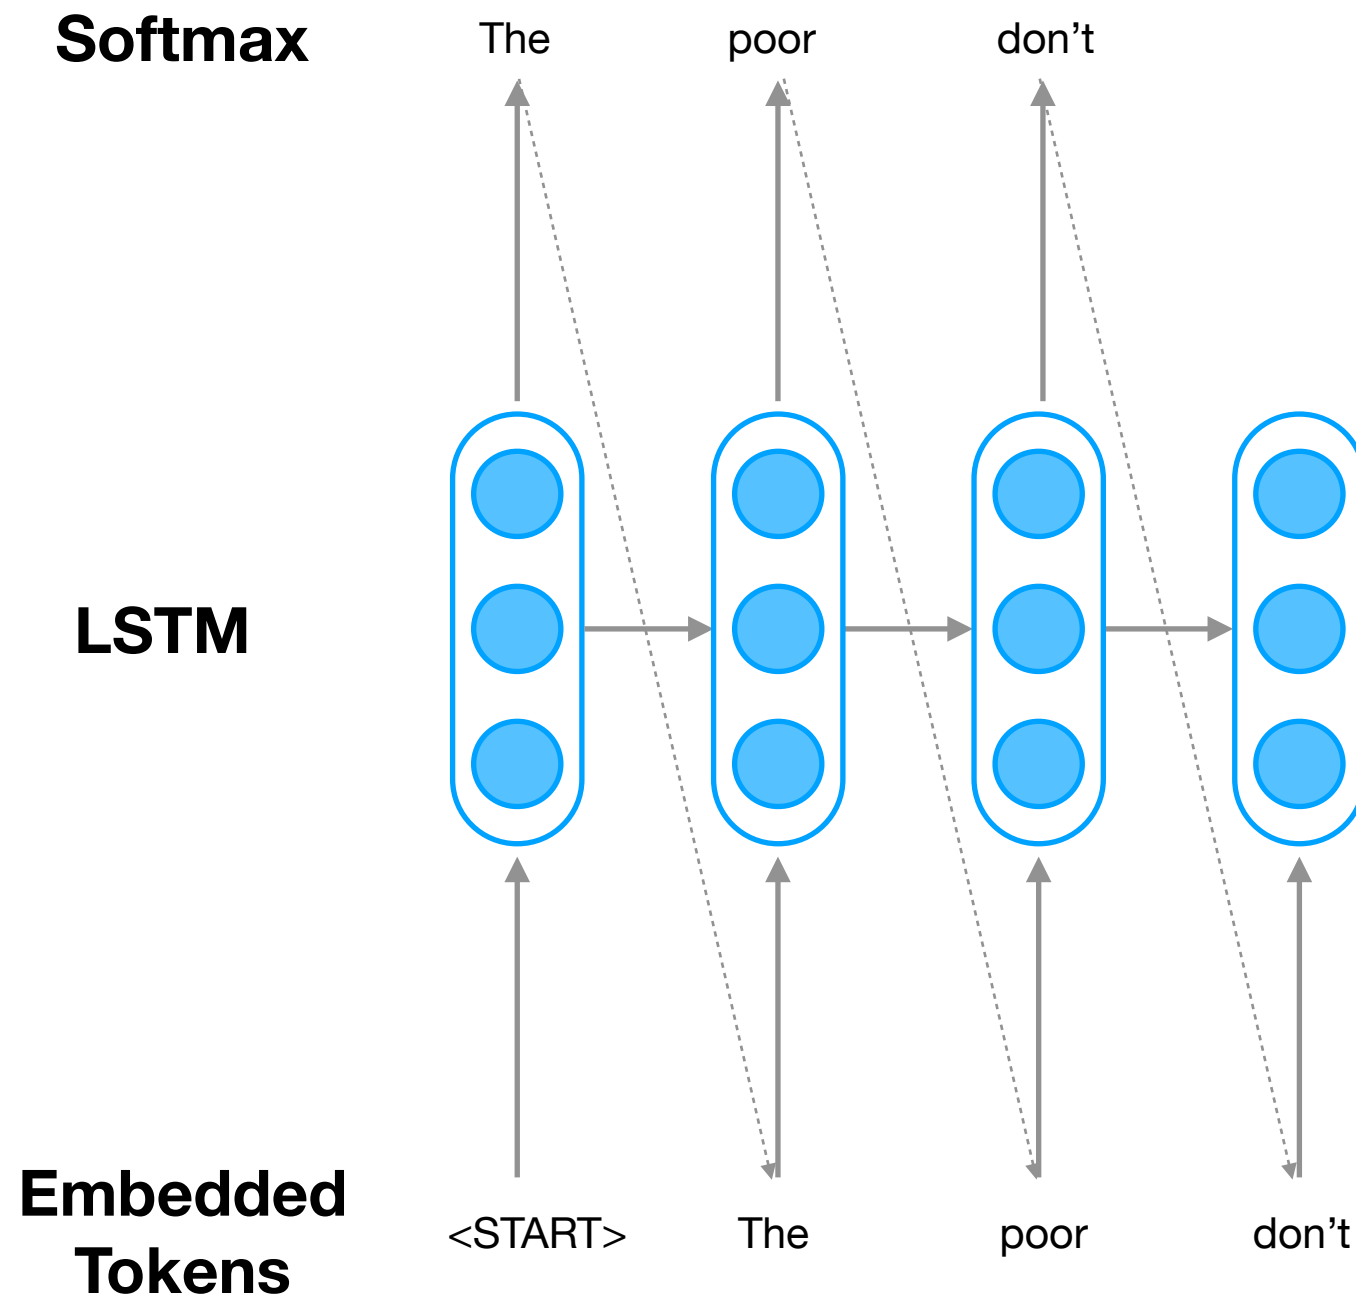
Language Model

Inference



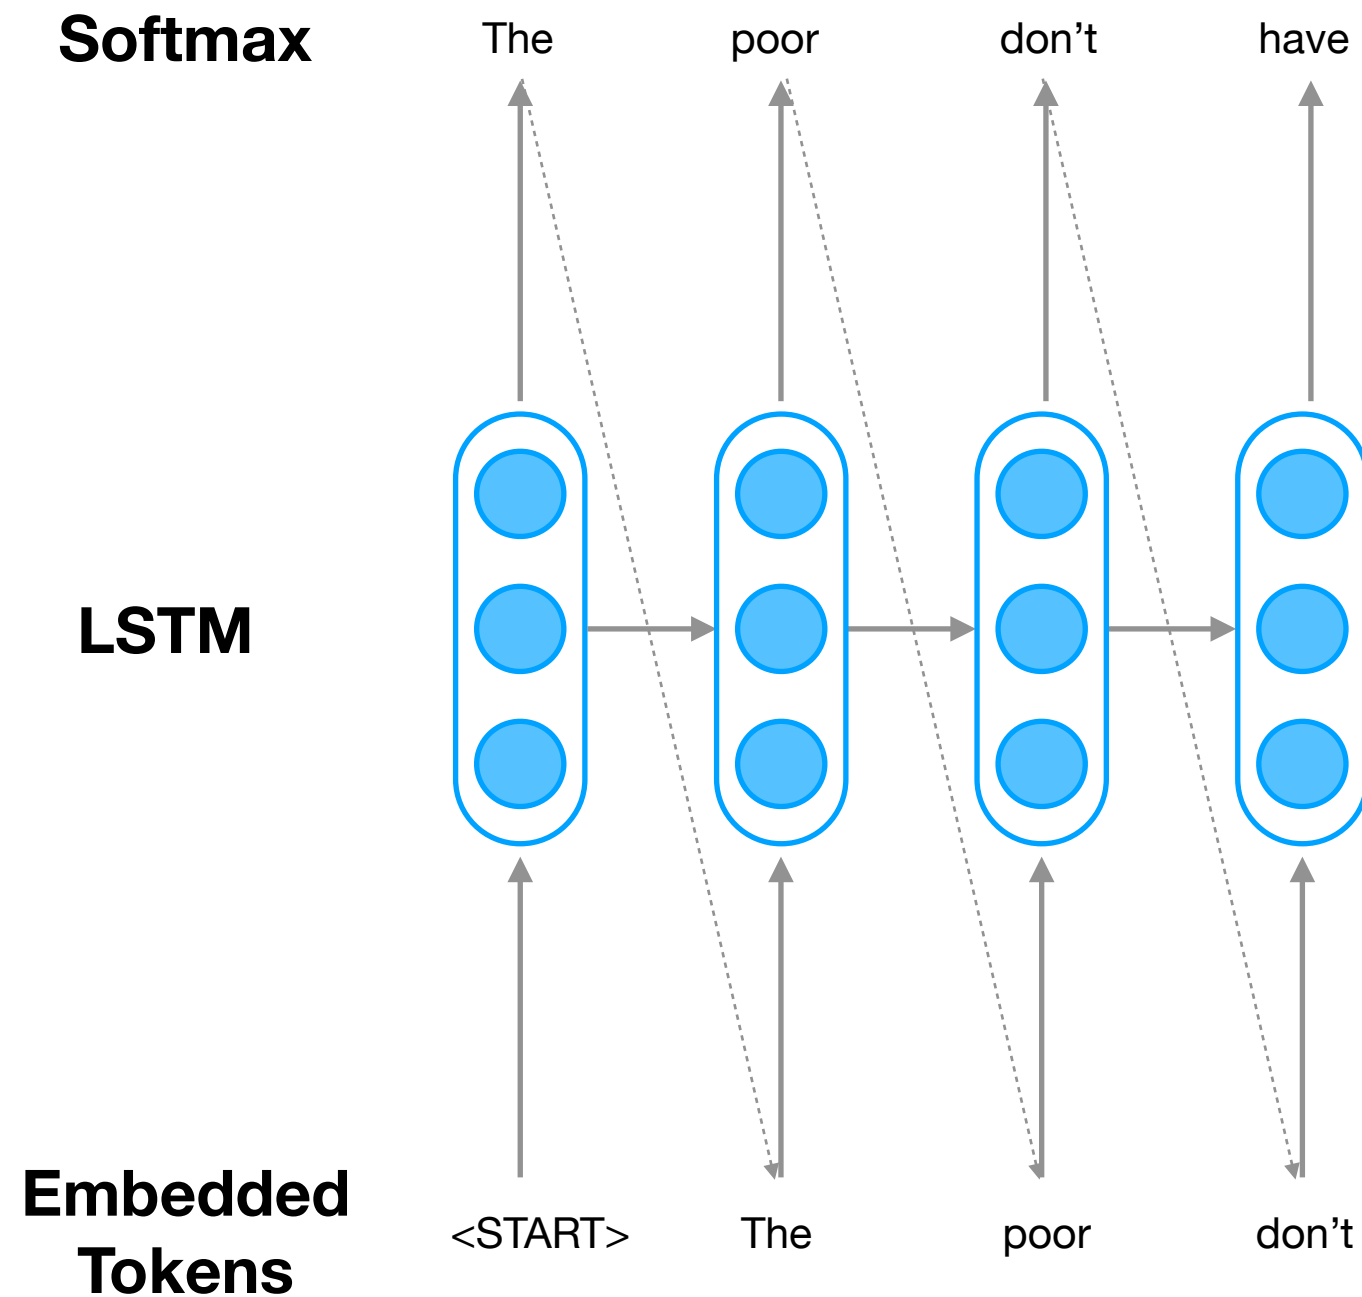
Language Model

Inference



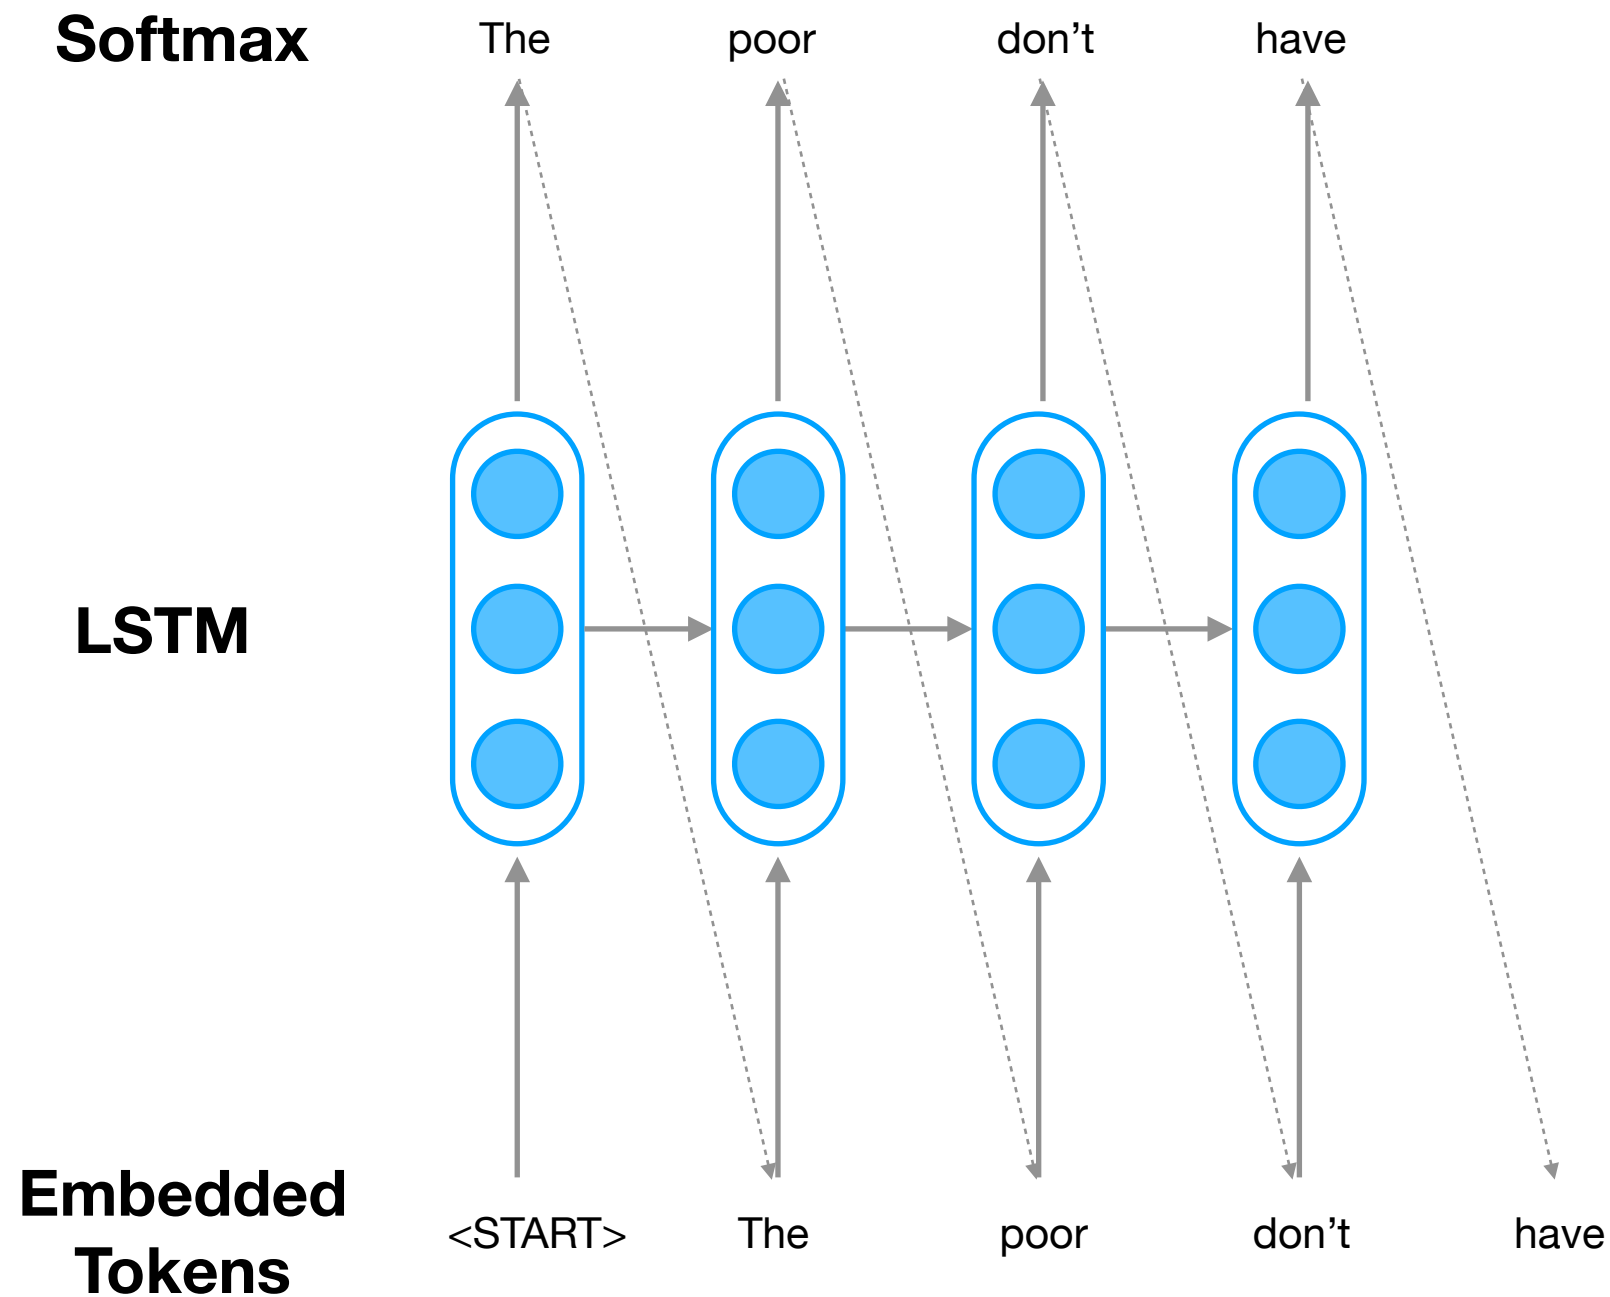
Language Model

Inference



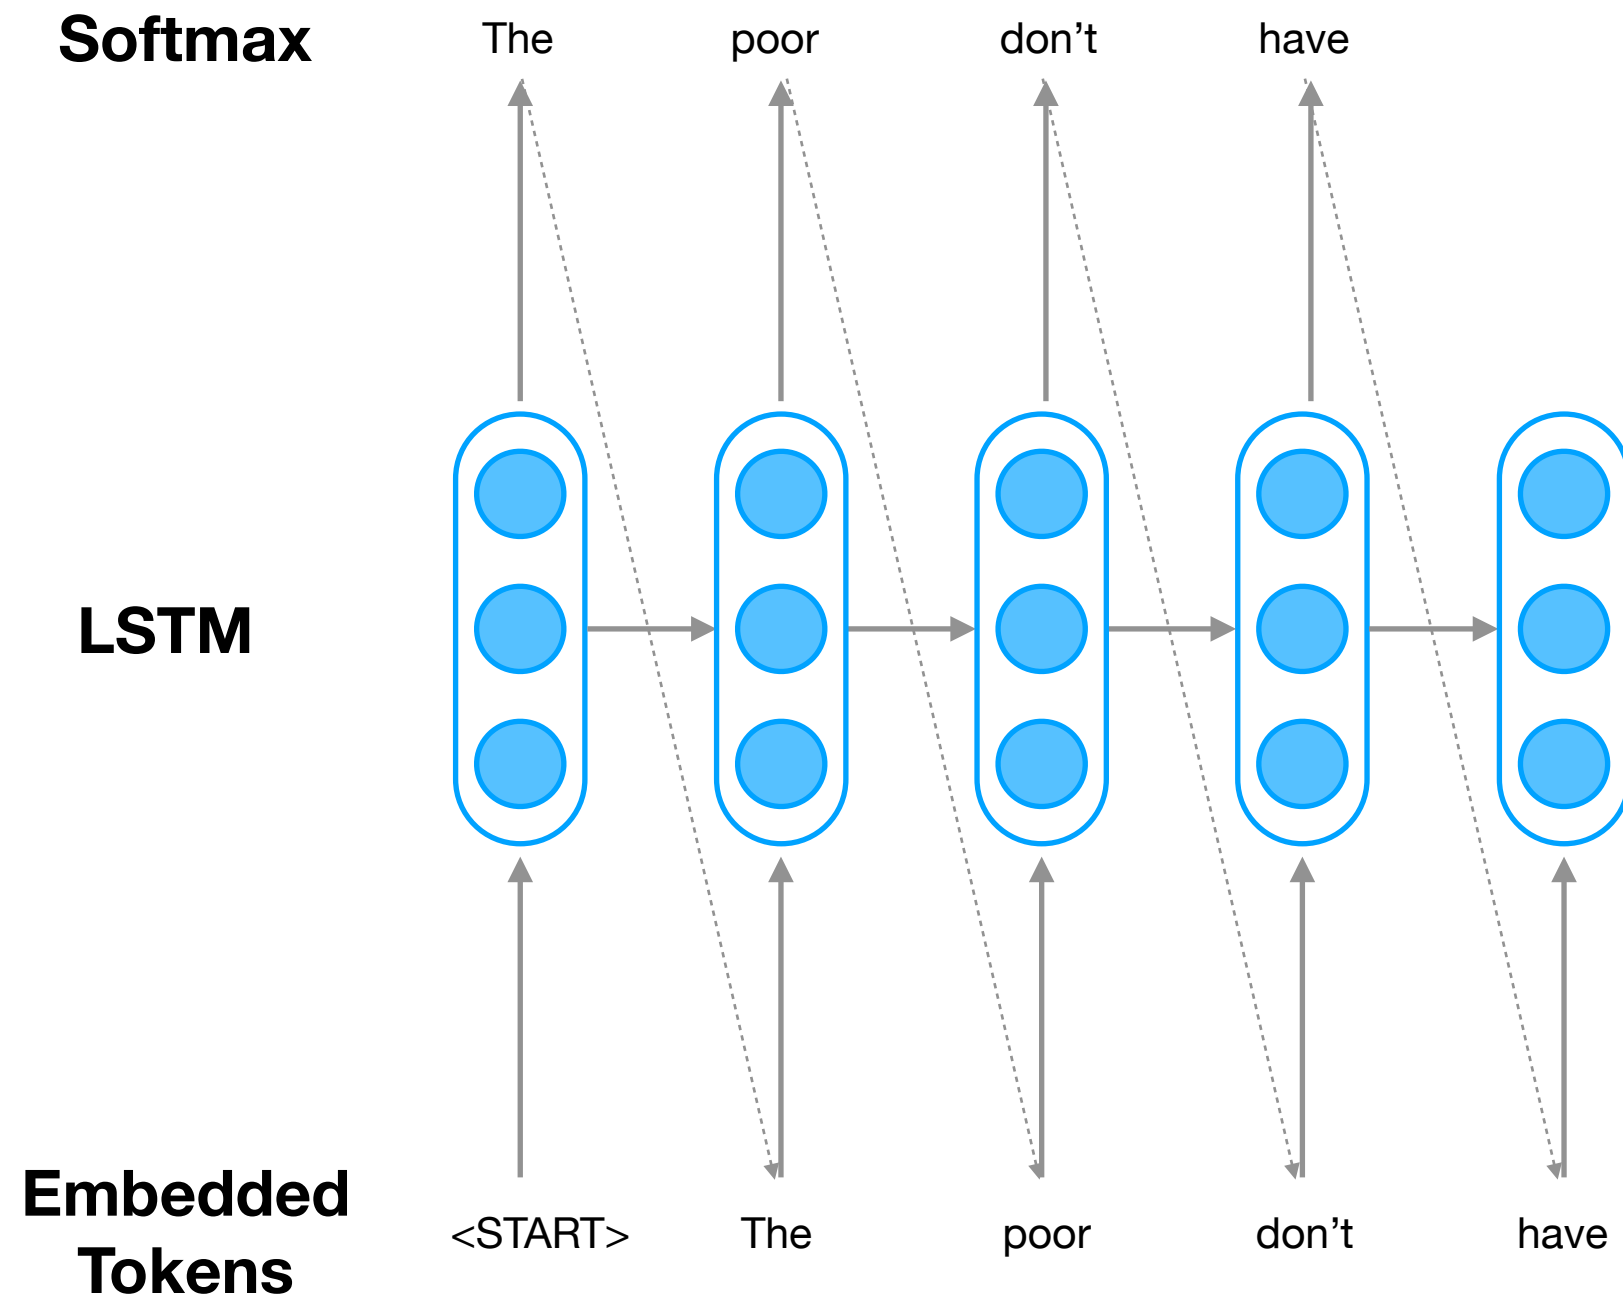
Language Model

Inference



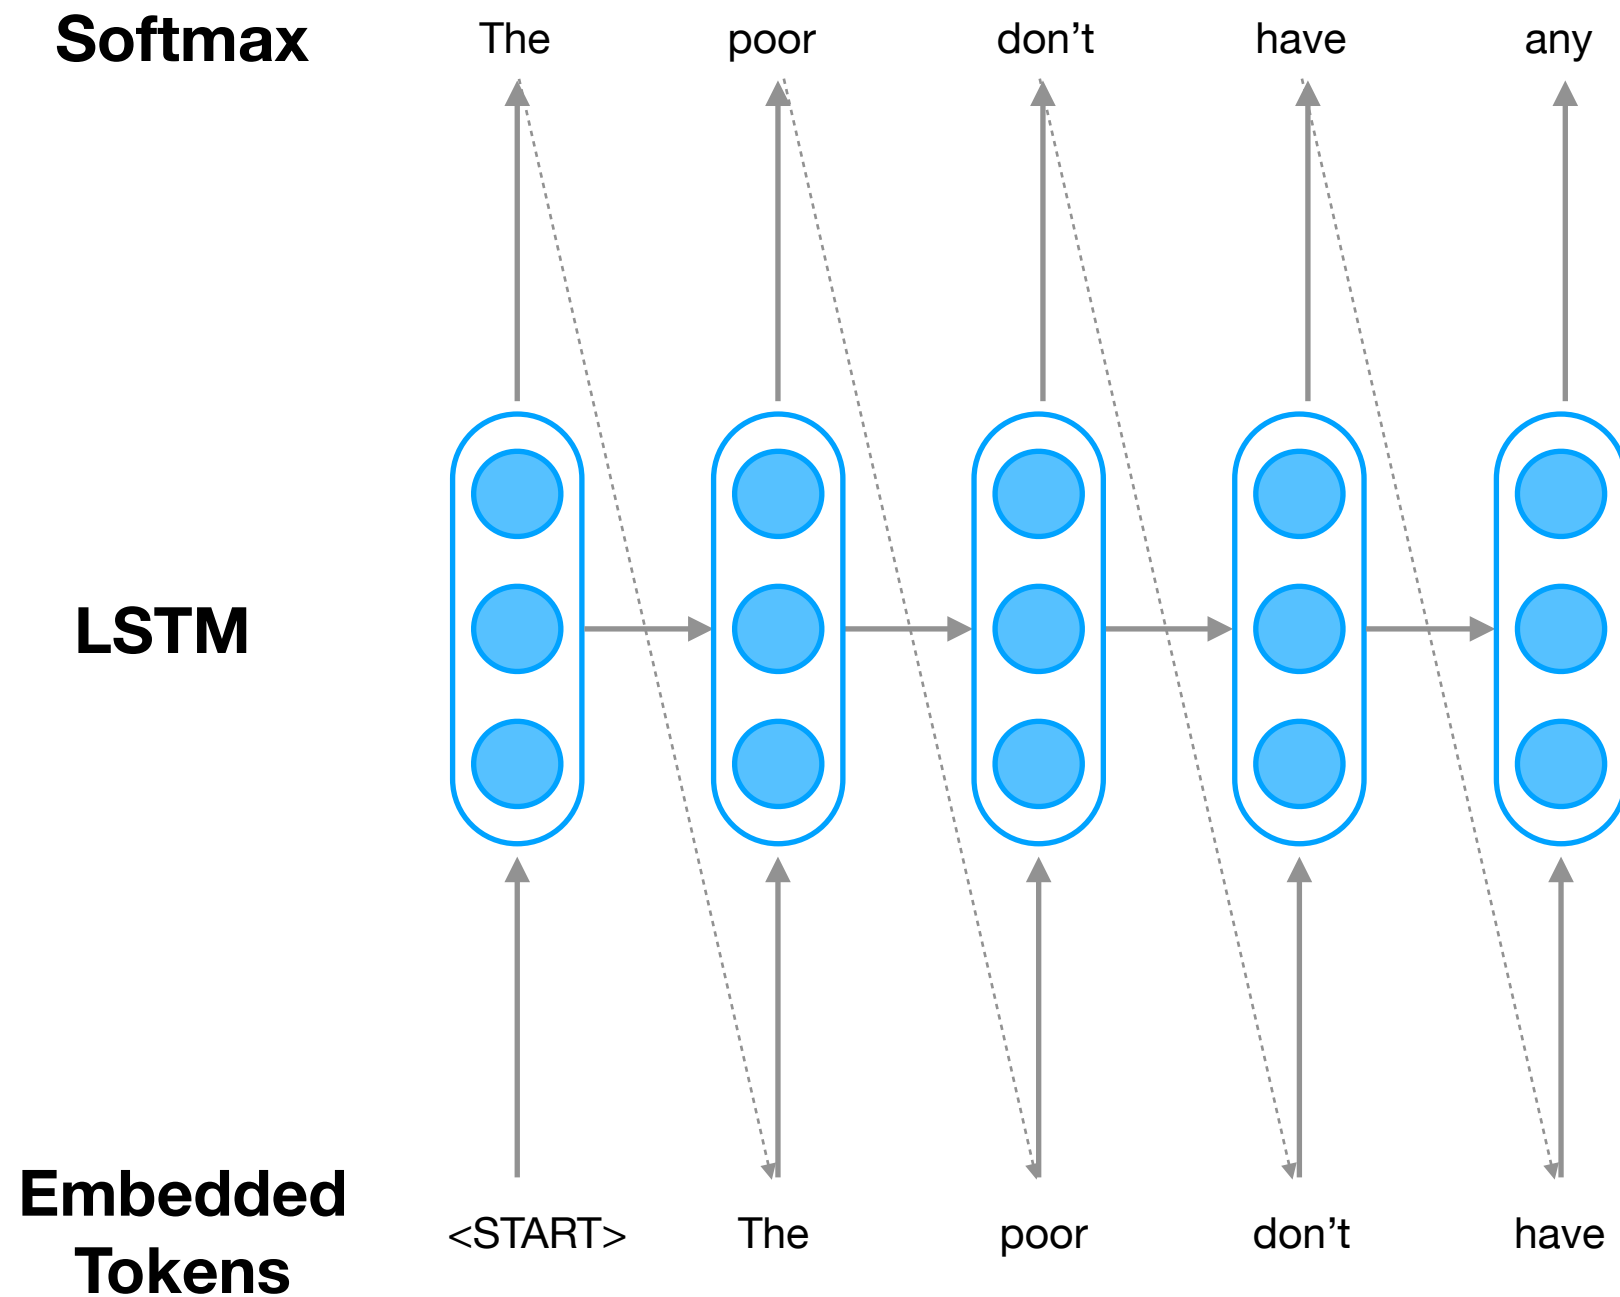
Language Model

Inference



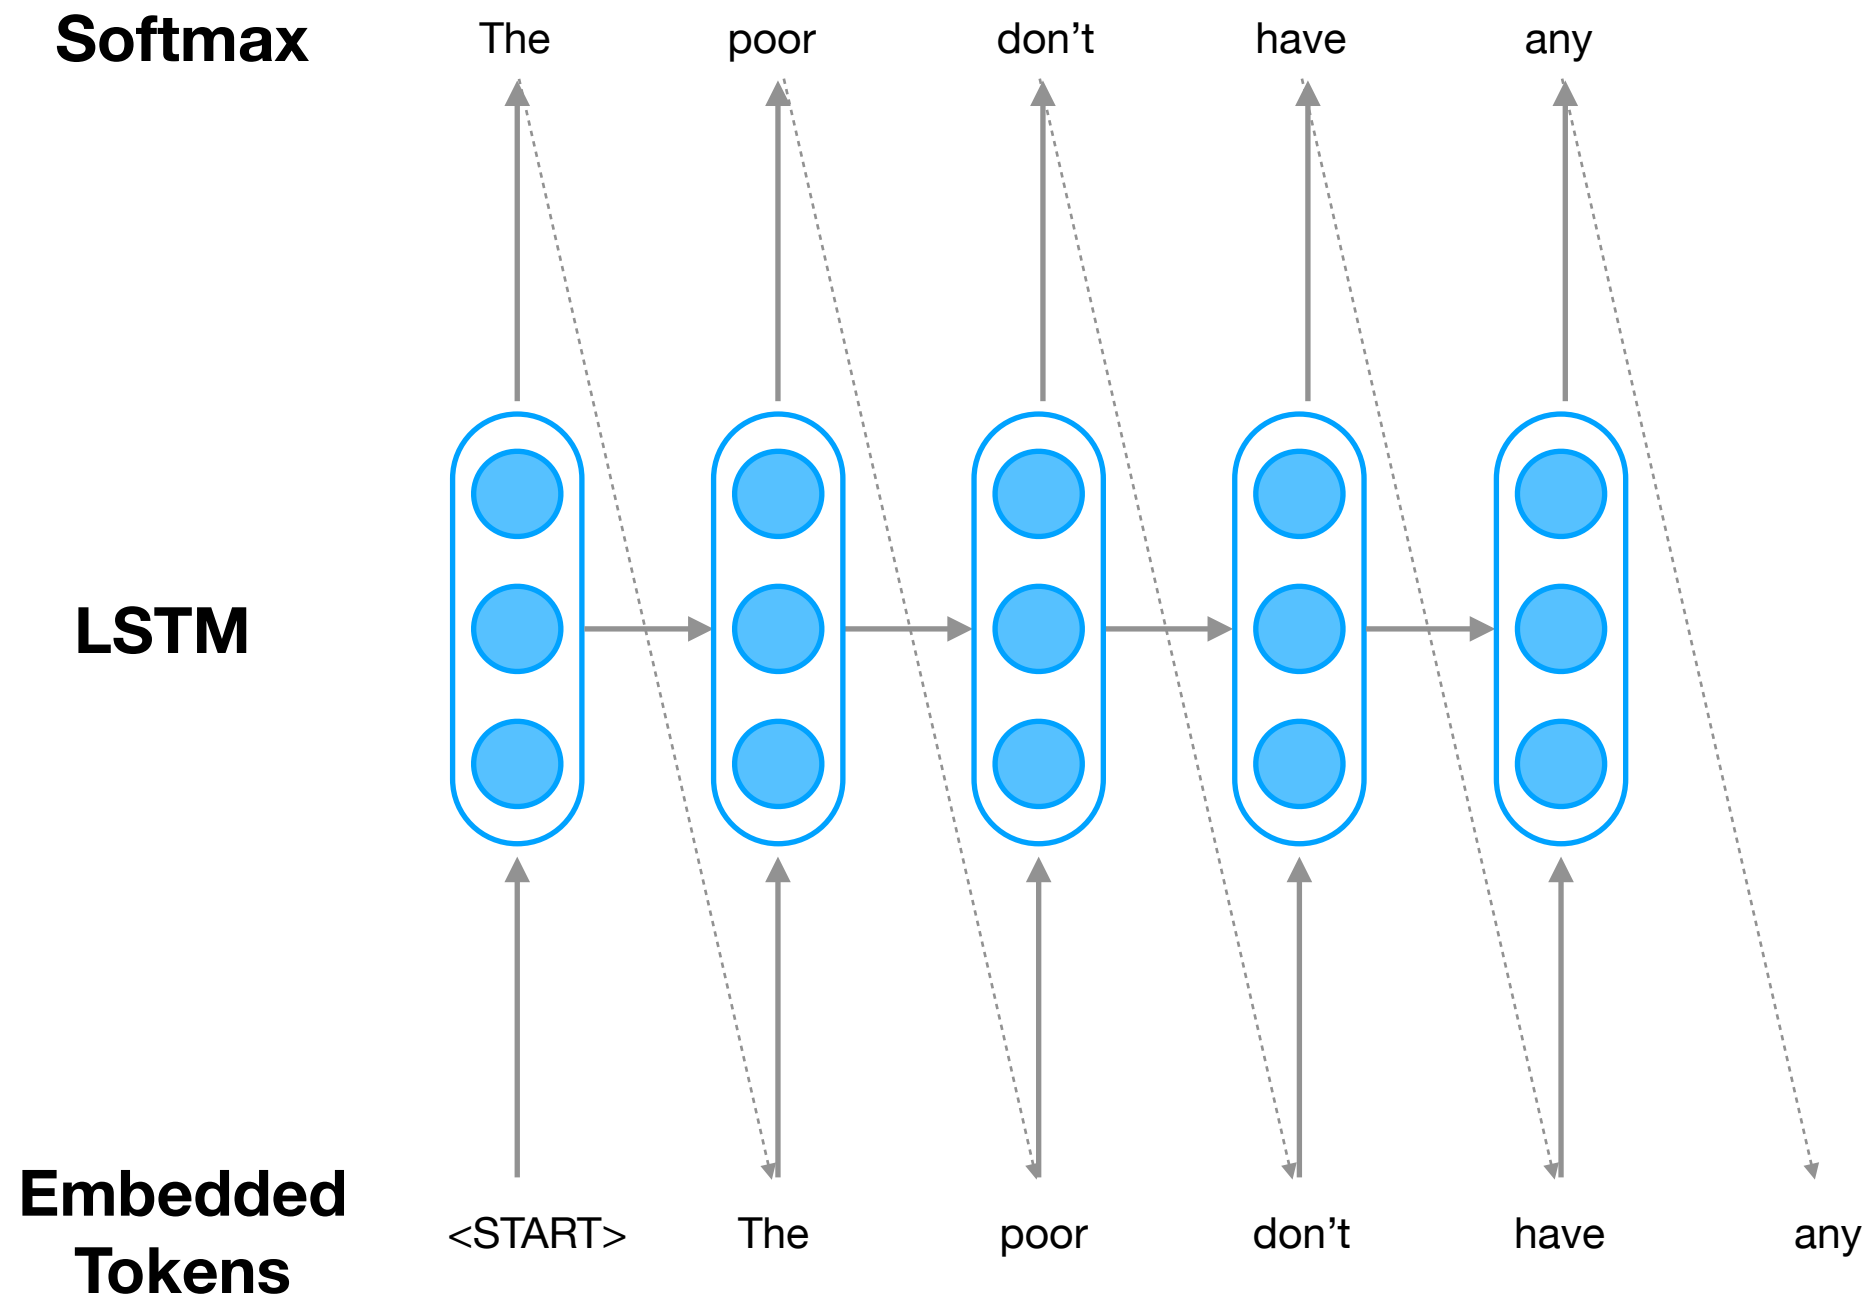
Language Model

Inference



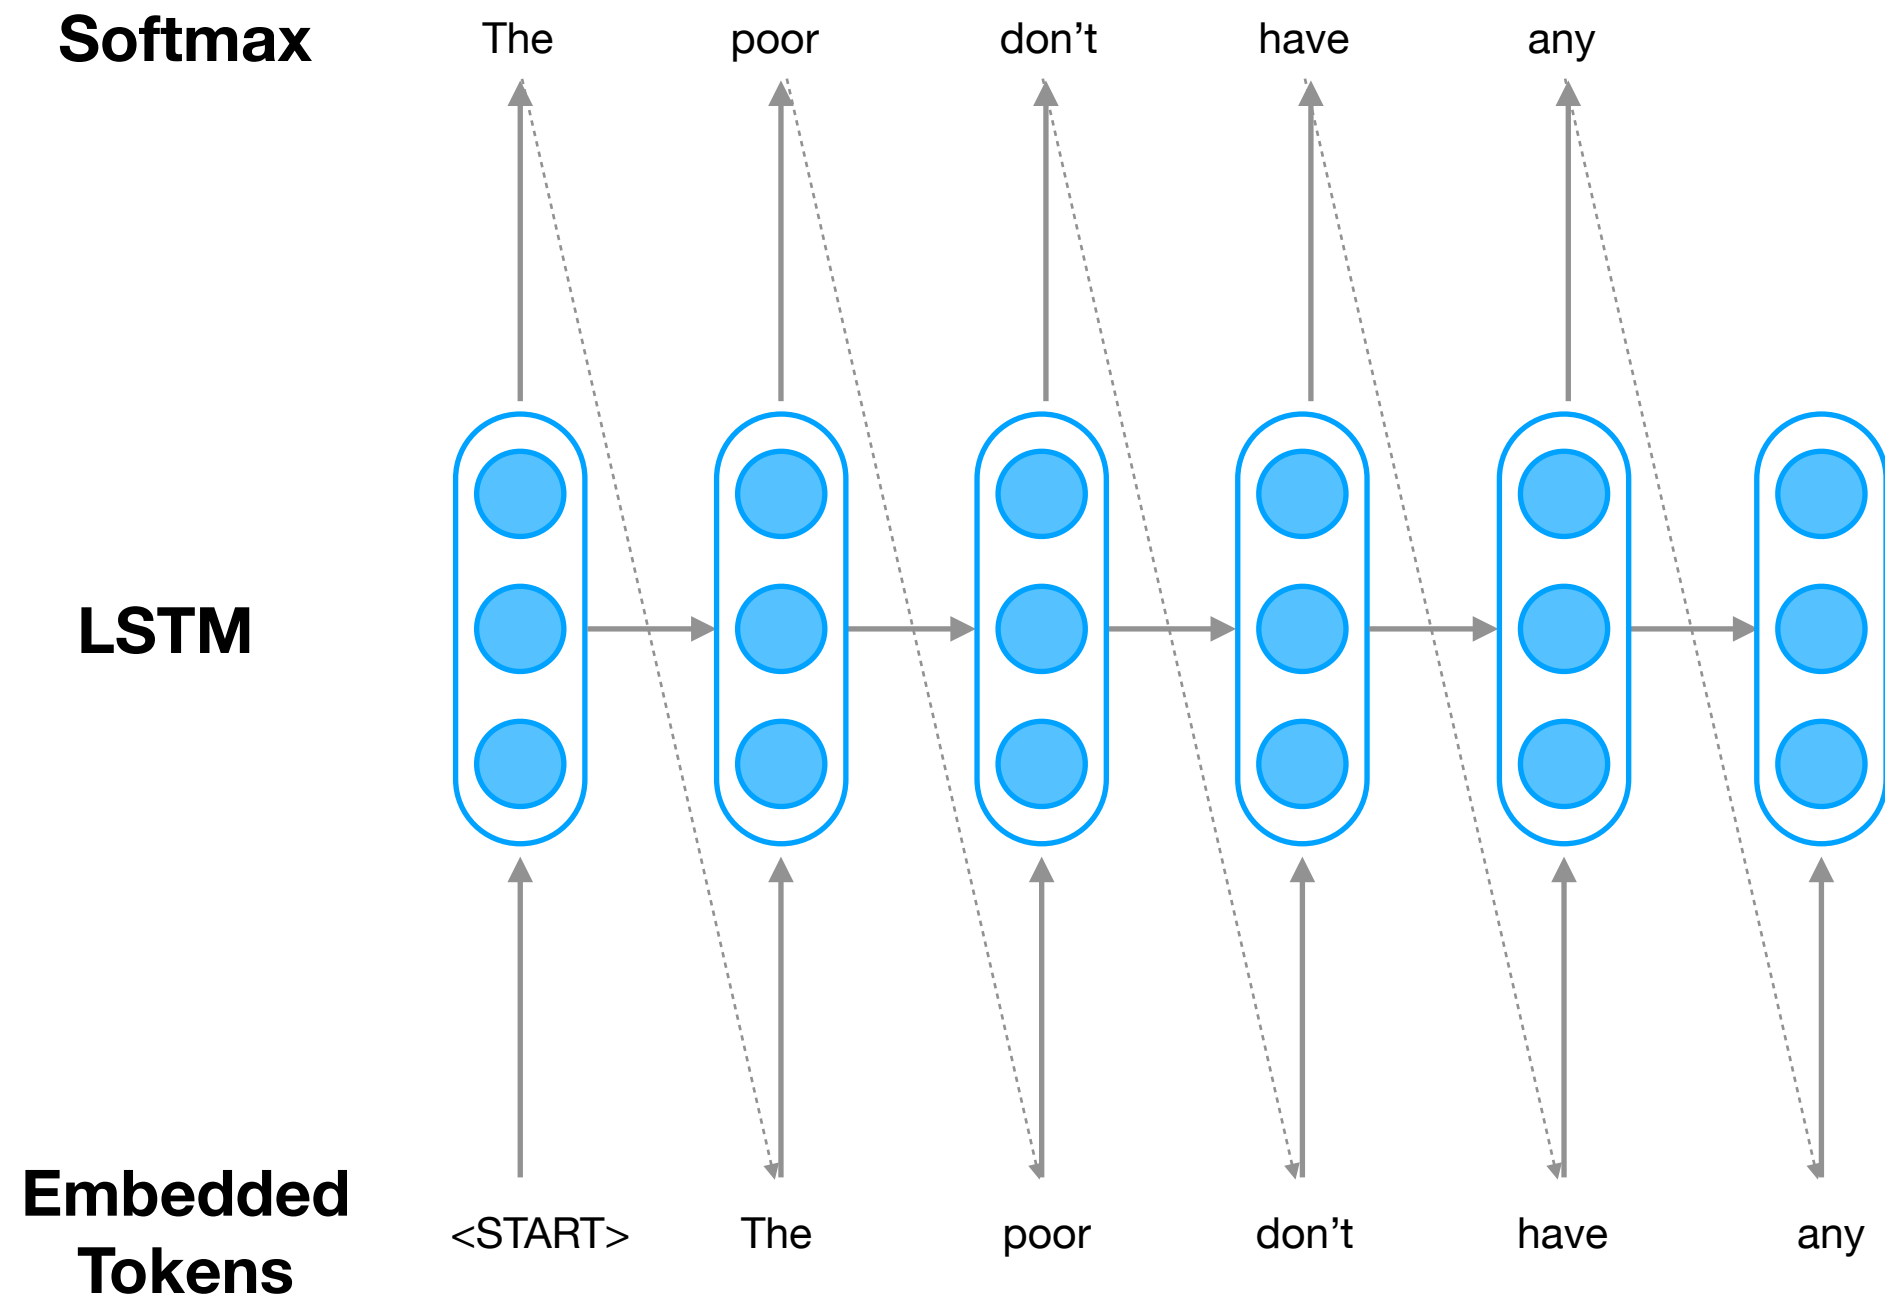
Language Model

Inference



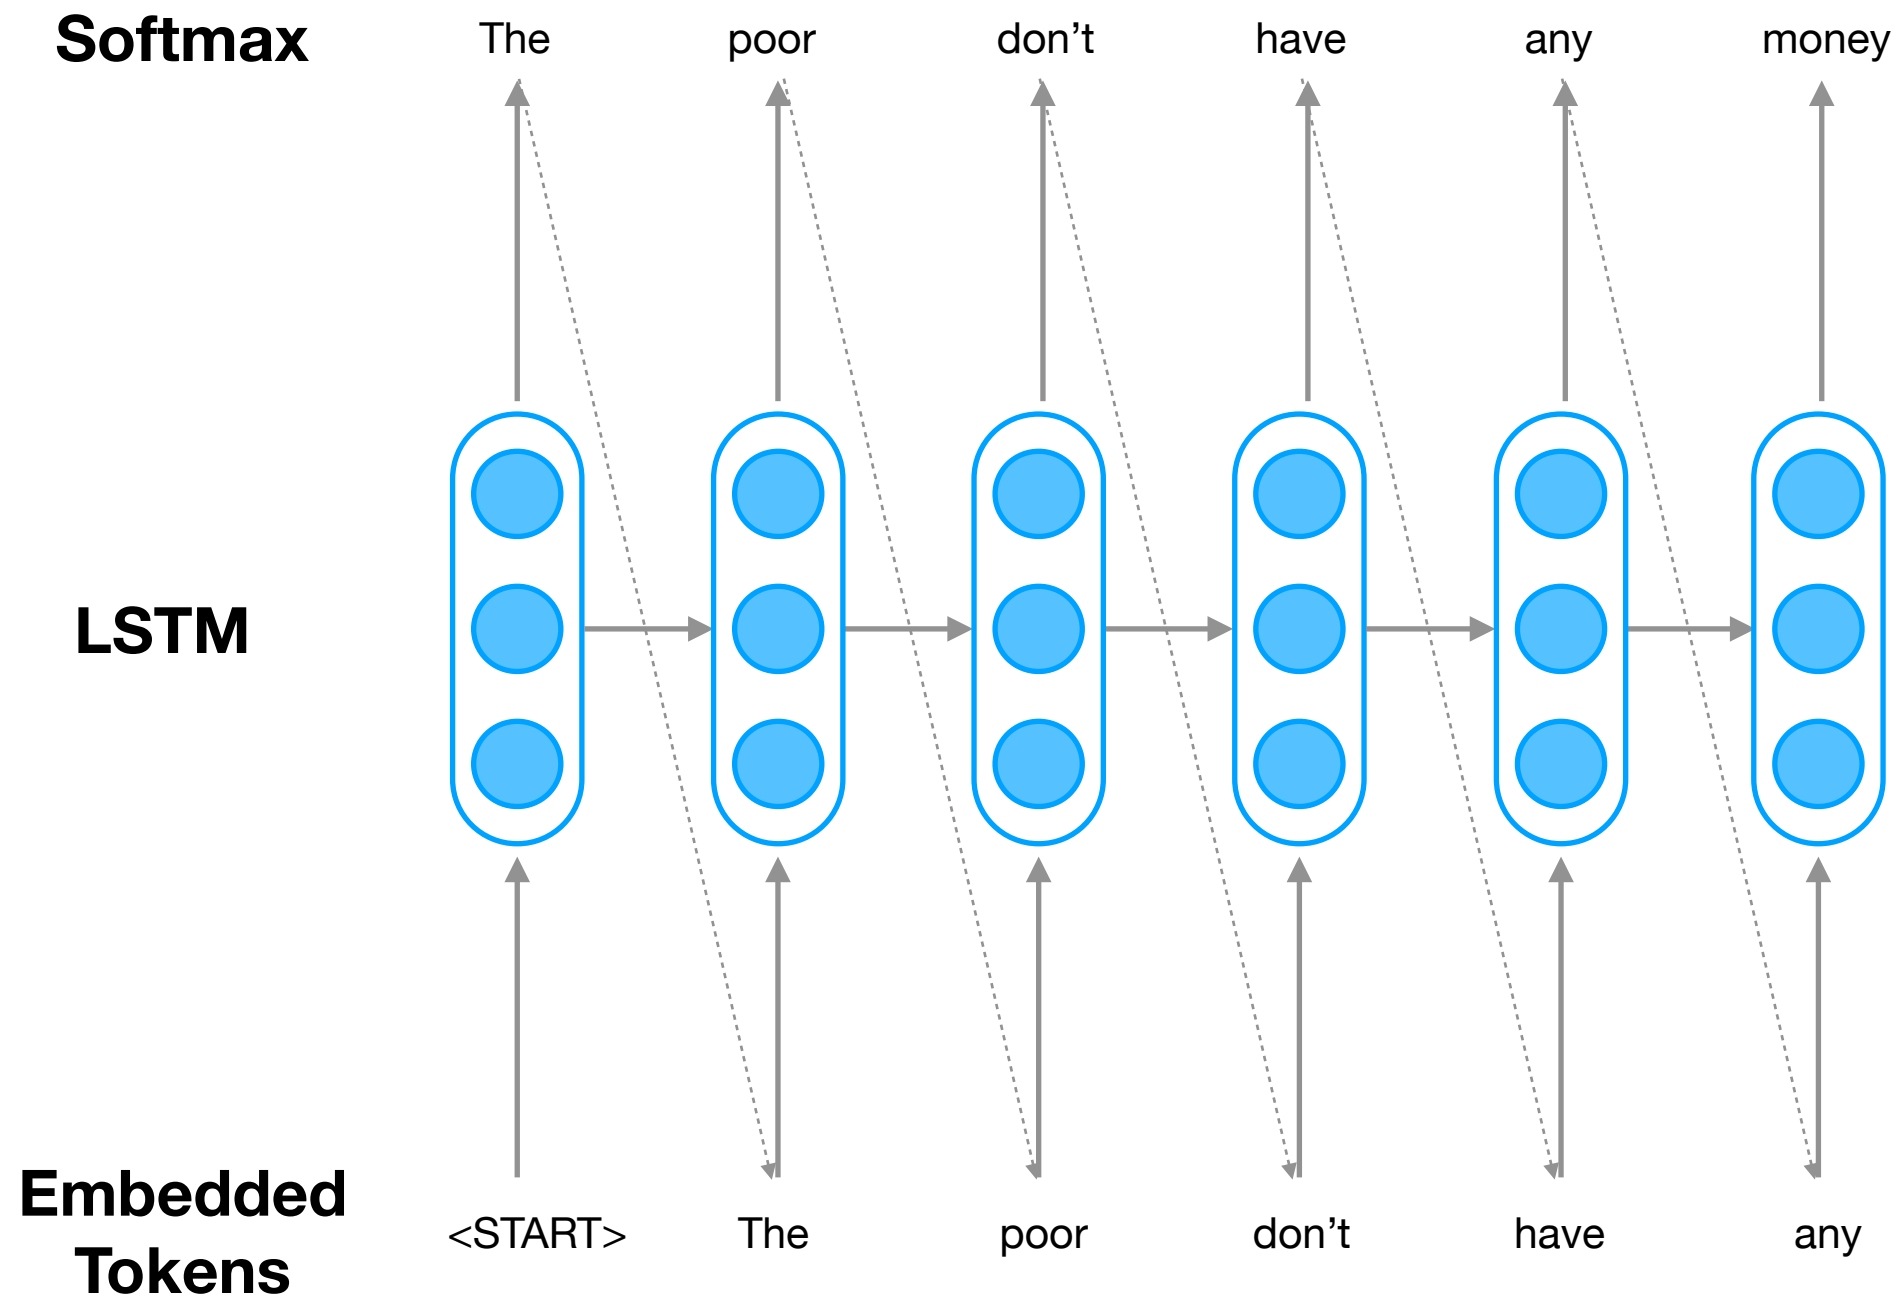
Language Model

Inference




Language Model

Inference




Language Model

Inference




Language Model

Inference



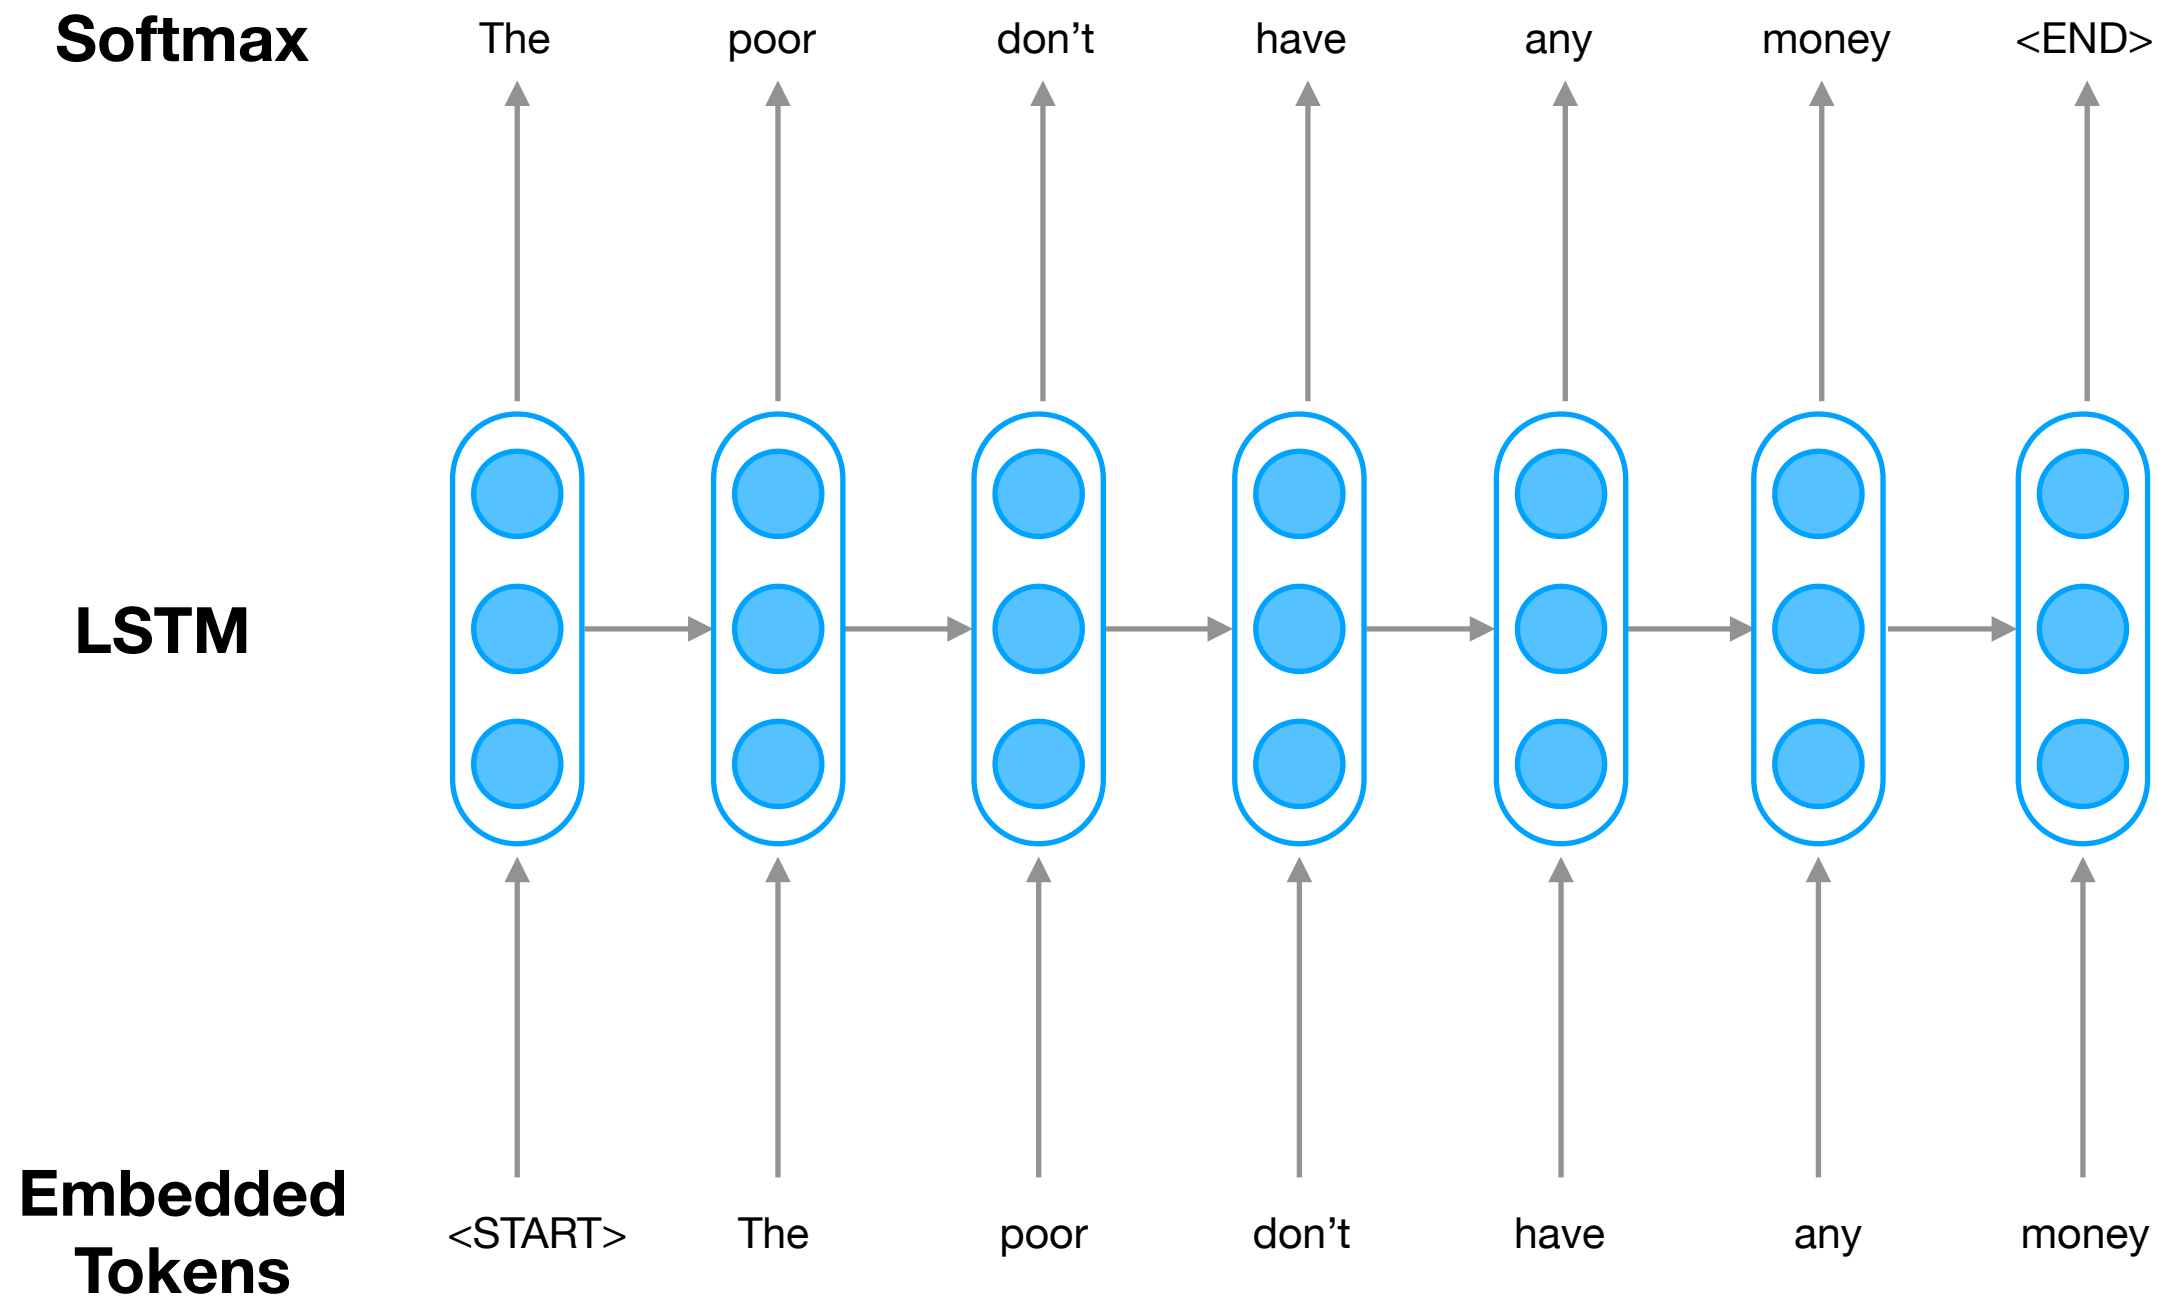
Language Model

Inference




Language Model

Training

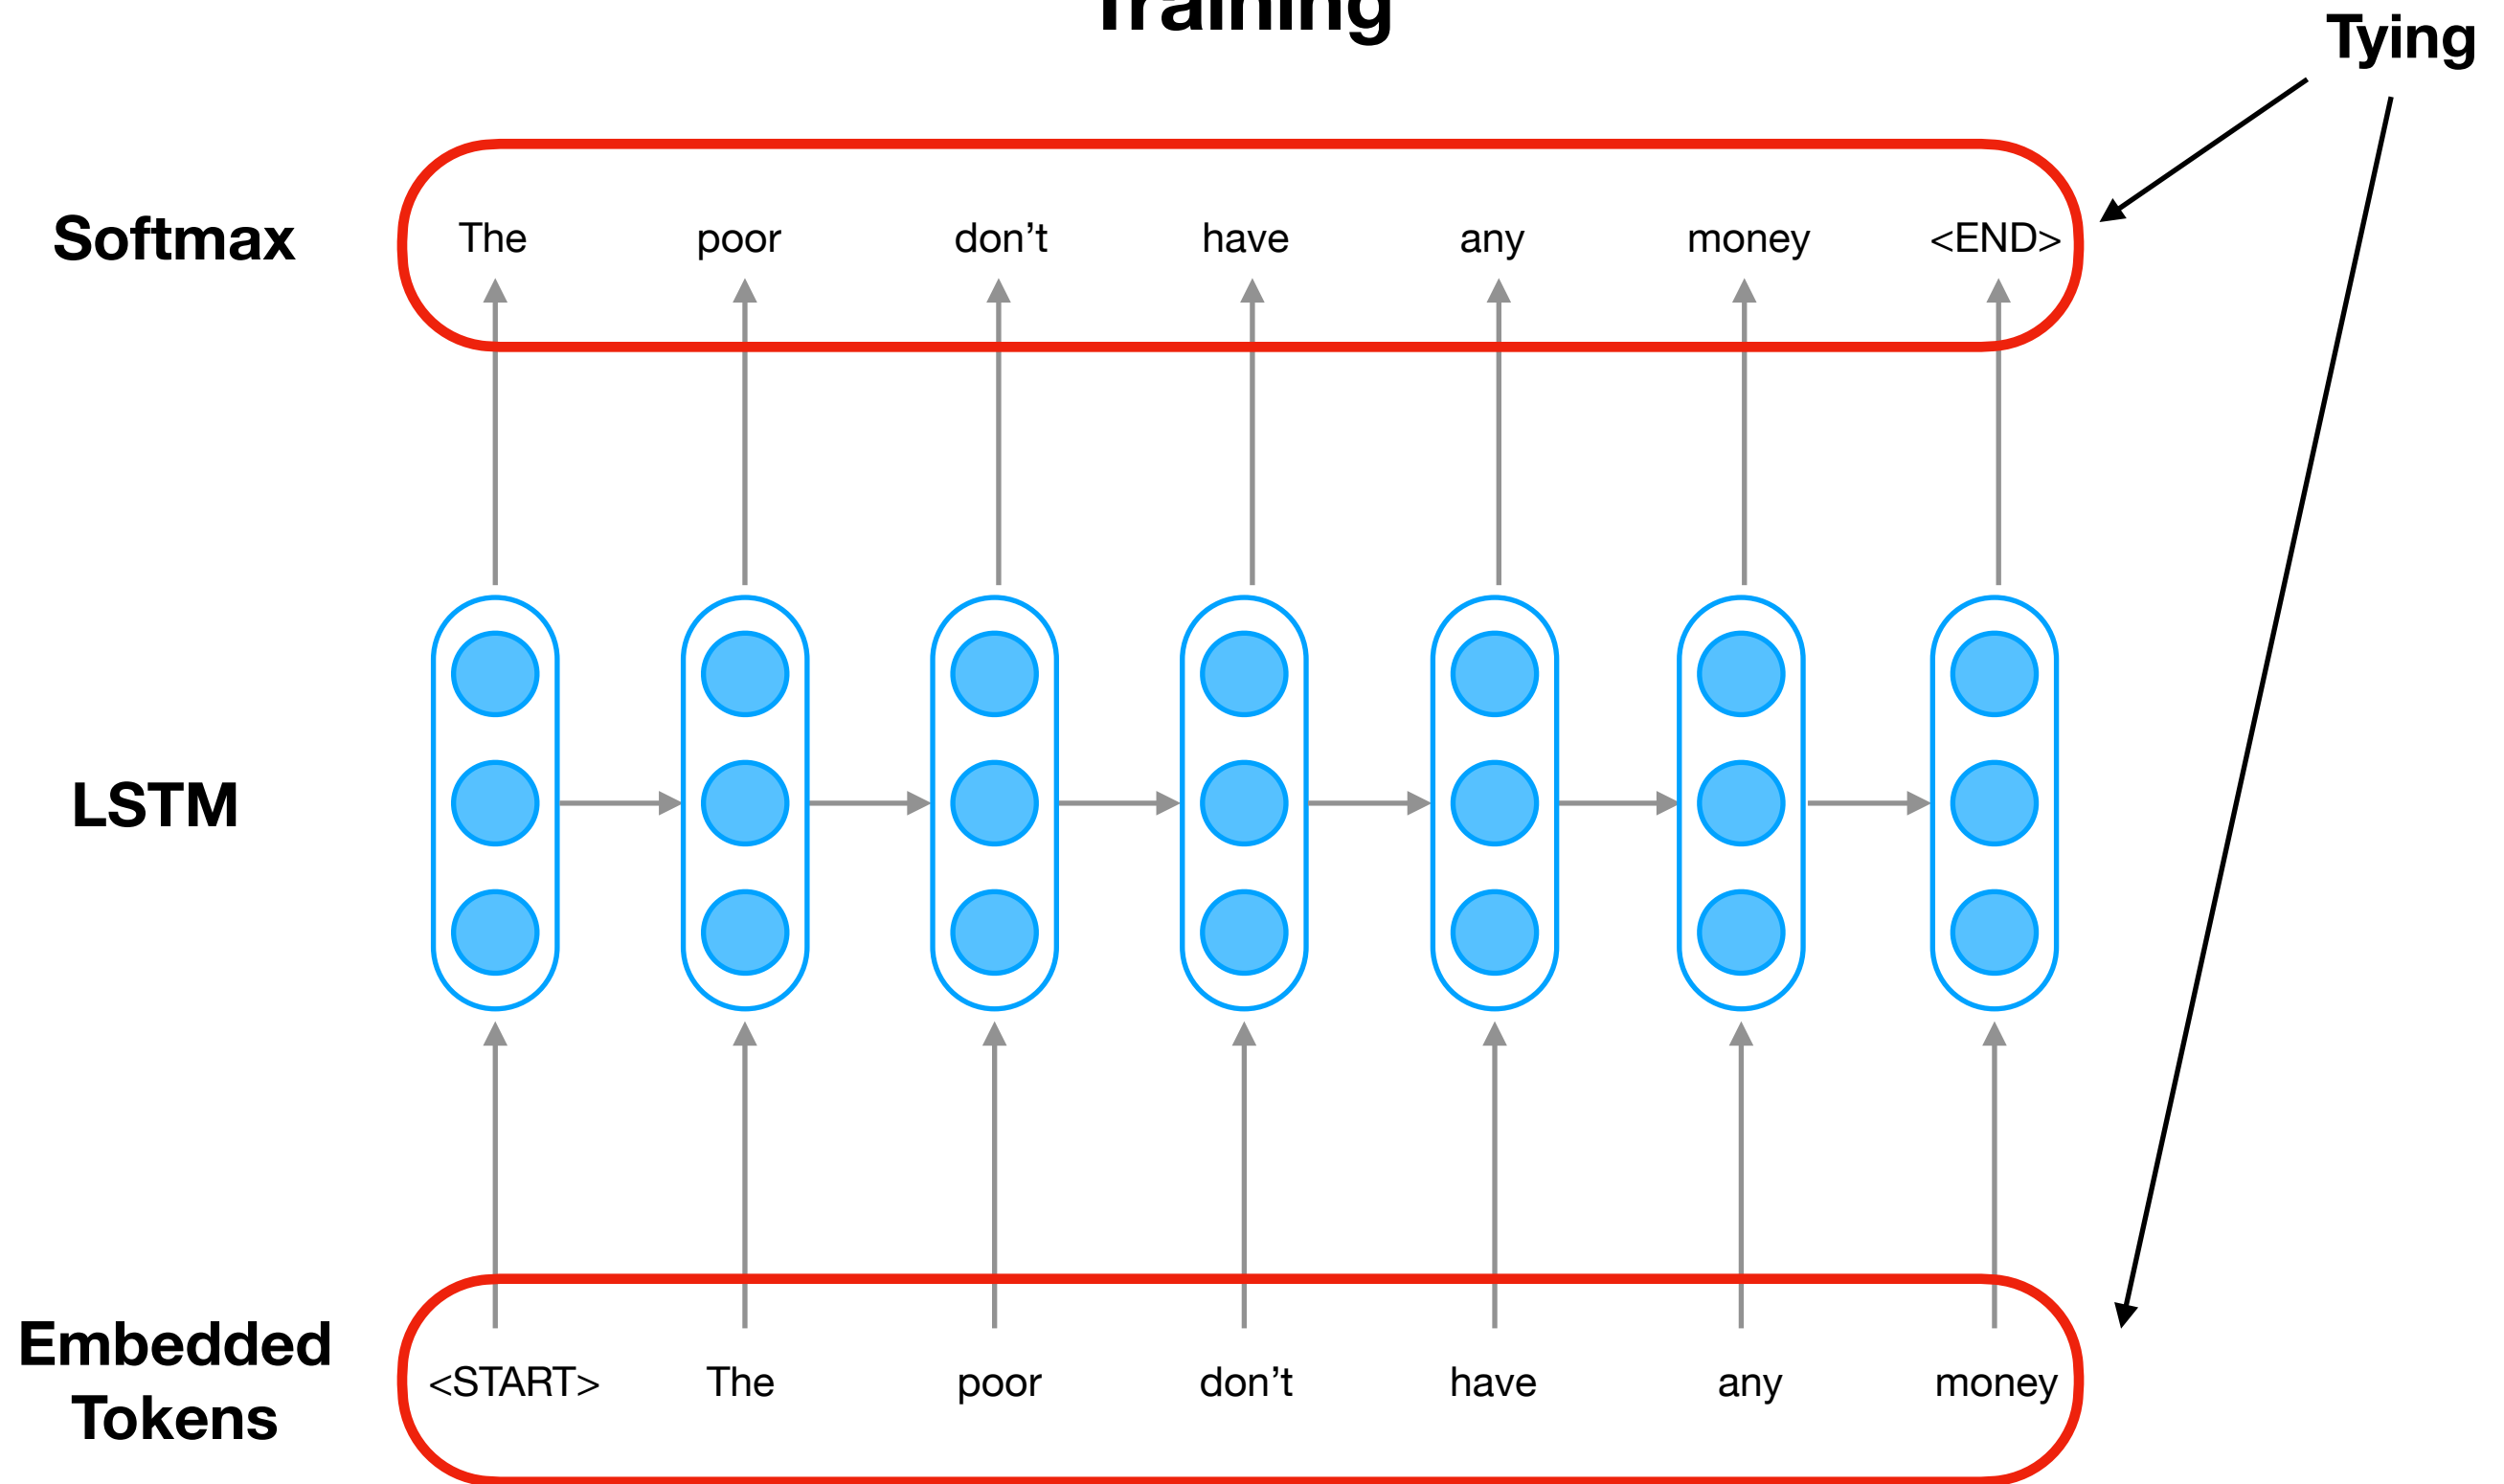


Language Model

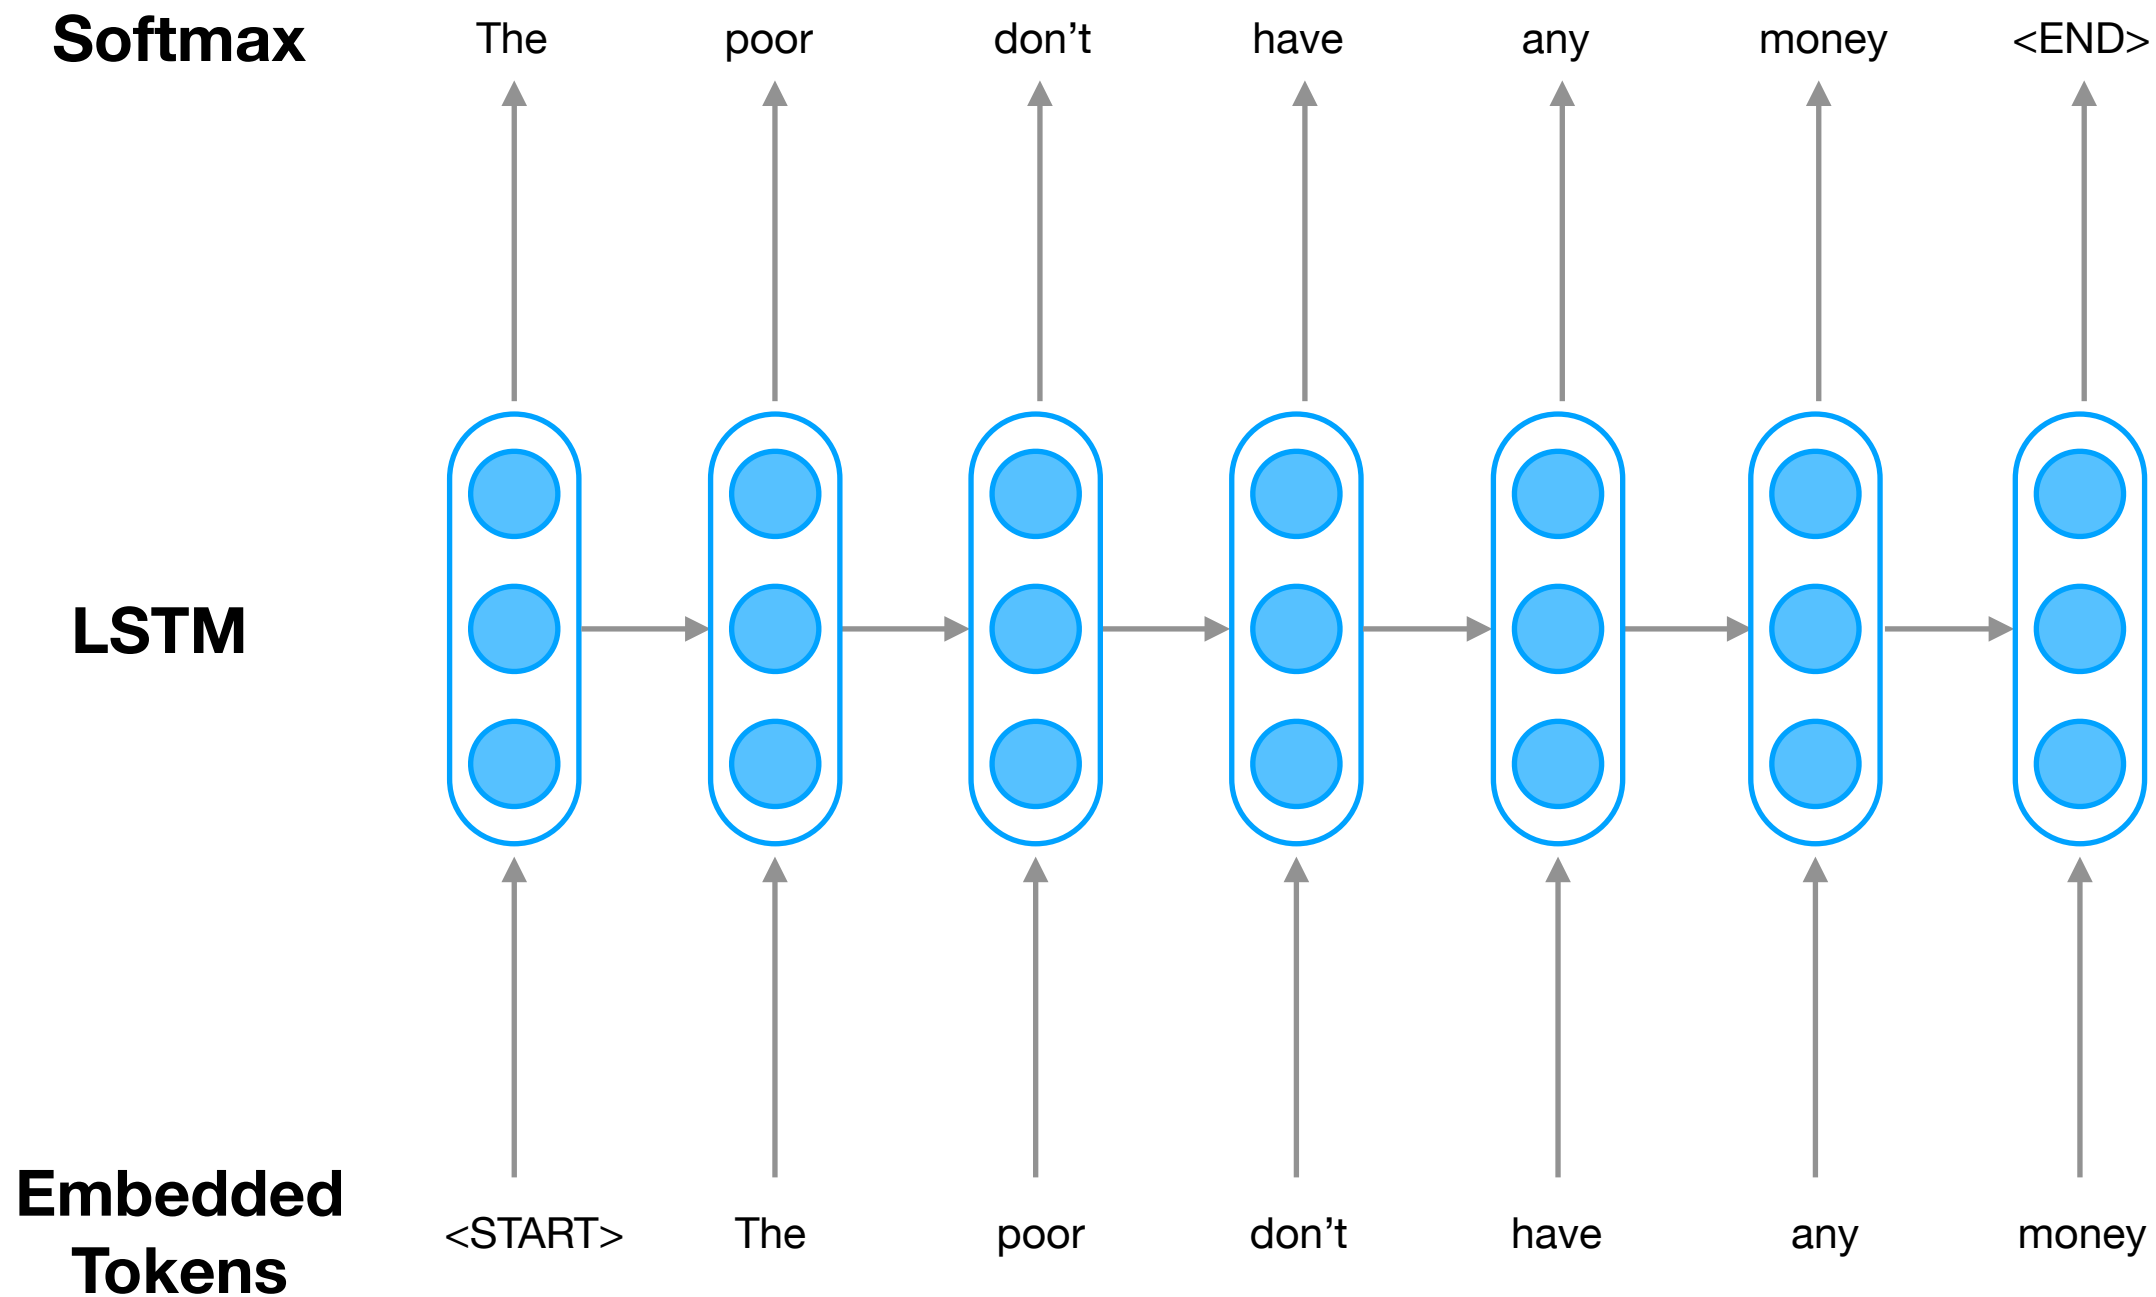


Language Model


Training



Teacher Forcing



Teacher Forcing



Language Model

<START> The poor don't have any money <END>

Language Model

X

<START>

Y

The

Language Model

X

<START>

The

Y

poor

Language Model

X

<START>

The

poor

Y

don't

Language Model

X

<START>

The

poor

don't

Y

have

Language Model

X

<START>

The

poor

don't

have

Y

any

Language Model

X	<START>	The	poor	don't	have	any	
Y							money

Language Model

X <START> The poor don't have any money

Y <END>

Language Model

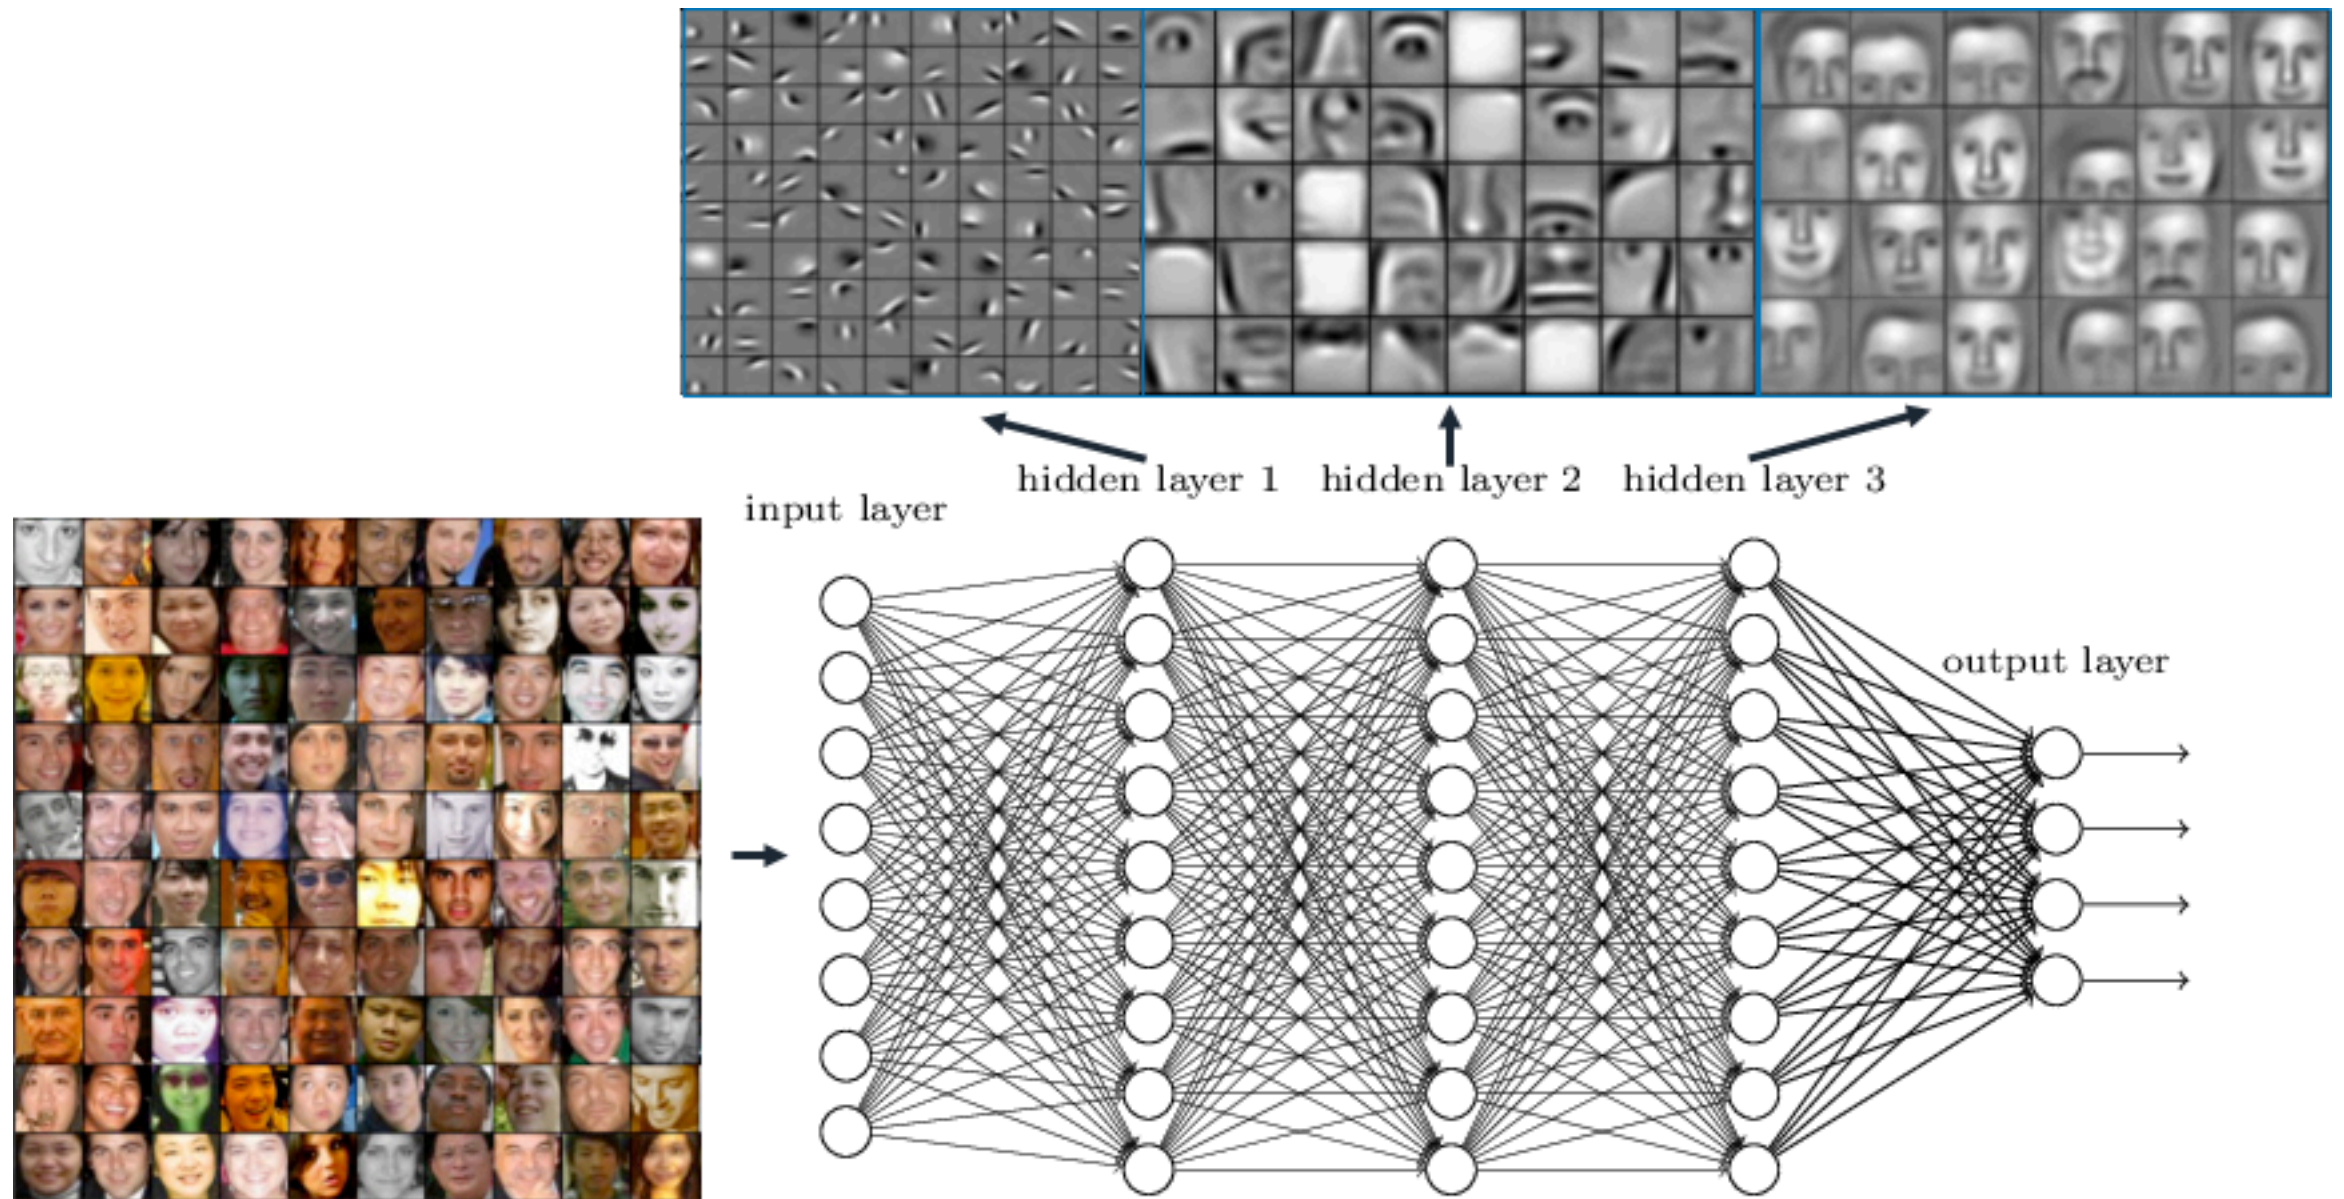
X <START> The poor don't have any money

Y <END>


Sum of Losses

Transfer Learning


CNN Intuition




NLP LM Transfer Learning




NLP LM Transfer Learning



NLP LM Transfer Learning



NLP LM Transfer Learning




NLP LM Transfer Learning

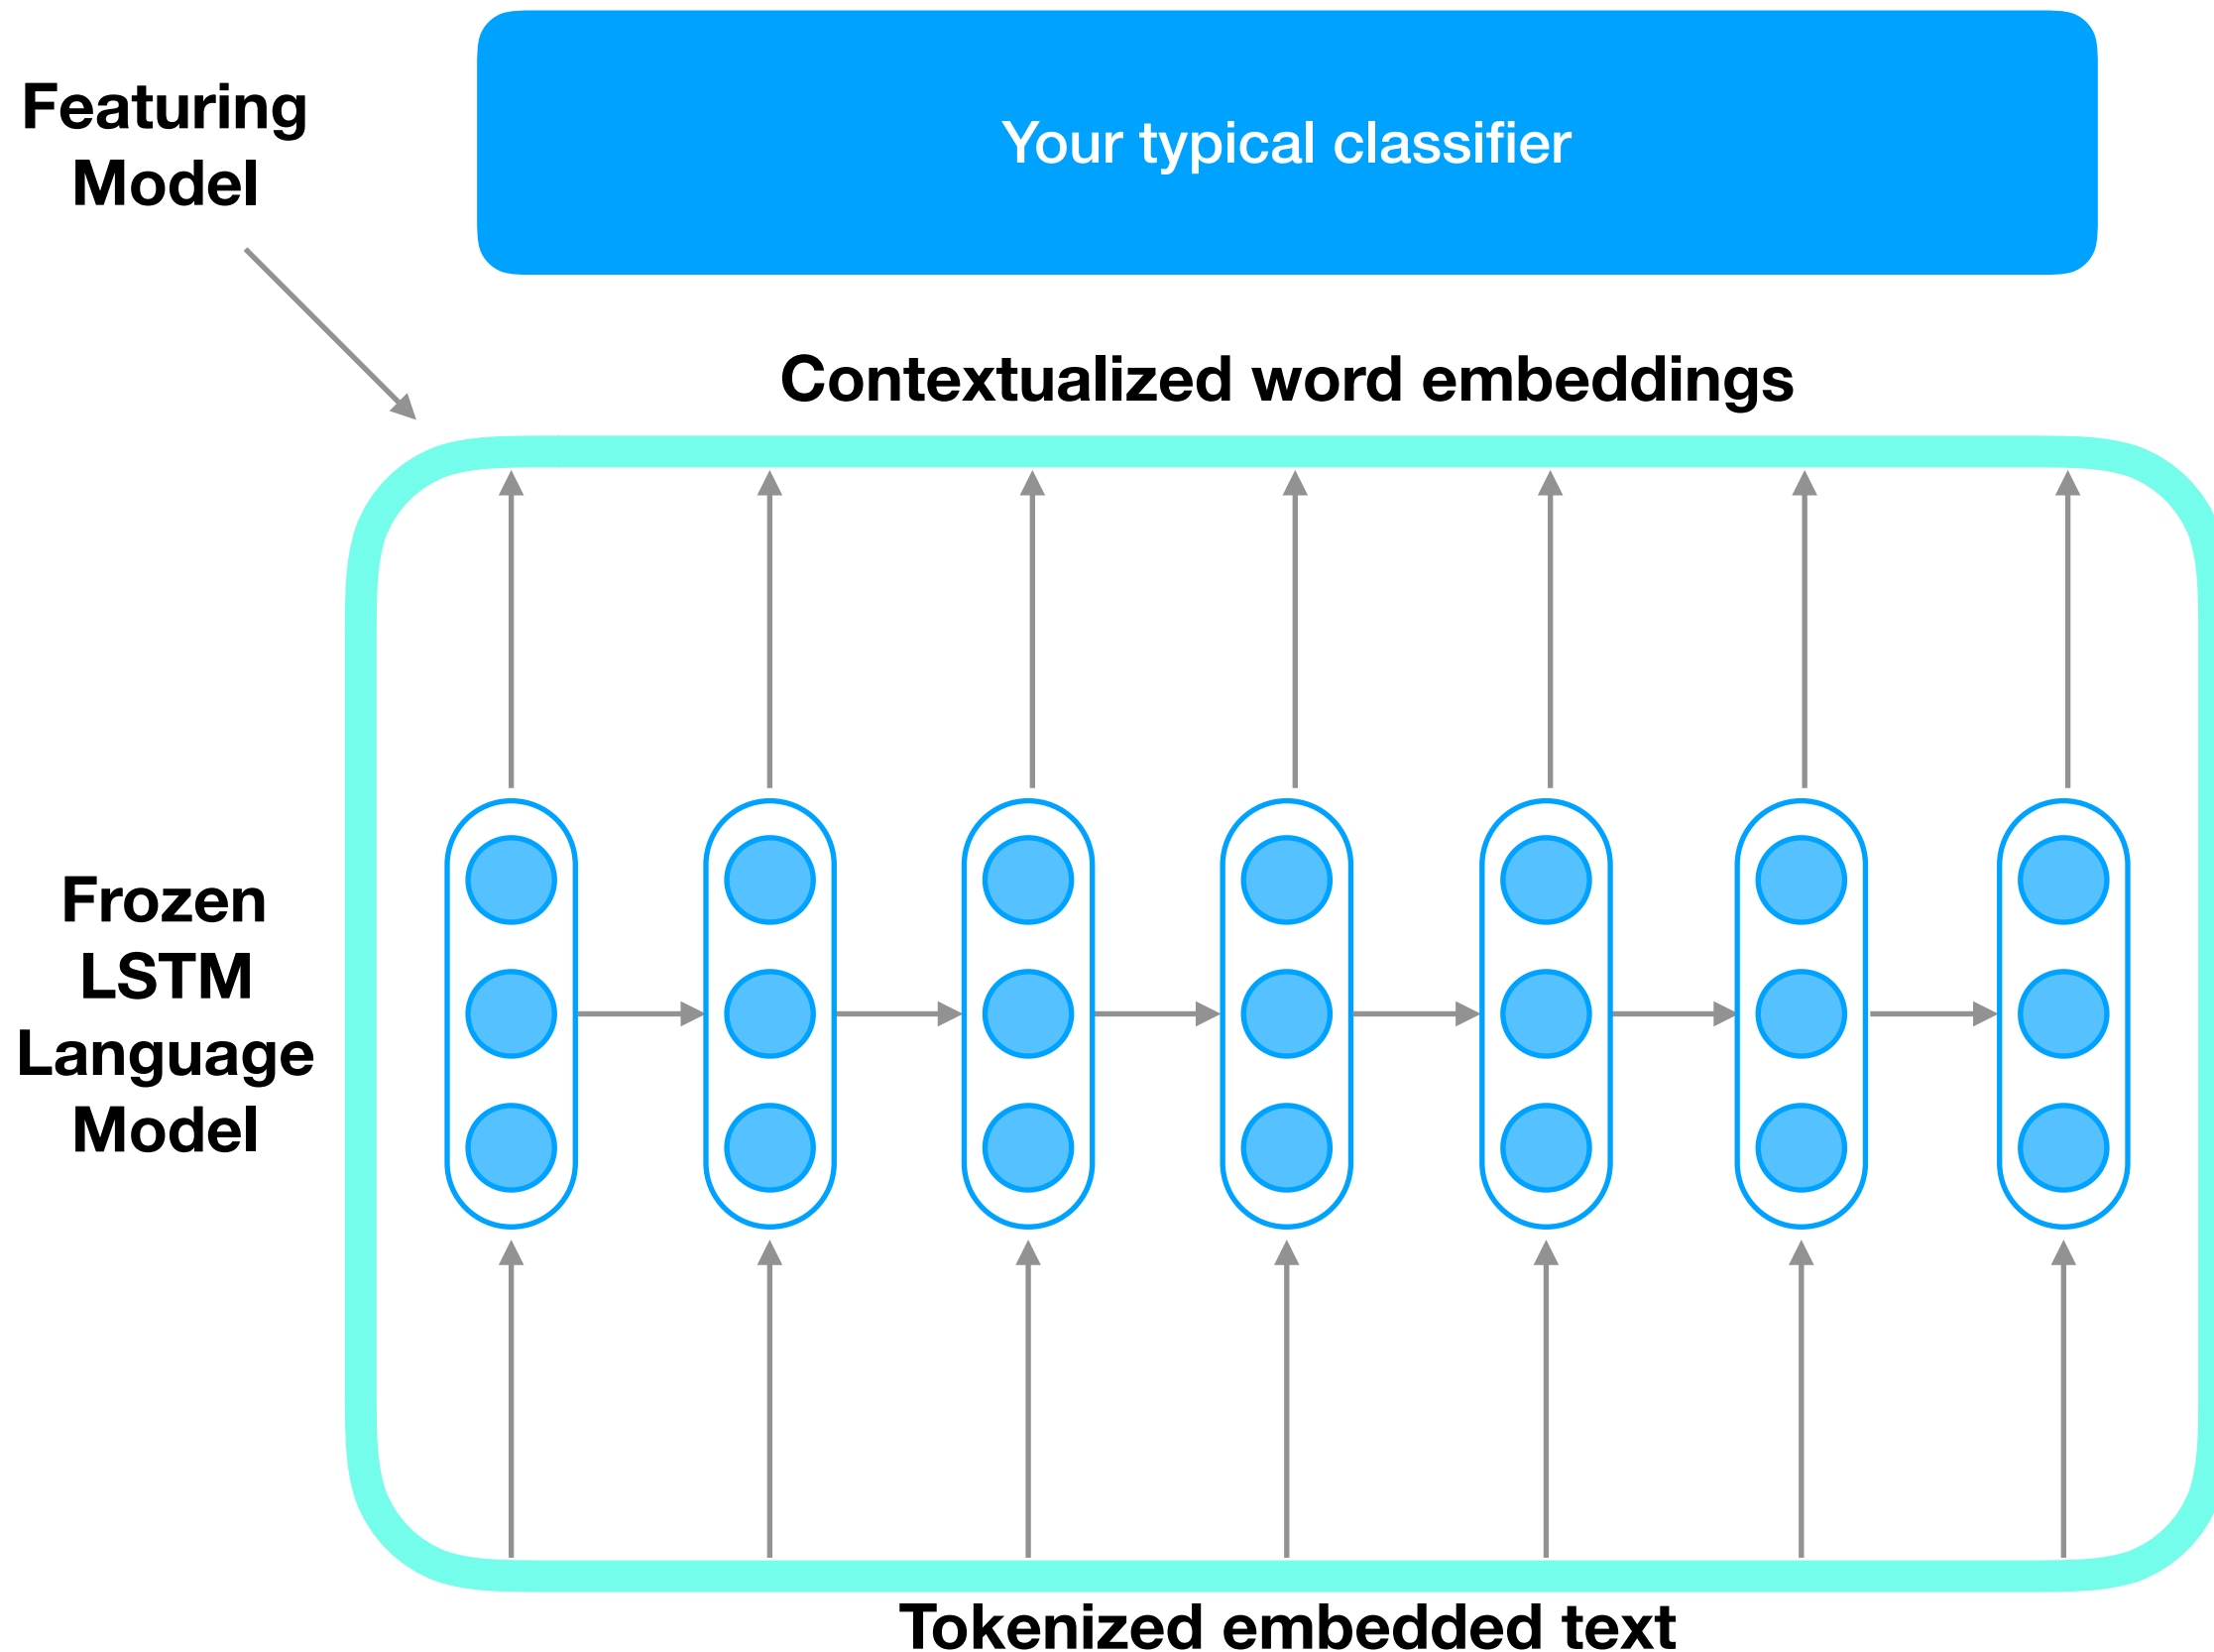
Your typical classifier

Contextualized word embeddings

Frozen
LSTM
Language
Model



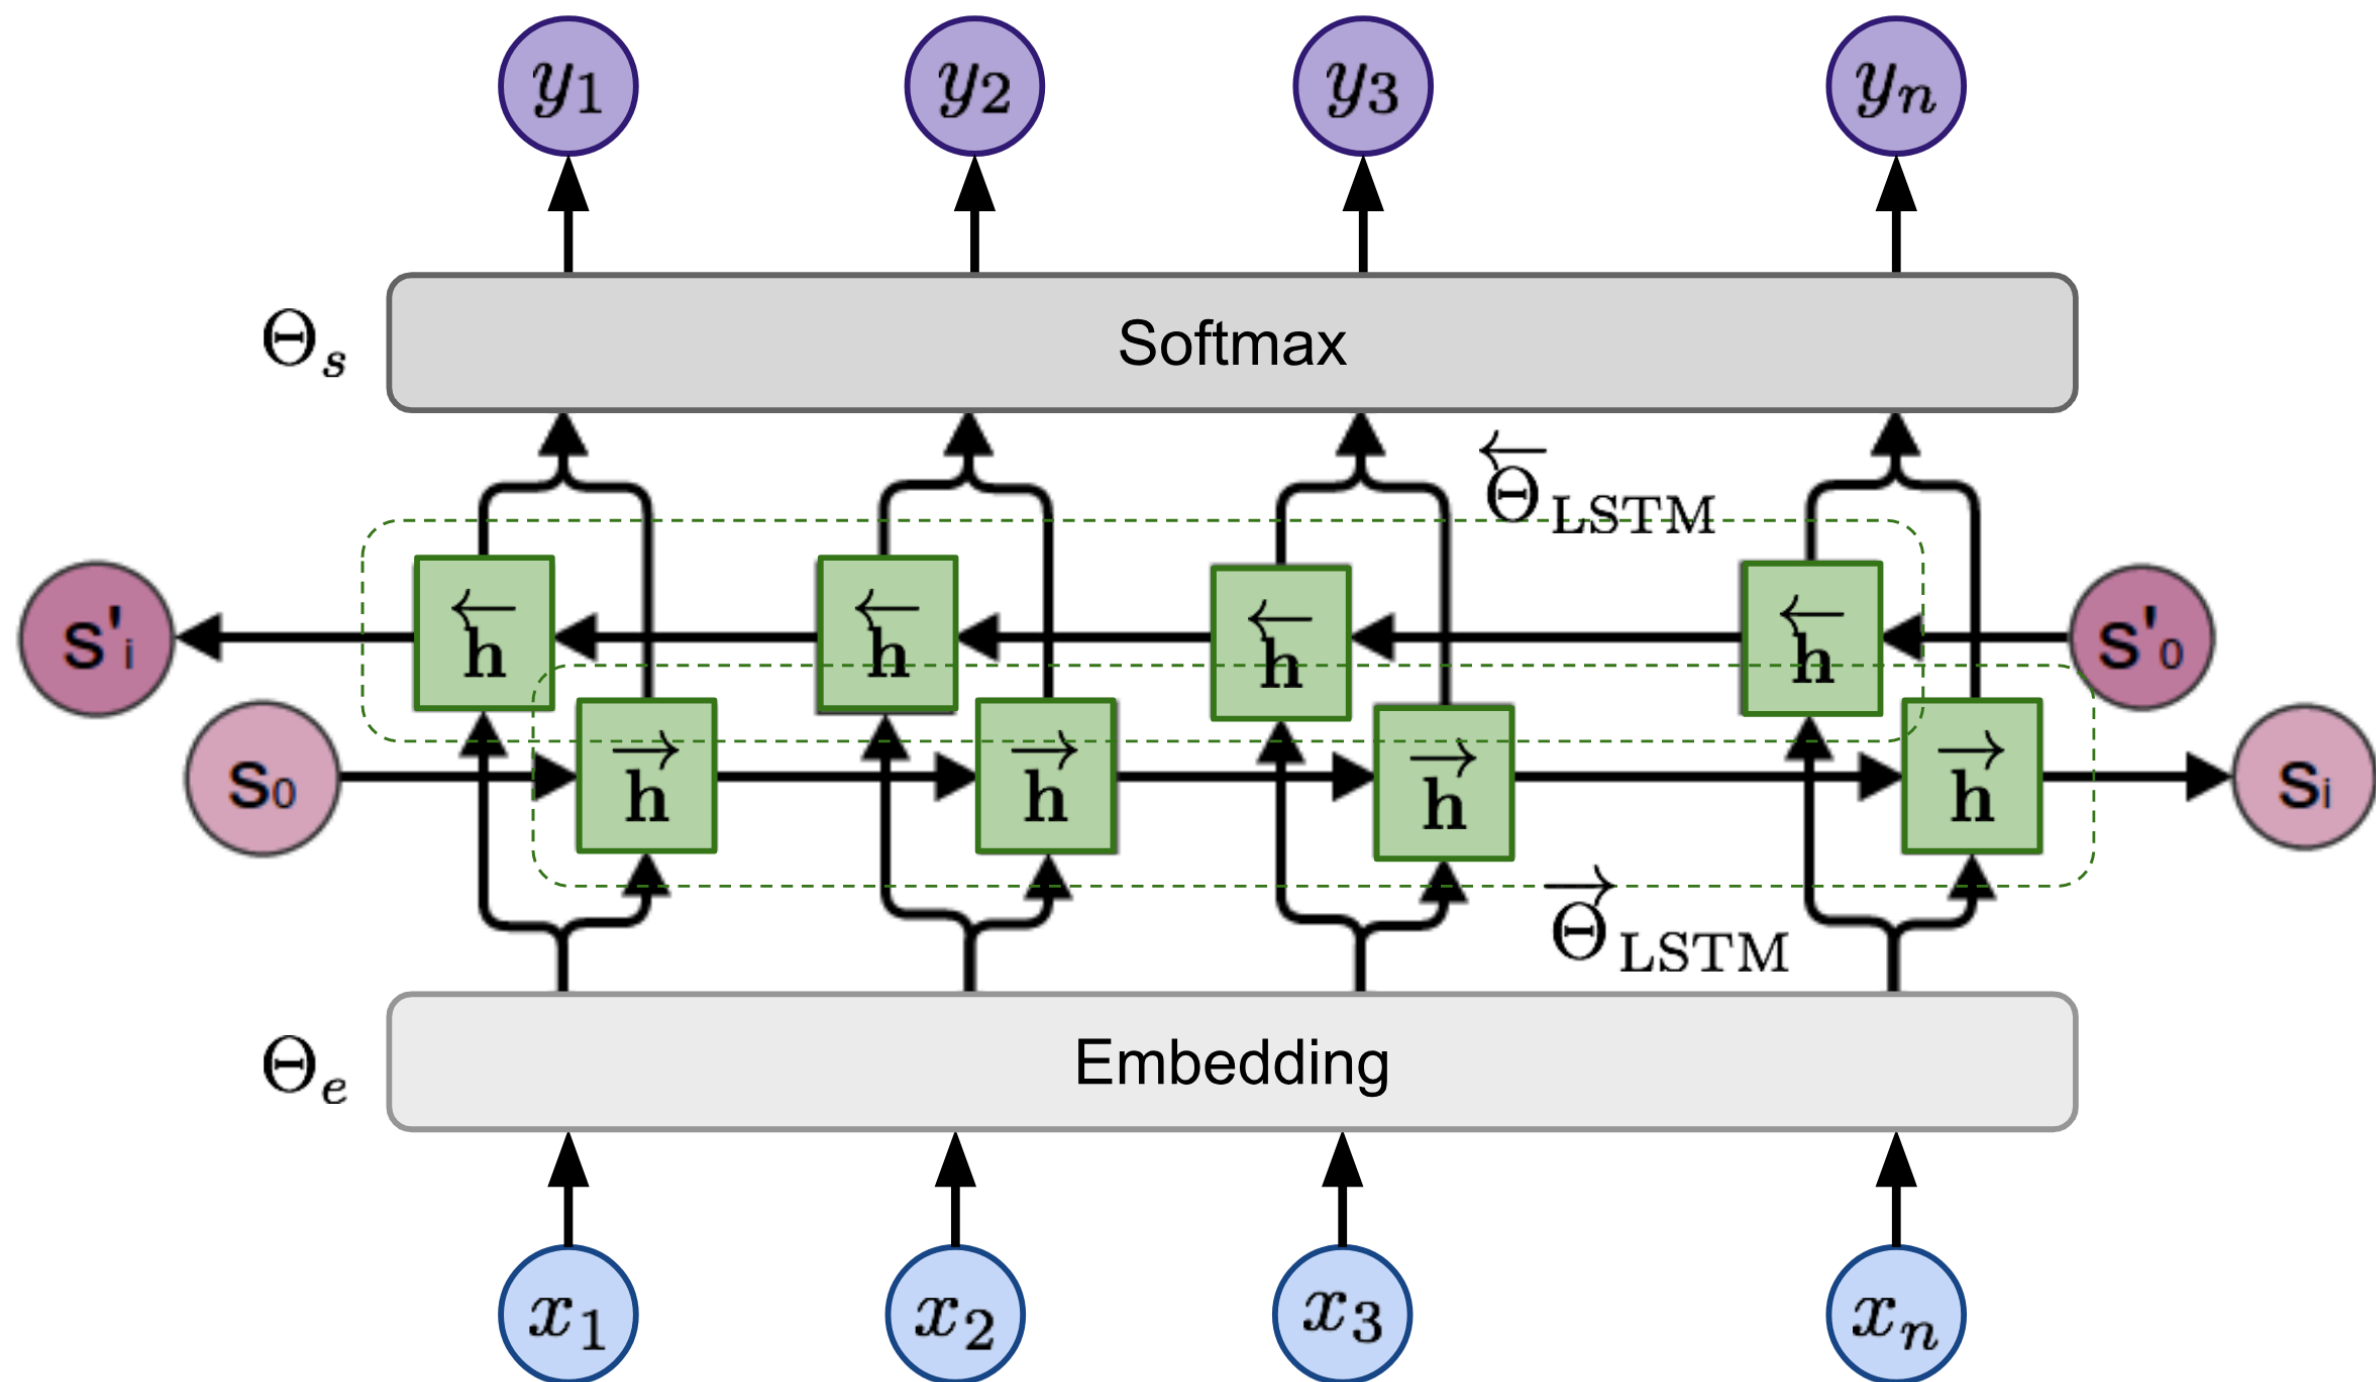
NLP LM Transfer Learning



ELMo




ELMo




ELMo

2 Forward Layers
2 Backward Layers



ELMo



ELMo

Source

<START>

The

poor

don't

have

any

money

<END>

ELMo

Source

<START>

The

poor

don't

have

any

money

<END>

Forward

Backward

ELMo

Source

<START>

The

poor

don't

have

any

money

<END>

Forward

<START>

Backward

<END>

ELMo

Source

<START> The poor don't have any money <END>

Forward

<START> The

Backward

<END> money

ELMo

Source <START> The poor don't have any money <END>

Forward <START> The poor

Backward <END> money any

ELMo

Source <START> The poor don't have any money <END>

Forward <START> The poor don't

Backward <END> money any have

ELMo

Source <START> The poor don't have any money <END>

Forward <START> The poor don't have

Backward <END> money any have don't

ELMo

Source <START> The poor don't have any money <END>

Forward <START> The poor don't have any

Backward <END> money any have don't poor

ELMo

Source <START> The poor don't have any money <END>

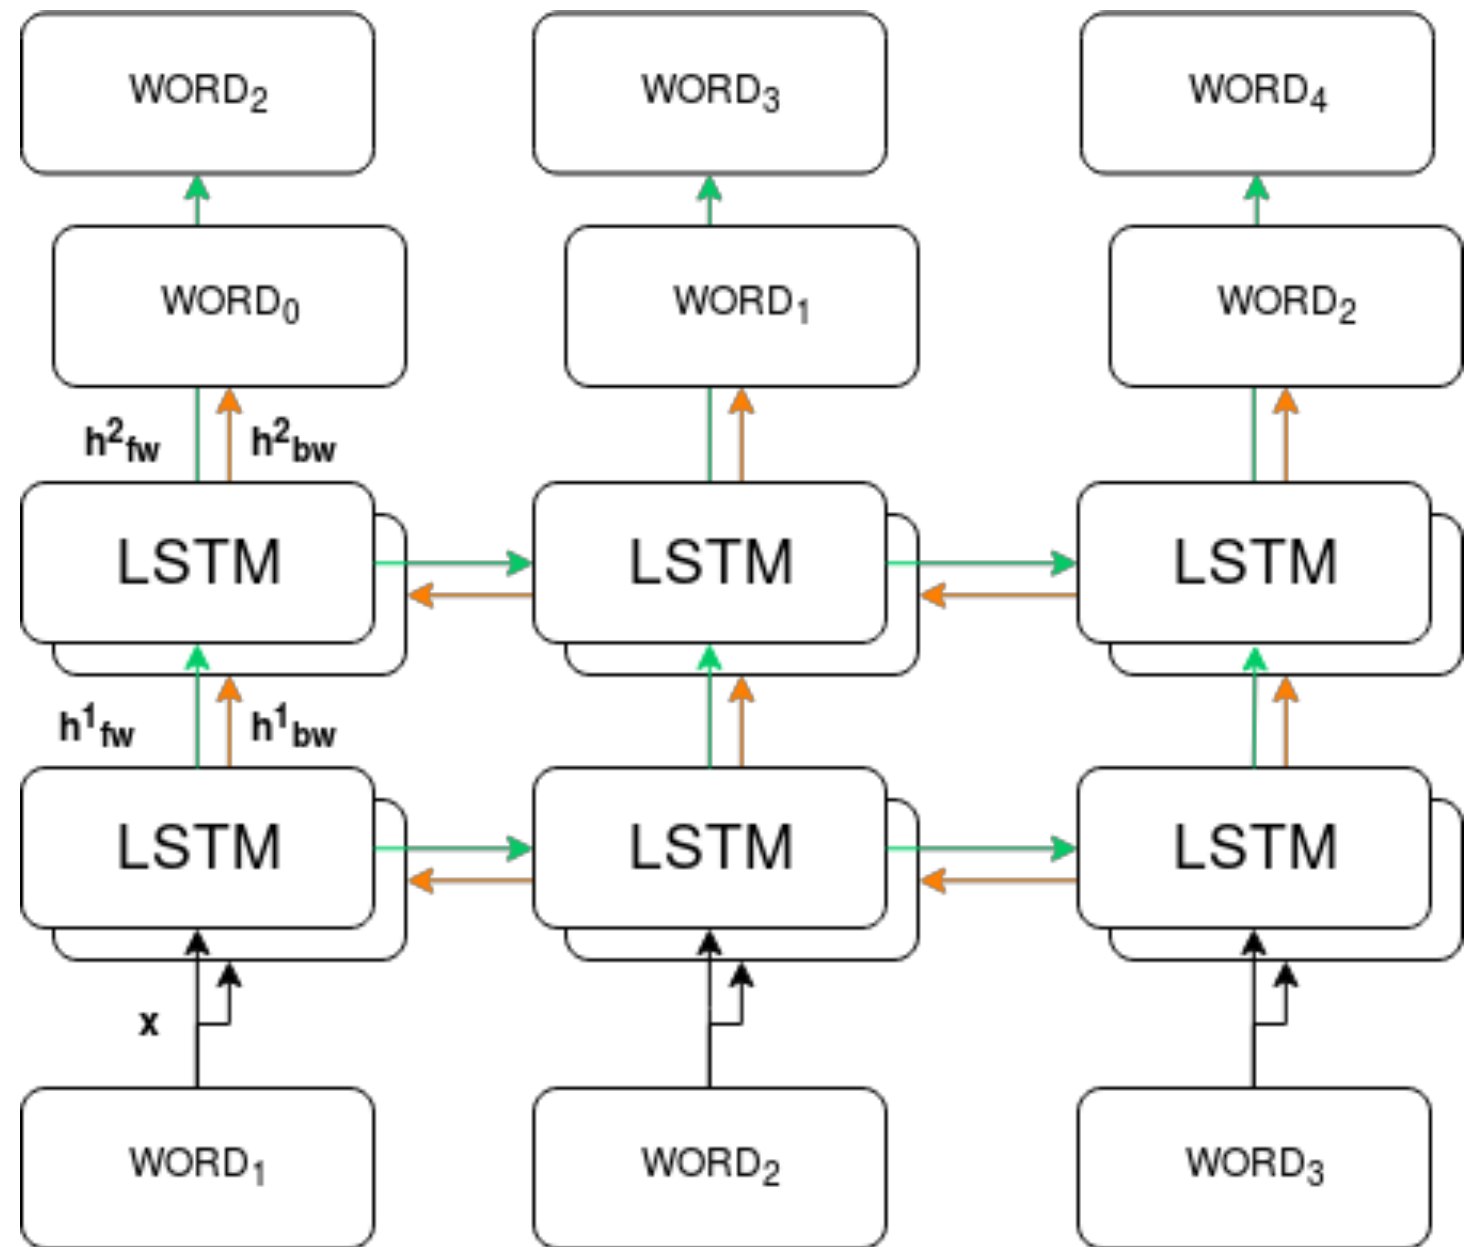
Forward <START> The poor don't have any money

Backward <END> money any have don't poor The


ELMo

Source	<START>	The	poor	don't	have	any	money	<END>
Forward	<START>	The	poor	don't	have	any	money	<END>
Backward	<END>	money	any	have	don't	poor	The	<START>

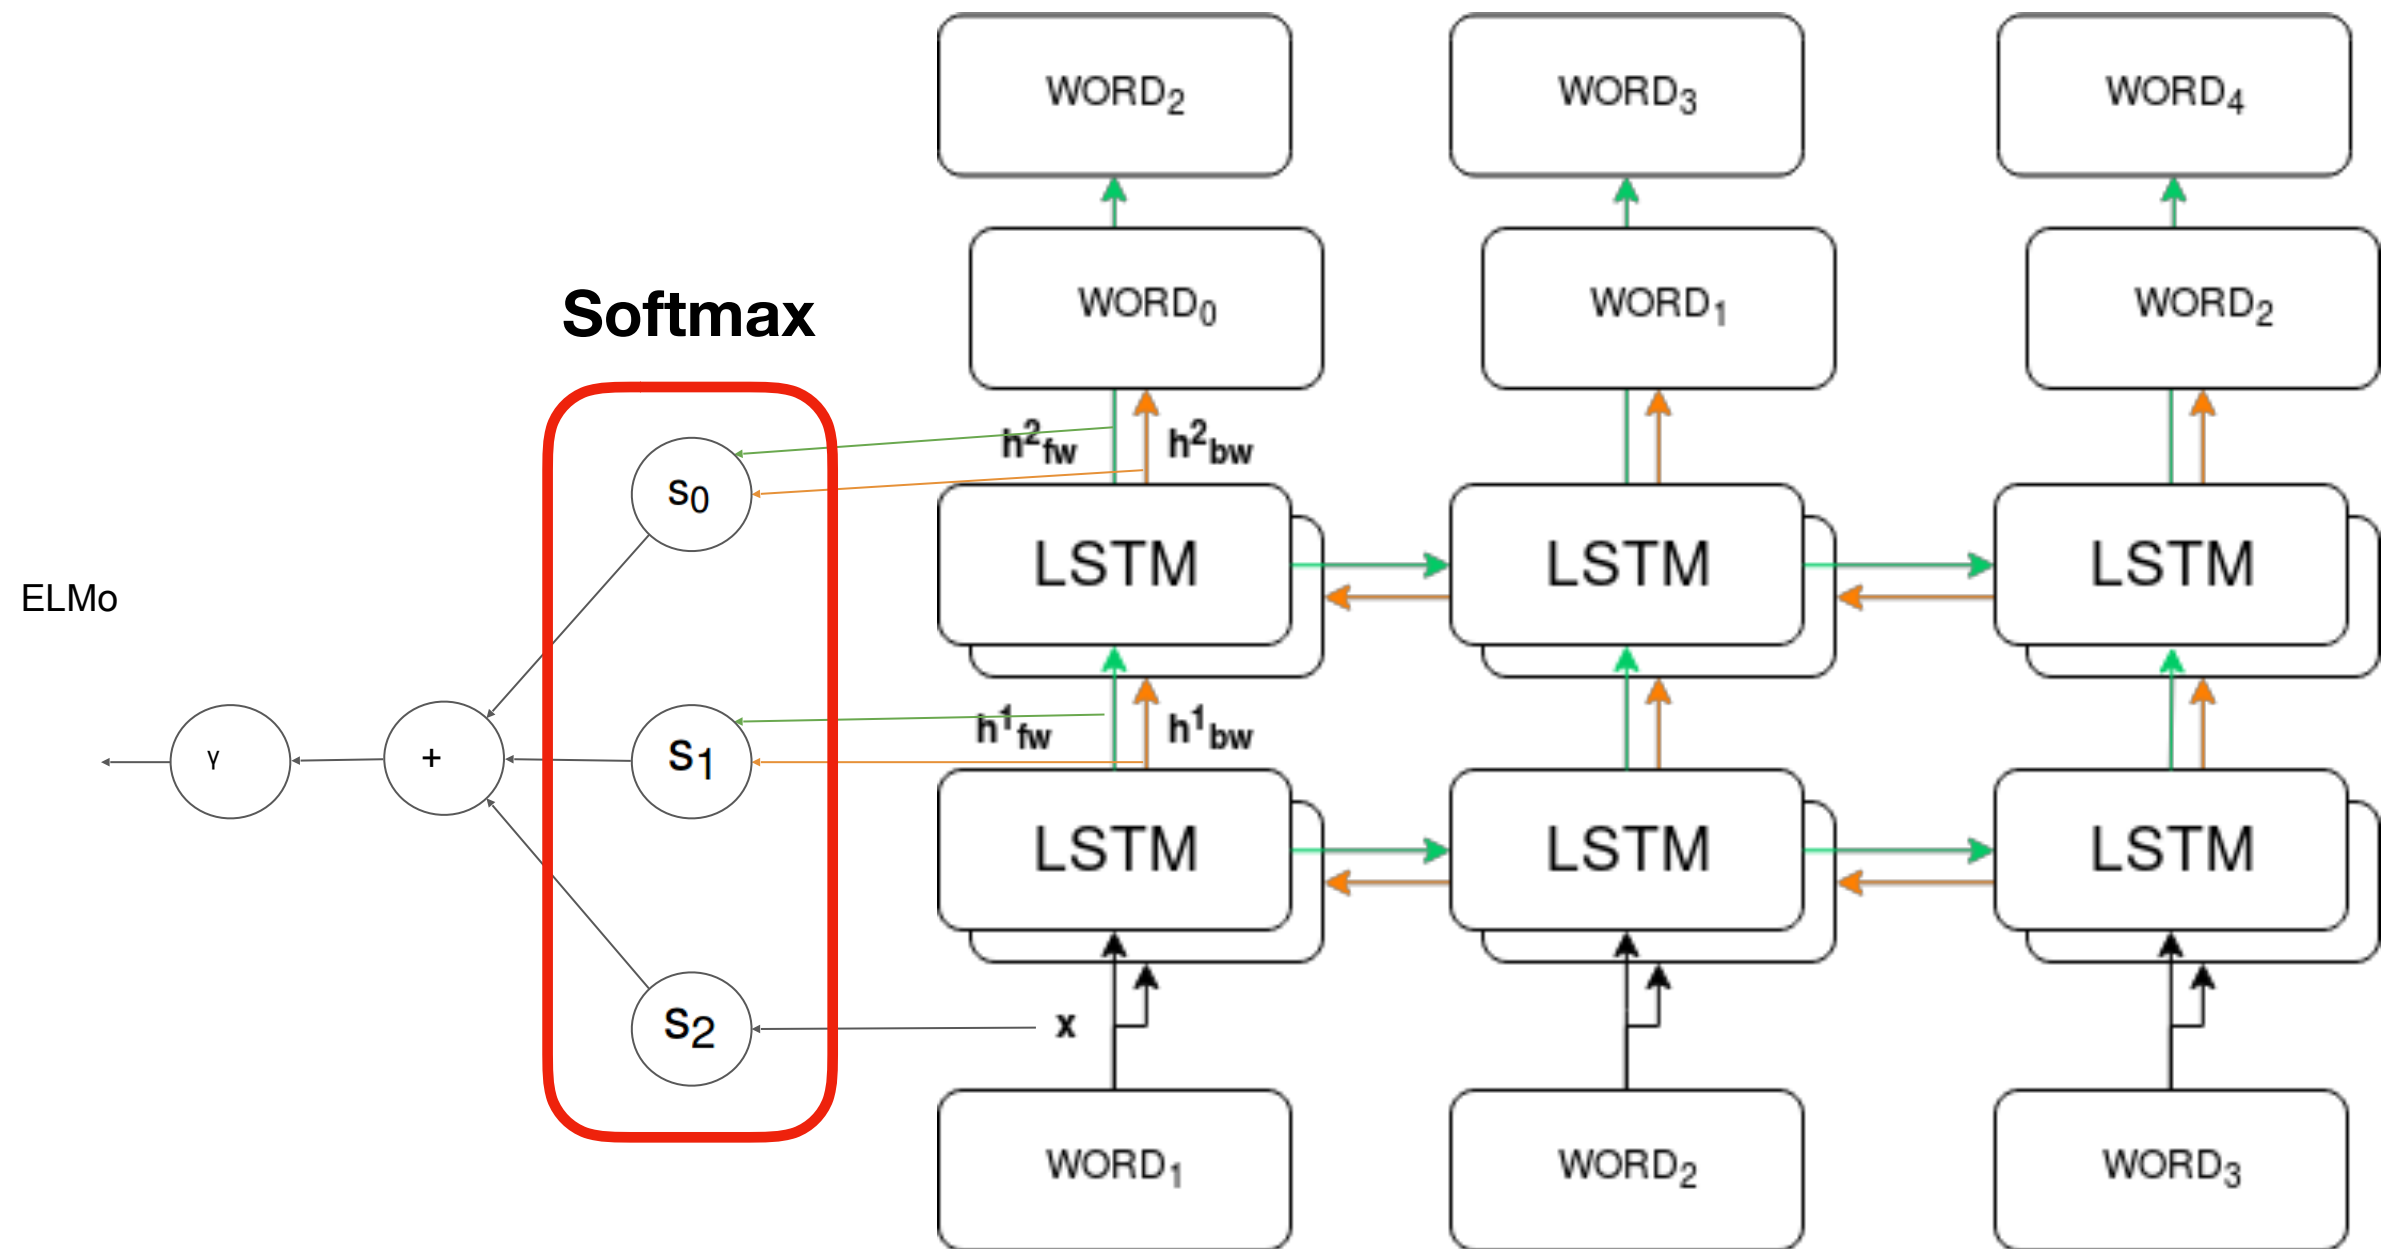
ELMo



ELMo



ELMo



ELMo

Results





Figure 2: Visualization of softmax normalized biLM layer weights across tasks and ELMo locations. Normalized weights less than $1/3$ are hatched with horizontal lines and those greater than $2/3$ are speckled.

Tokenization

Закон Ципфа



Tokenization

Word level

- + Small text length
- Big vocabulary size
- OOV

Character level

- Long text
- + Small vocabulary size
- + Almost no OOV

Tokenization

Word level

- + Small text length
- Big vocabulary size
- OOV

Character level

- Long text
- + Small vocabulary size
- + Almost no OOV

1. Word level

i'm a second year student in an ivy league school ->

["i'm", 'a', 'second', 'year', 'student', 'in', 'an', 'ivy', 'league', 'school']

2. Character level

['i', "'", 'm', "'", 'a', "'", 's', 'e', 'c', 'o', 'n', 'd', "'", 'y', 'e', 'a', 'r', "'", 's', 't', 'u', 'd', 'e', 'n', 't', "'", 'i', 'n', "'", 'a', 'n', "'", 'i', 'v', 'y', "'", 'l', 'e', 'a', 'g', 'u', 'e', "'", 's', 'c', 'h', 'o', 'o', 'l']

ВРЕ

I saw a girl with a telescope. ->

['__I', '__saw', '__a', '__girl', '__with', '__a', '__', 'te', 'le', 's', 'c', 'o', 'pe', '.']

опубликовано видео убитого саудовского журналиста джамаля хашкуджи ->

['__опубликовано', '__видео', '__убитого', '__саудов', 'ского', '__журналиста',
['__джама', 'ля', '__ха', 'шку', 'джи']

BPE

Algorithm 1 Learn BPE operations

```
import re, collections

def get_stats(vocab):
    pairs = collections.defaultdict(int)
    for word, freq in vocab.items():
        symbols = word.split()
        for i in range(len(symbols)-1):
            pairs[symbols[i], symbols[i+1]] += freq
    return pairs

def merge_vocab(pair, v_in):
    v_out = {}
    bigram = re.escape(' '.join(pair))
    p = re.compile(r'(?!\S)' + bigram + r'(?!\S)')
    for word in v_in:
        w_out = p.sub(''.join(pair), word)
        v_out[w_out] = v_in[word]
    return v_out

vocab = {'l o w </w>' : 5, 'l o w e r </w>' : 2,
        'n e w e s t </w>':6, 'w i d e s t </w>':3}
num_merges = 10
for i in range(num_merges):
    pairs = get_stats(vocab)
    best = max(pairs, key=pairs.get)
    vocab = merge_vocab(best, vocab)
    print(best)
```

r ·	→	r·
l o	→	lo
l o w	→	low
e r·	→	er·

Figure 1: BPE merge operations learned from dictionary {‘low’, ‘lowest’, ‘newer’, ‘wider’}.

BPE

Algorithm 1 Learn BPE operations

```
import re, collections

def get_stats(vocab):
    pairs = collections.defaultdict(int)
    for word, freq in vocab.items():
        symbols = word.split()
        for i in range(len(symbols)-1):
            pairs[symbols[i], symbols[i+1]] += freq
    return pairs

def merge_vocab(pair, v_in):
    v_out = {}
    bigram = re.escape(' '.join(pair))
    p = re.compile(r'(?!\S)' + bigram + r'(?!\S)')
    for word in v_in:
        w_out = p.sub(''.join(pair), word)
        v_out[w_out] = v_in[word]
    return v_out

vocab = {'l o w </w>' : 5, 'l o w e r </w>' : 2,
        'n e w e s t </w>':6, 'w i d e s t </w>':3}
num_merges = 10
for i in range(num_merges):
    pairs = get_stats(vocab)
    best = max(pairs, key=pairs.get)
    vocab = merge_vocab(best, vocab)
    print(best)
```

- learning
 - word:freq : {low:5, lowest:2, newer:6, wider:3}
 - marge & count
 1. 'r' '</w>' : 9 → marge'r</w>'
 2. 'e' 'r</w>' : 9 → marge'er</w>'
 3. 'l' 'o' : 7 → marge'lo'
 4. 'lo' 'w' : 7 → marge'low'

→ OOV : 'lower' segmented 'low er</w>'

r ·	→	r ·
l o	→	l o
l o w	→	l o w
e r ·	→	e r ·

Figure 1: BPE merge operations learned from dictionary {'low', 'lowest', 'newer', 'wider'}.

BPE

Algorithm 1 Learn BPE operations

```
import re, collections

def get_stats(vocab):
    pairs = collections.defaultdict(int)
    for word, freq in vocab.items():
        symbols = word.split()
        for i in range(len(symbols)-1):
            pairs[symbols[i], symbols[i+1]] += freq
    return pairs

def merge_vocab(pair, v_in):
    v_out = {}
    bigram = re.escape(' '.join(pair))
    p = re.compile(r'(?!\S)' + bigram + r'(?!\S)')
    for word in v_in:
        w_out = p.sub(''.join(pair), word)
        v_out[w_out] = v_in[word]
    return v_out

vocab = {'l o w </w>' : 5, 'l o w e r </w>' : 2,
        'n e w e s t </w>' : 6, 'w i d e s t </w>' : 3}
num_merges = 10
for i in range(num_merges):
    pairs = get_stats(vocab)
    best = max(pairs, key=pairs.get)
    vocab = merge_vocab(best, vocab)
    print(best)
```

r ·	→	r·
l o	→	lo
l o w	→	low
e r·	→	er·

- learning
 - word:freq : {low:5, lowest:2, newer:6, wider:3}
 - marge & count
 1. 'r' '</w>' : 9 → marge'r</w>'
 2. 'e' 'r</w>' : 9 → marge'er</w>'
 3. 'l' 'o' : 7 → marge'lo'
 4. 'lo' 'w' : 7 → marge'low'

→ OOV : 'lower' segmented 'low er</w>'

Vocabulary sizes:

5000, 10000, 15000, ..., 50000

Figure 1: BPE merge operations learned from dictionary {'low', 'lowest', 'newer', 'wider'}.

BPE

SentencePiece

build failing build passing coverage 98% issues 35 open code quality A pypi package 0.1.83 contributions welcome
License Apache 2.0

SentencePiece is an unsupervised text tokenizer and detokenizer mainly for Neural Network-based text generation systems where the vocabulary size is predetermined prior to the neural model training. SentencePiece implements **subword units** (e.g., **byte-pair-encoding (BPE)** [[Sennrich et al.](#)]) and **unigram language model** [[Kudo.](#)]) with the extension of direct training from raw sentences. SentencePiece allows us to make a purely end-to-end system that does not depend on language-specific pre/postprocessing.

This is not an official Google product.

YouTokenToMe

YouTokenToMe is an unsupervised text tokenizer focused on computational efficiency. It currently implements fast Byte Pair Encoding (BPE) [[Sennrich et al.](#)]. Our implementation is much faster in training and tokenization than both [fastBPE](#) and [SentencePiece](#). In some test cases, it is 90 times faster. Check out our [benchmark](#) results.

Key advantages:

- Multithreading for training and tokenization
- The algorithm has $O(N)$ complexity, where N is the length of training data
- Highly efficient implementation in C++
- Python wrapper and command-line interface

Sequence to sequence

Source sentence

Sequence to sequence

Les pauvres sont démunis

Source sentence


Sequence to sequence

Encoder


Les pauvres sont démunis

Source sentence


Sequence to sequence




Sequence to sequence




Sequence to sequence



Sequence to sequence

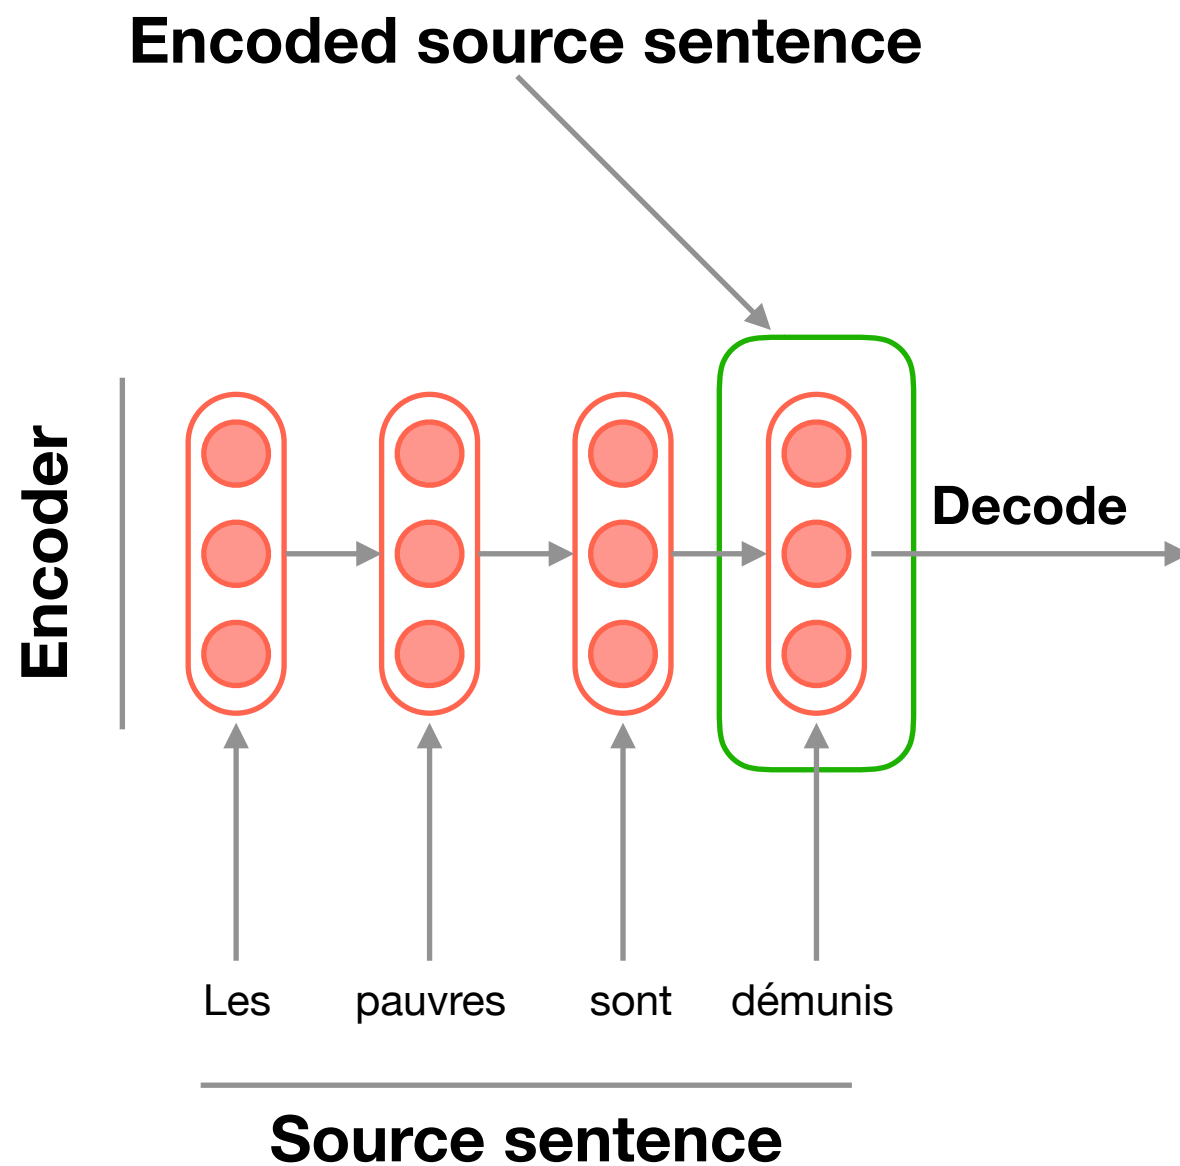


Sequence to sequence



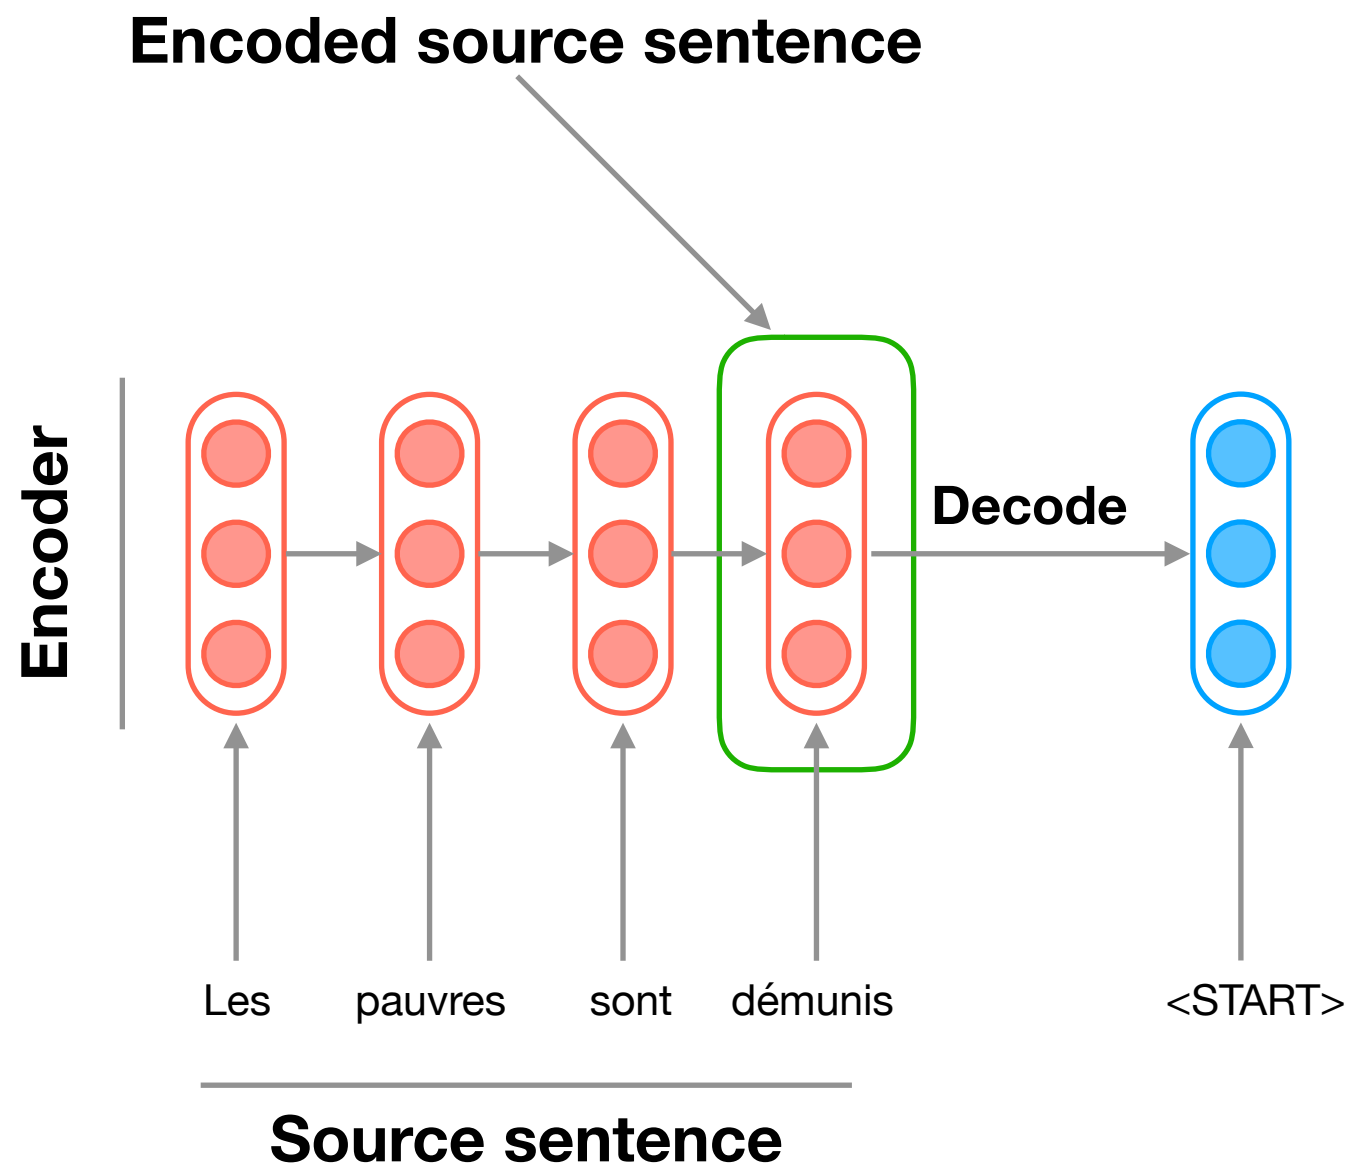
Sequence to sequence

Inference



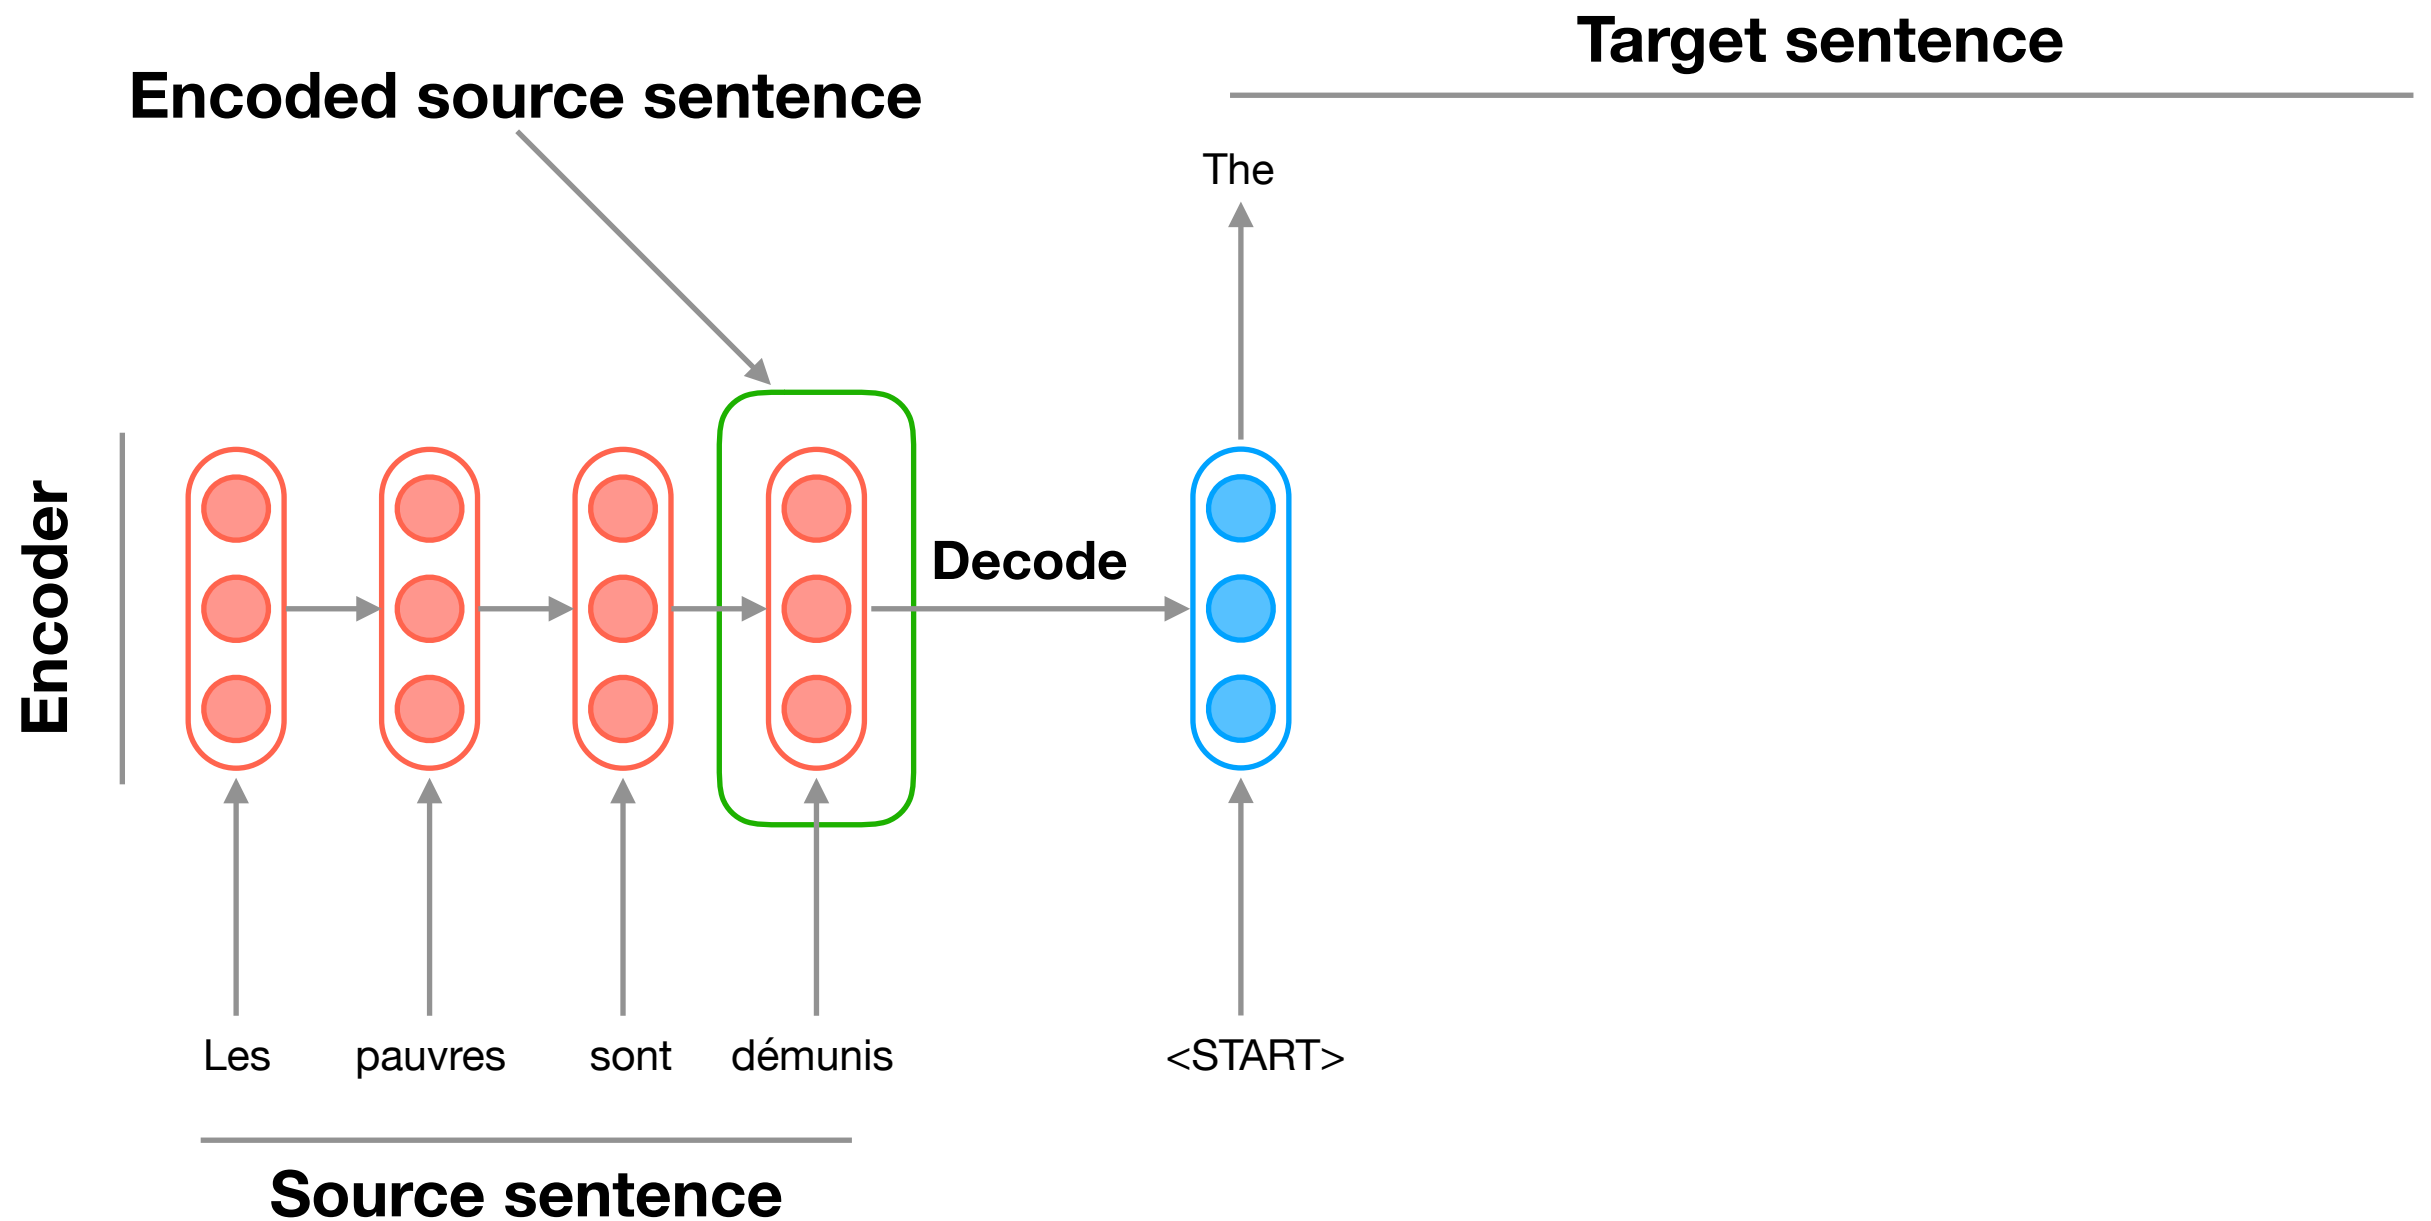
Sequence to sequence

Inference



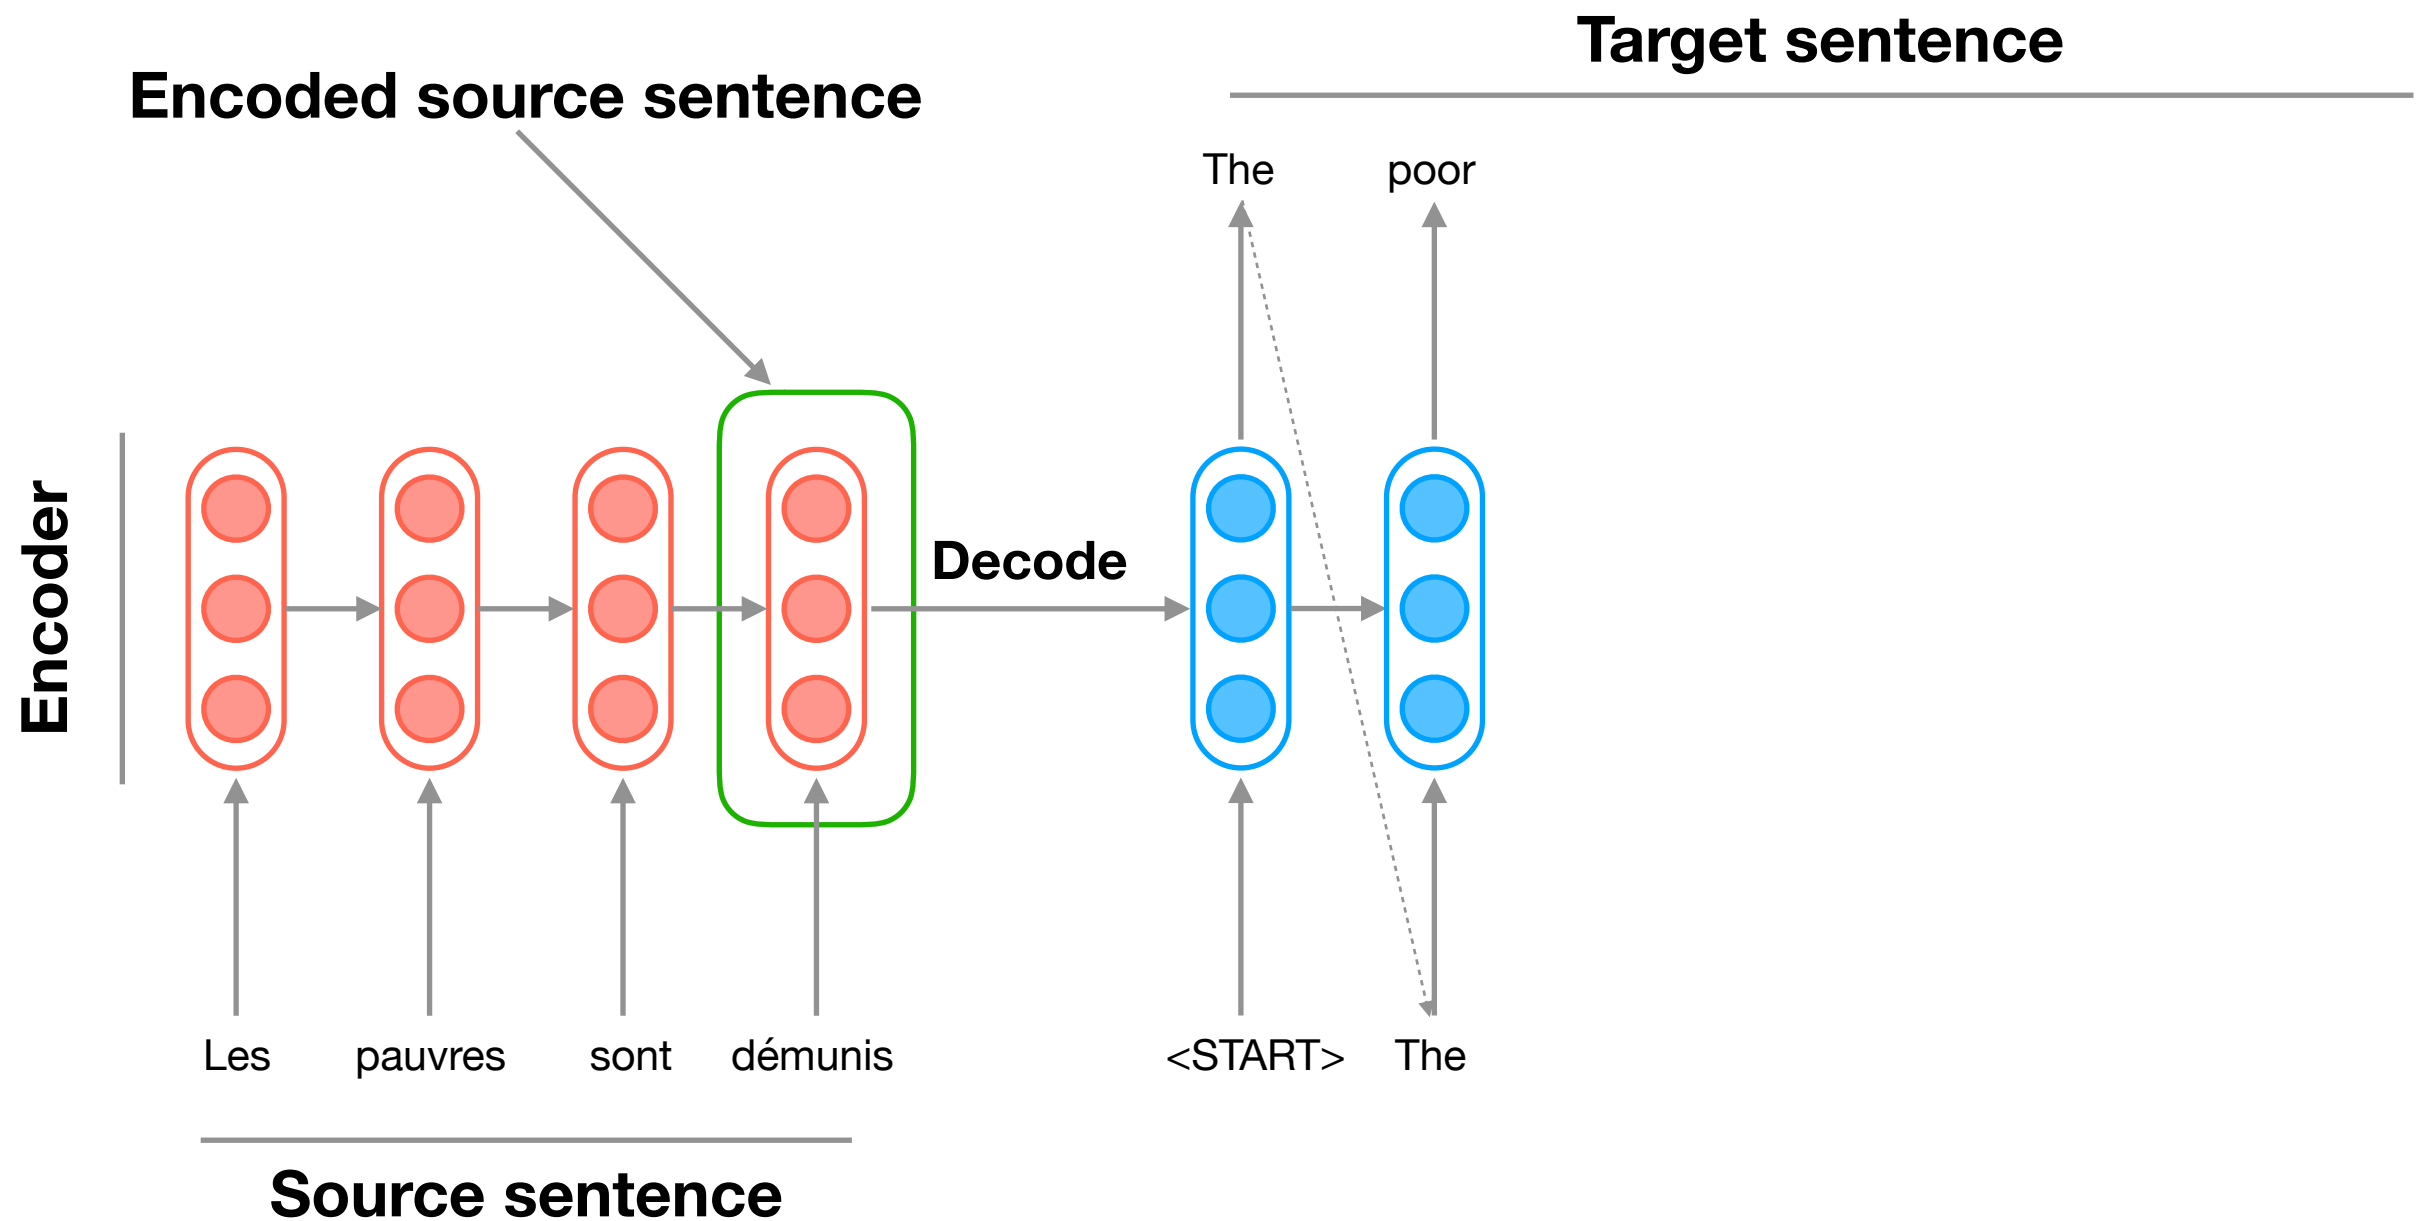
Sequence to sequence

Inference



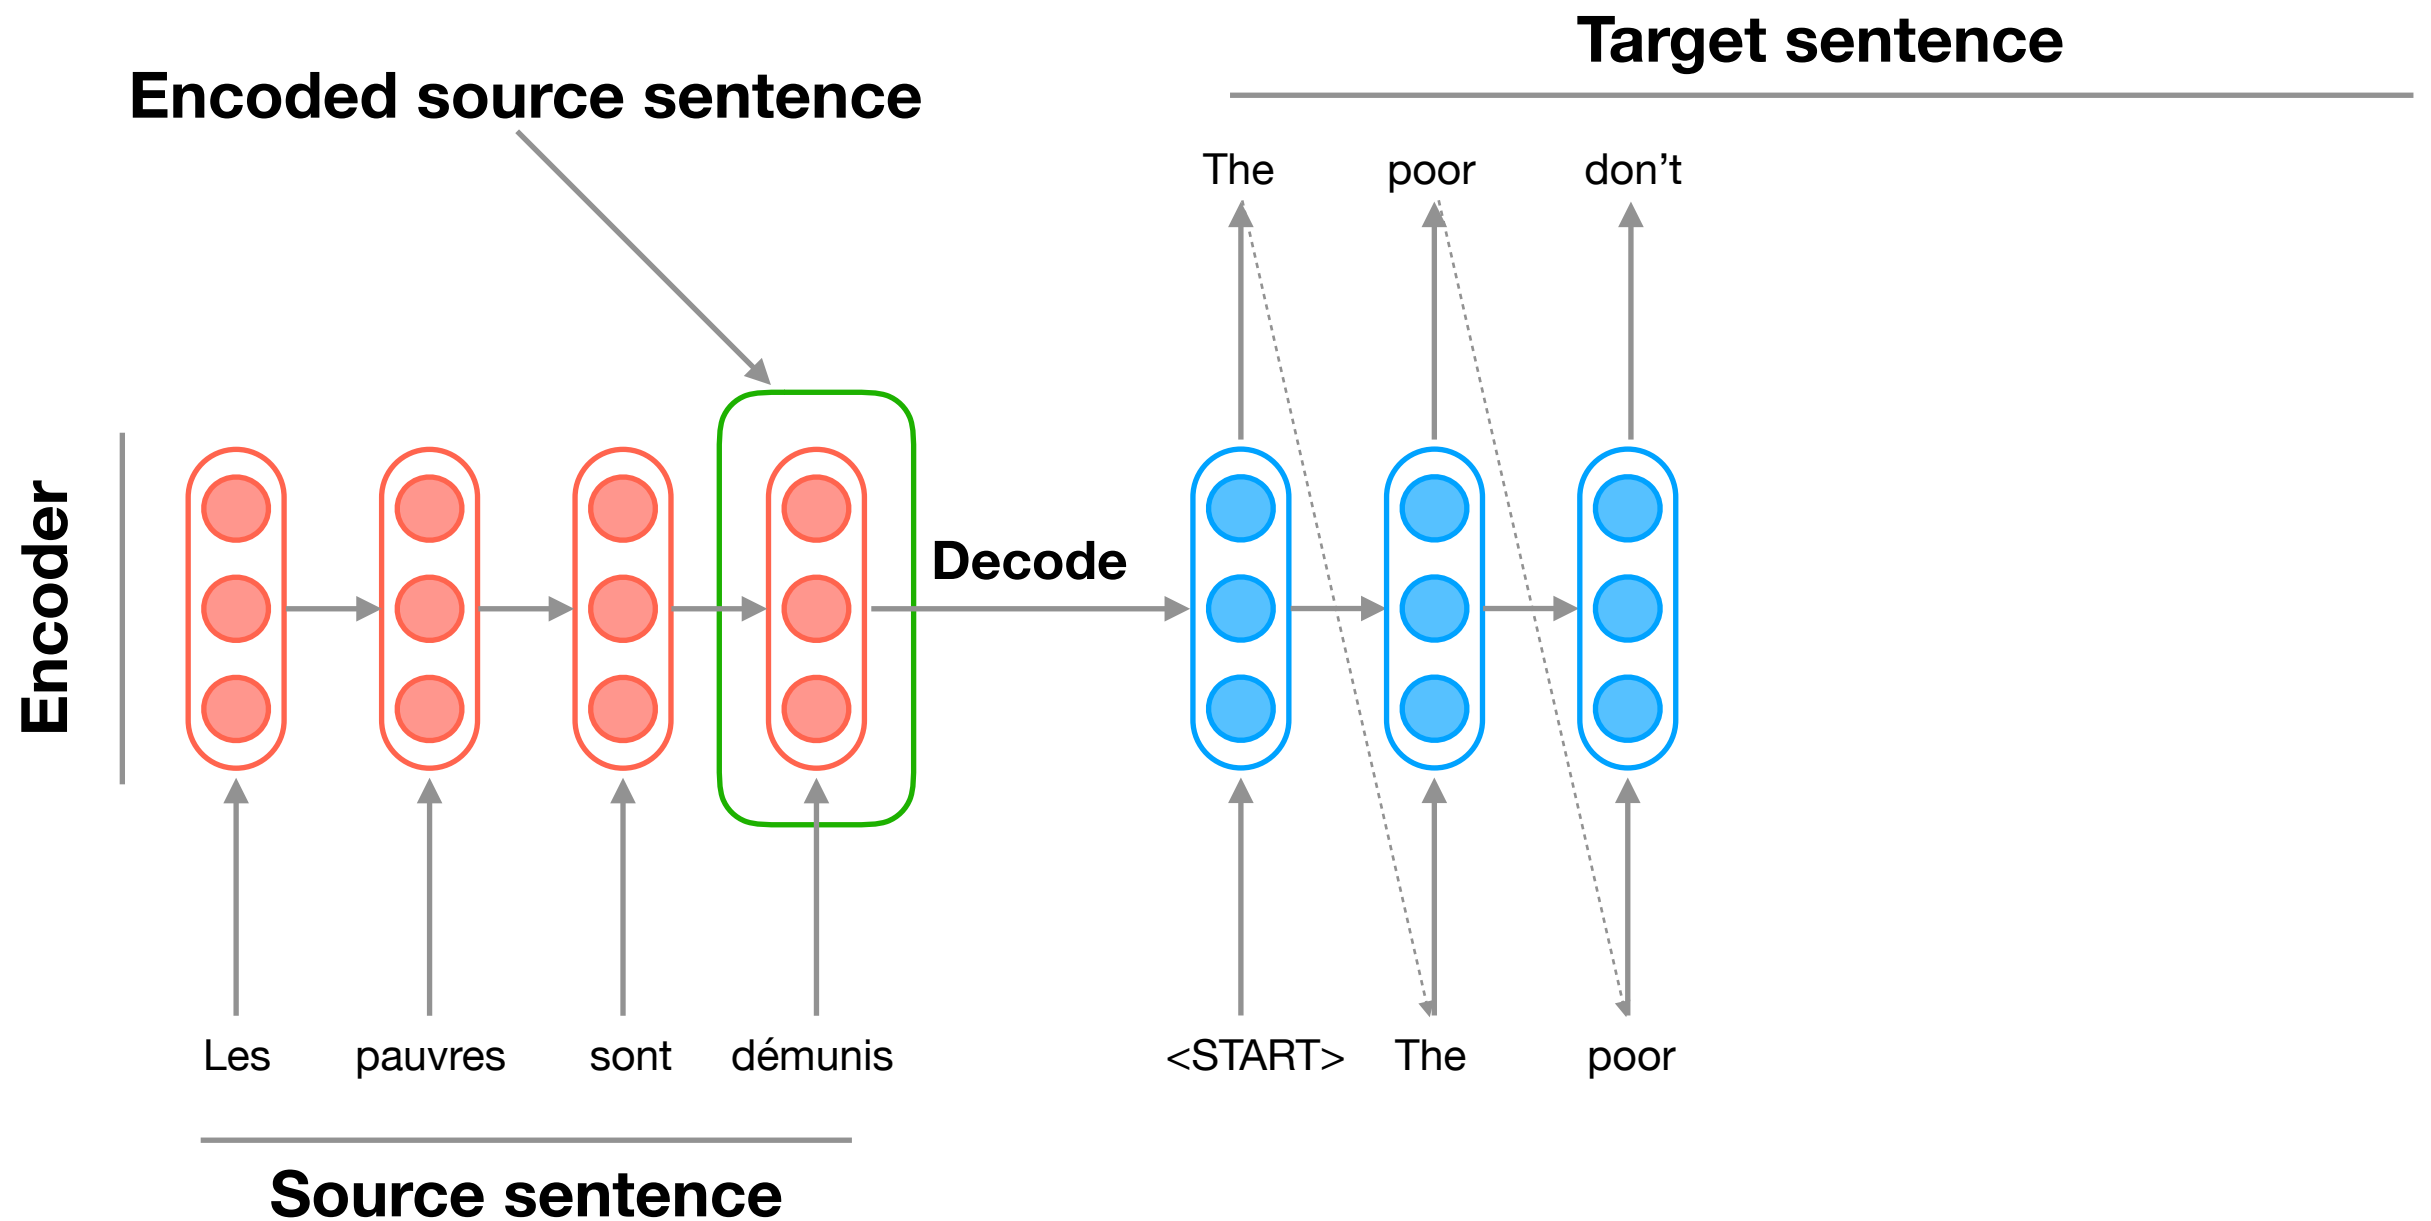
Sequence to sequence

Inference



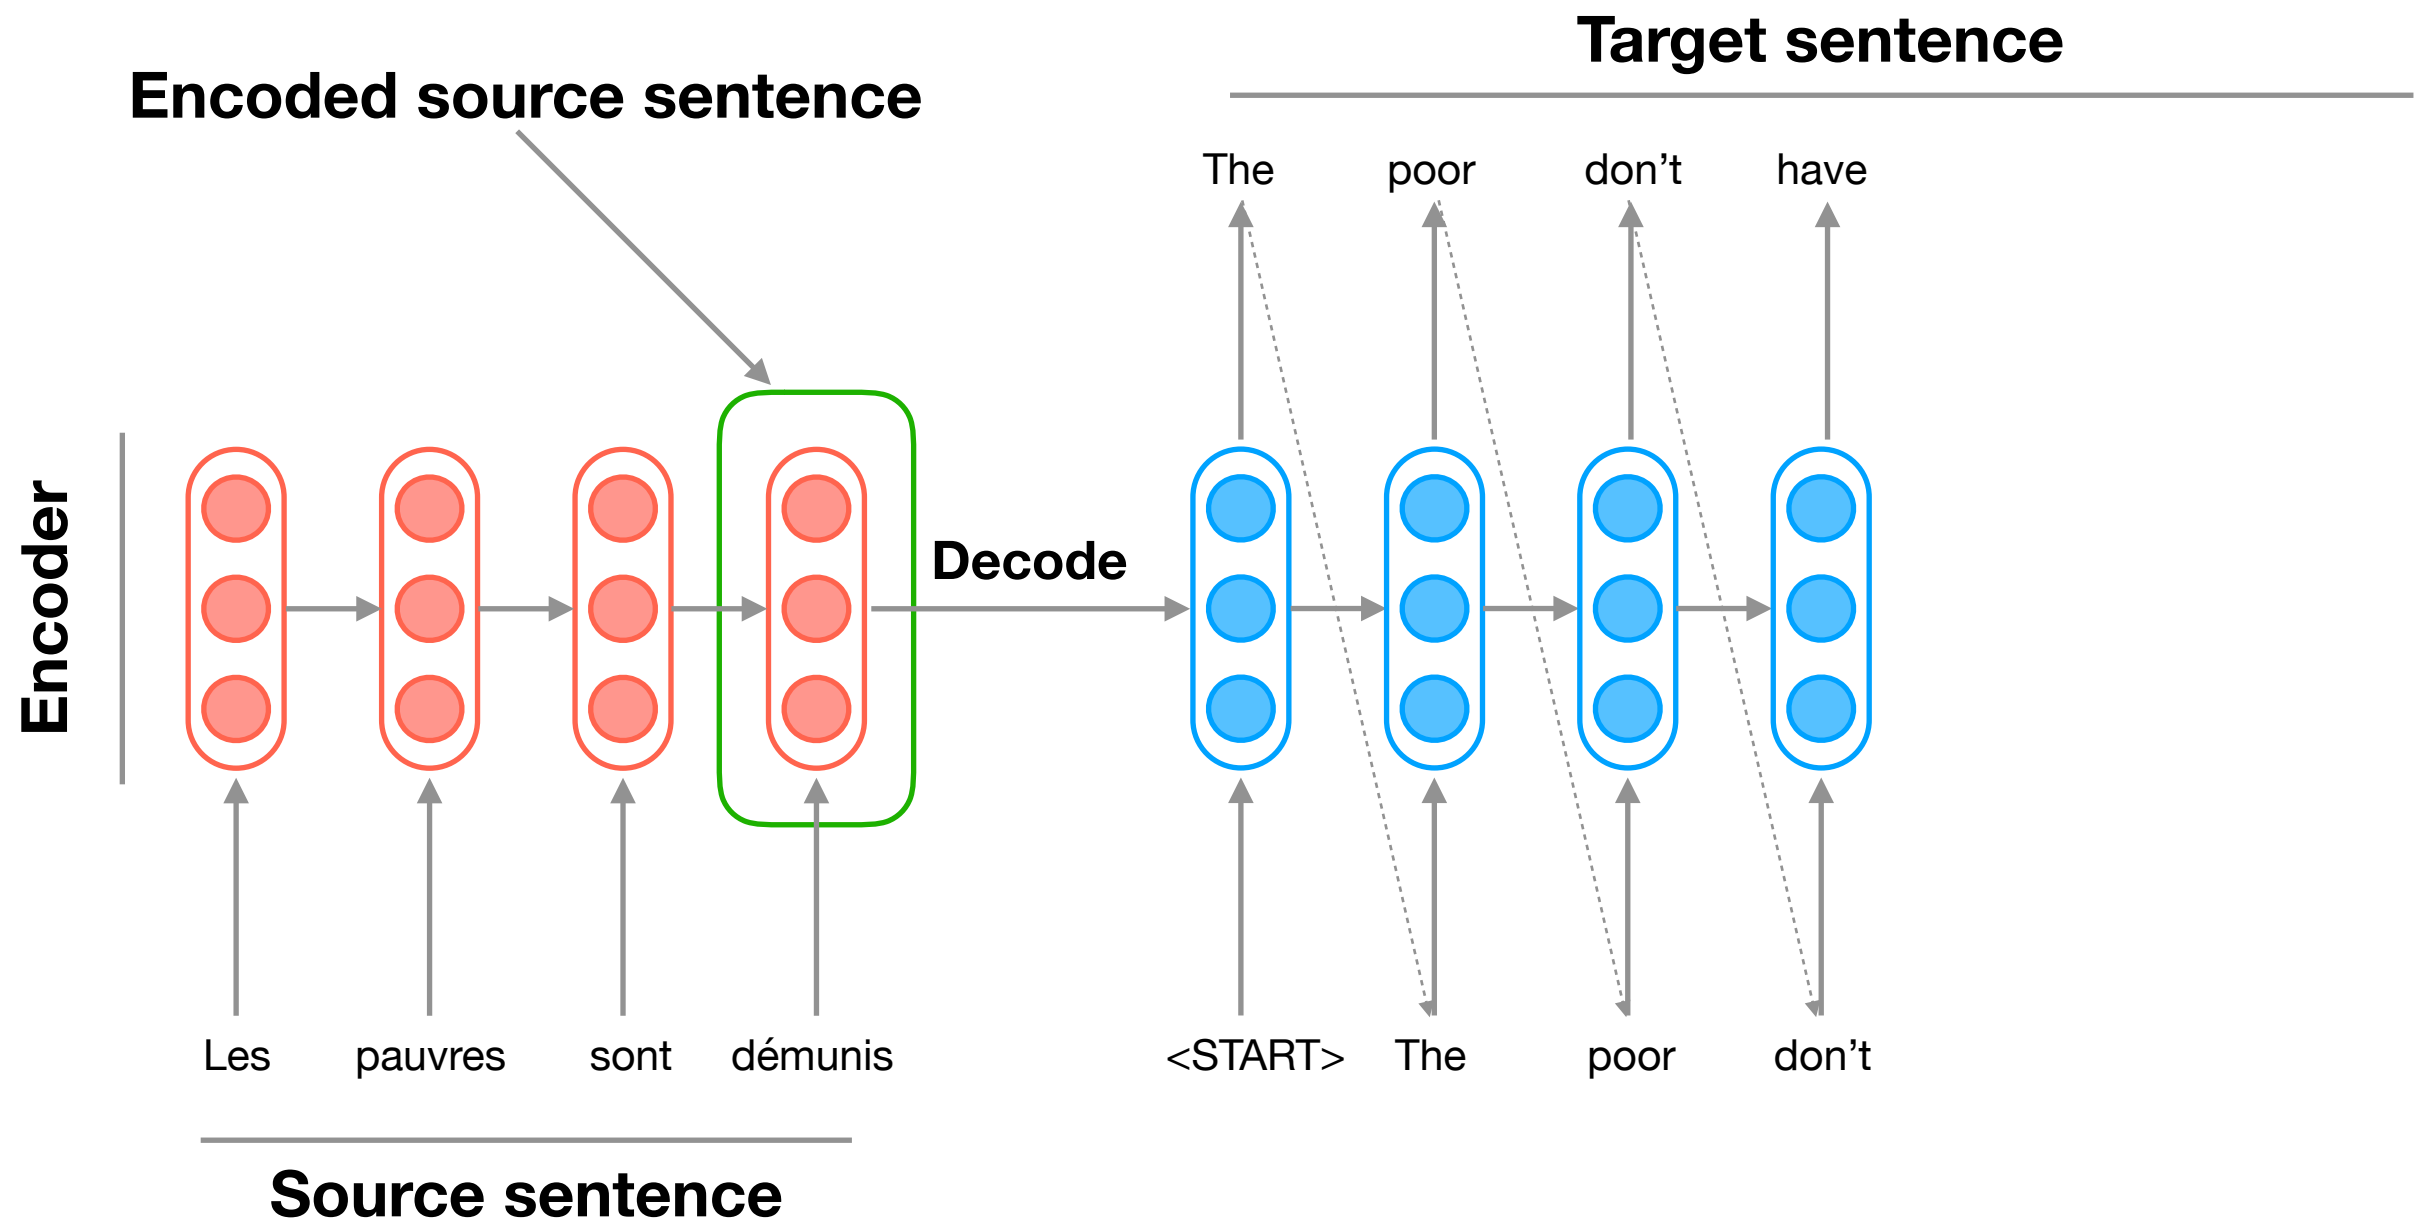
Sequence to sequence

Inference



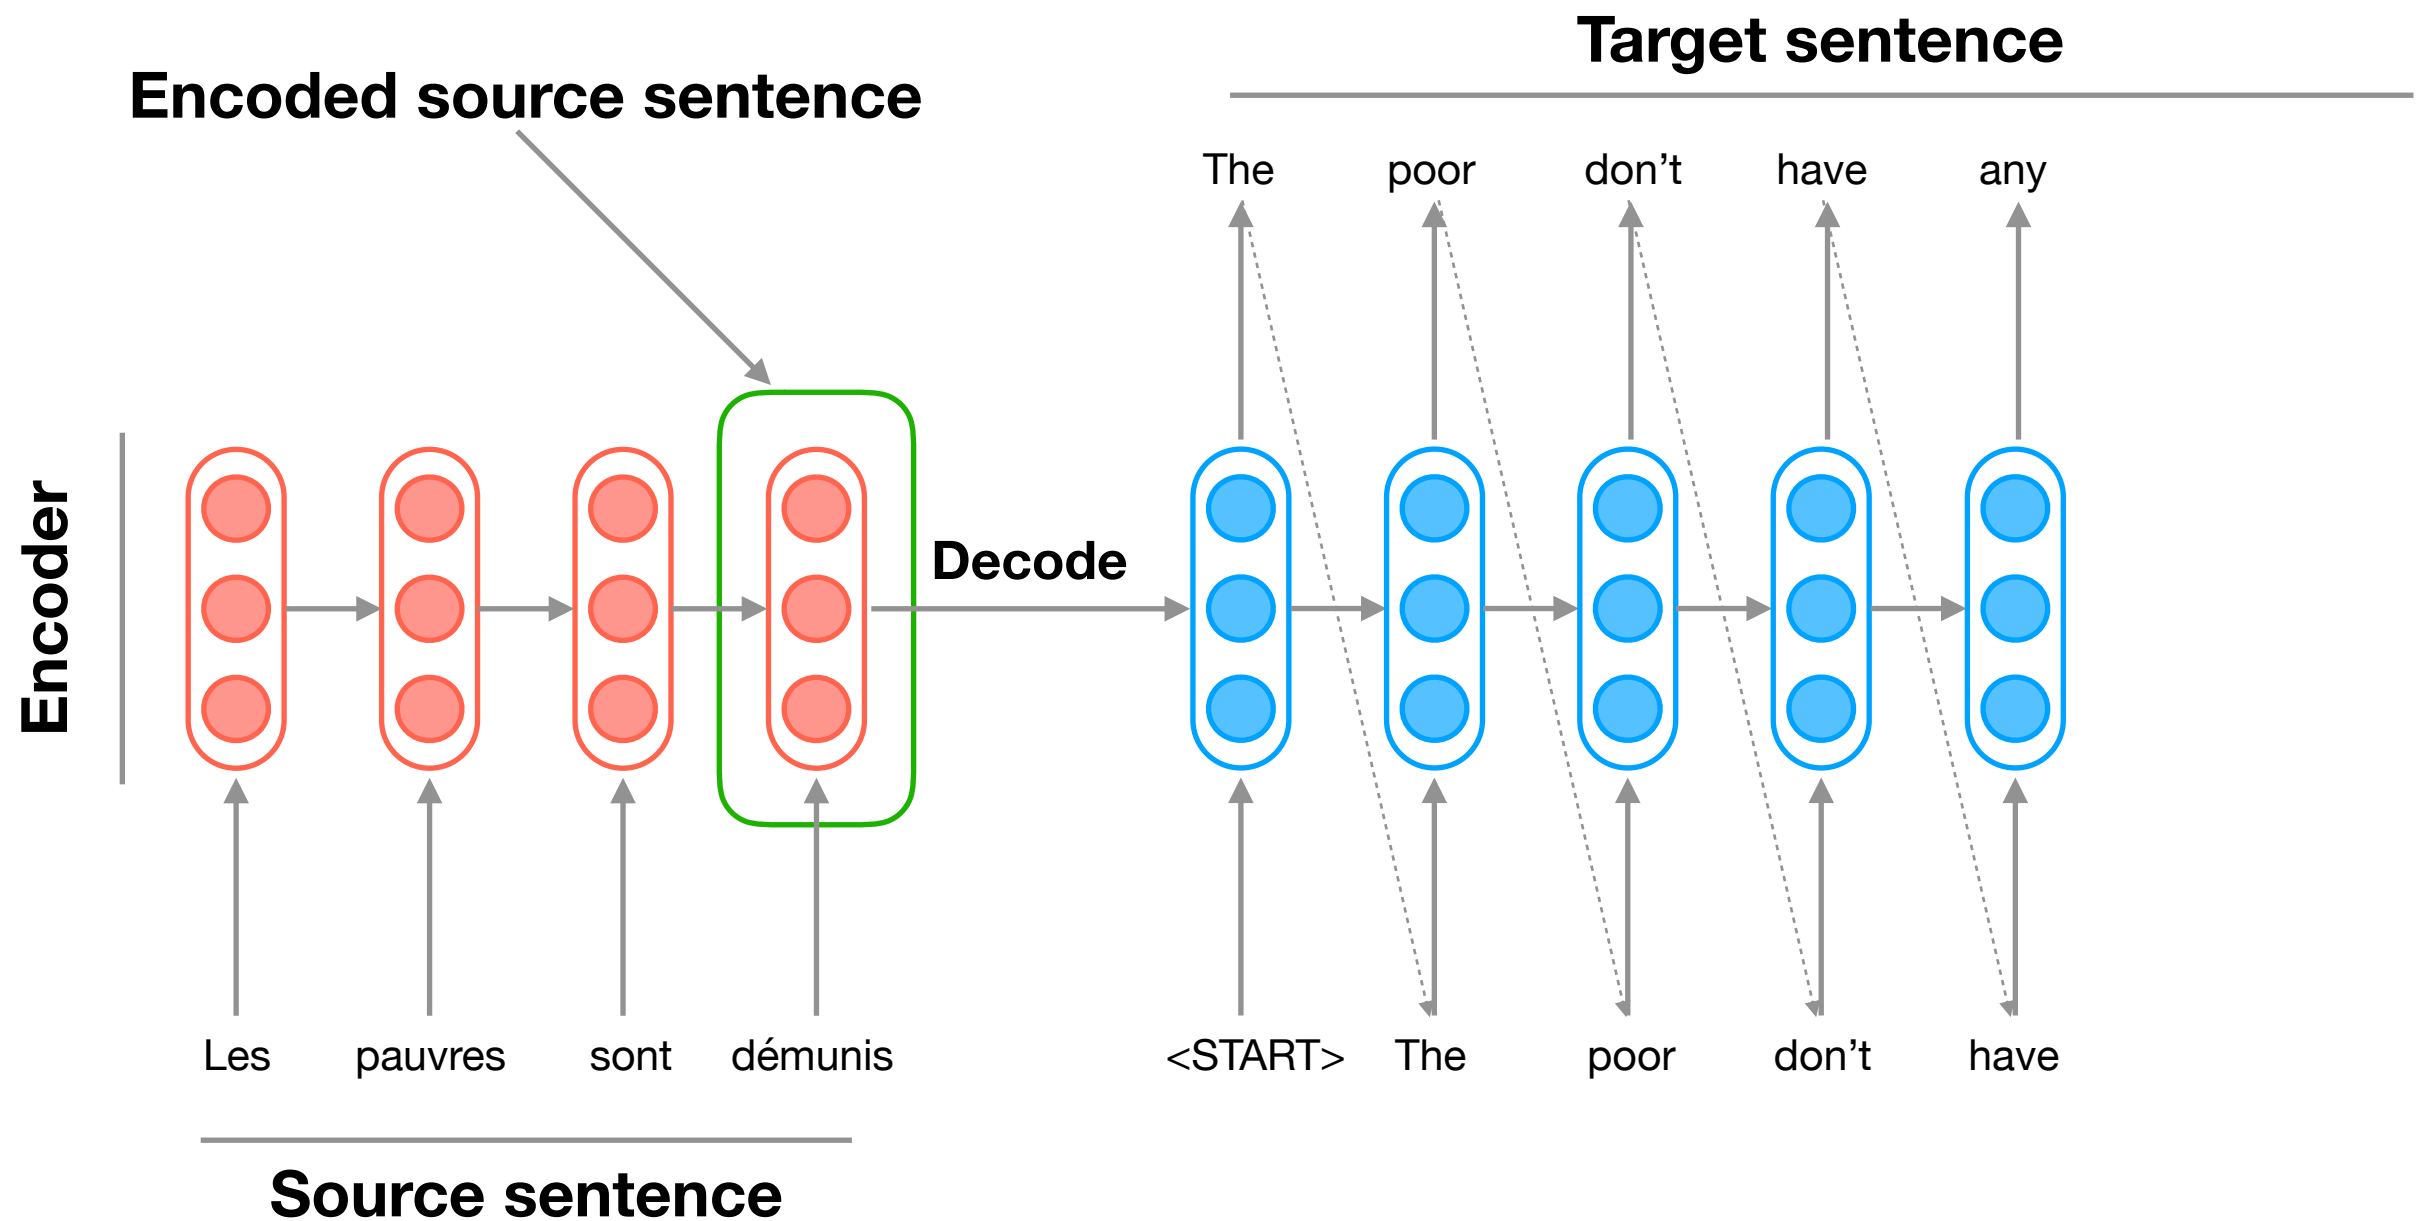
Sequence to sequence

Inference



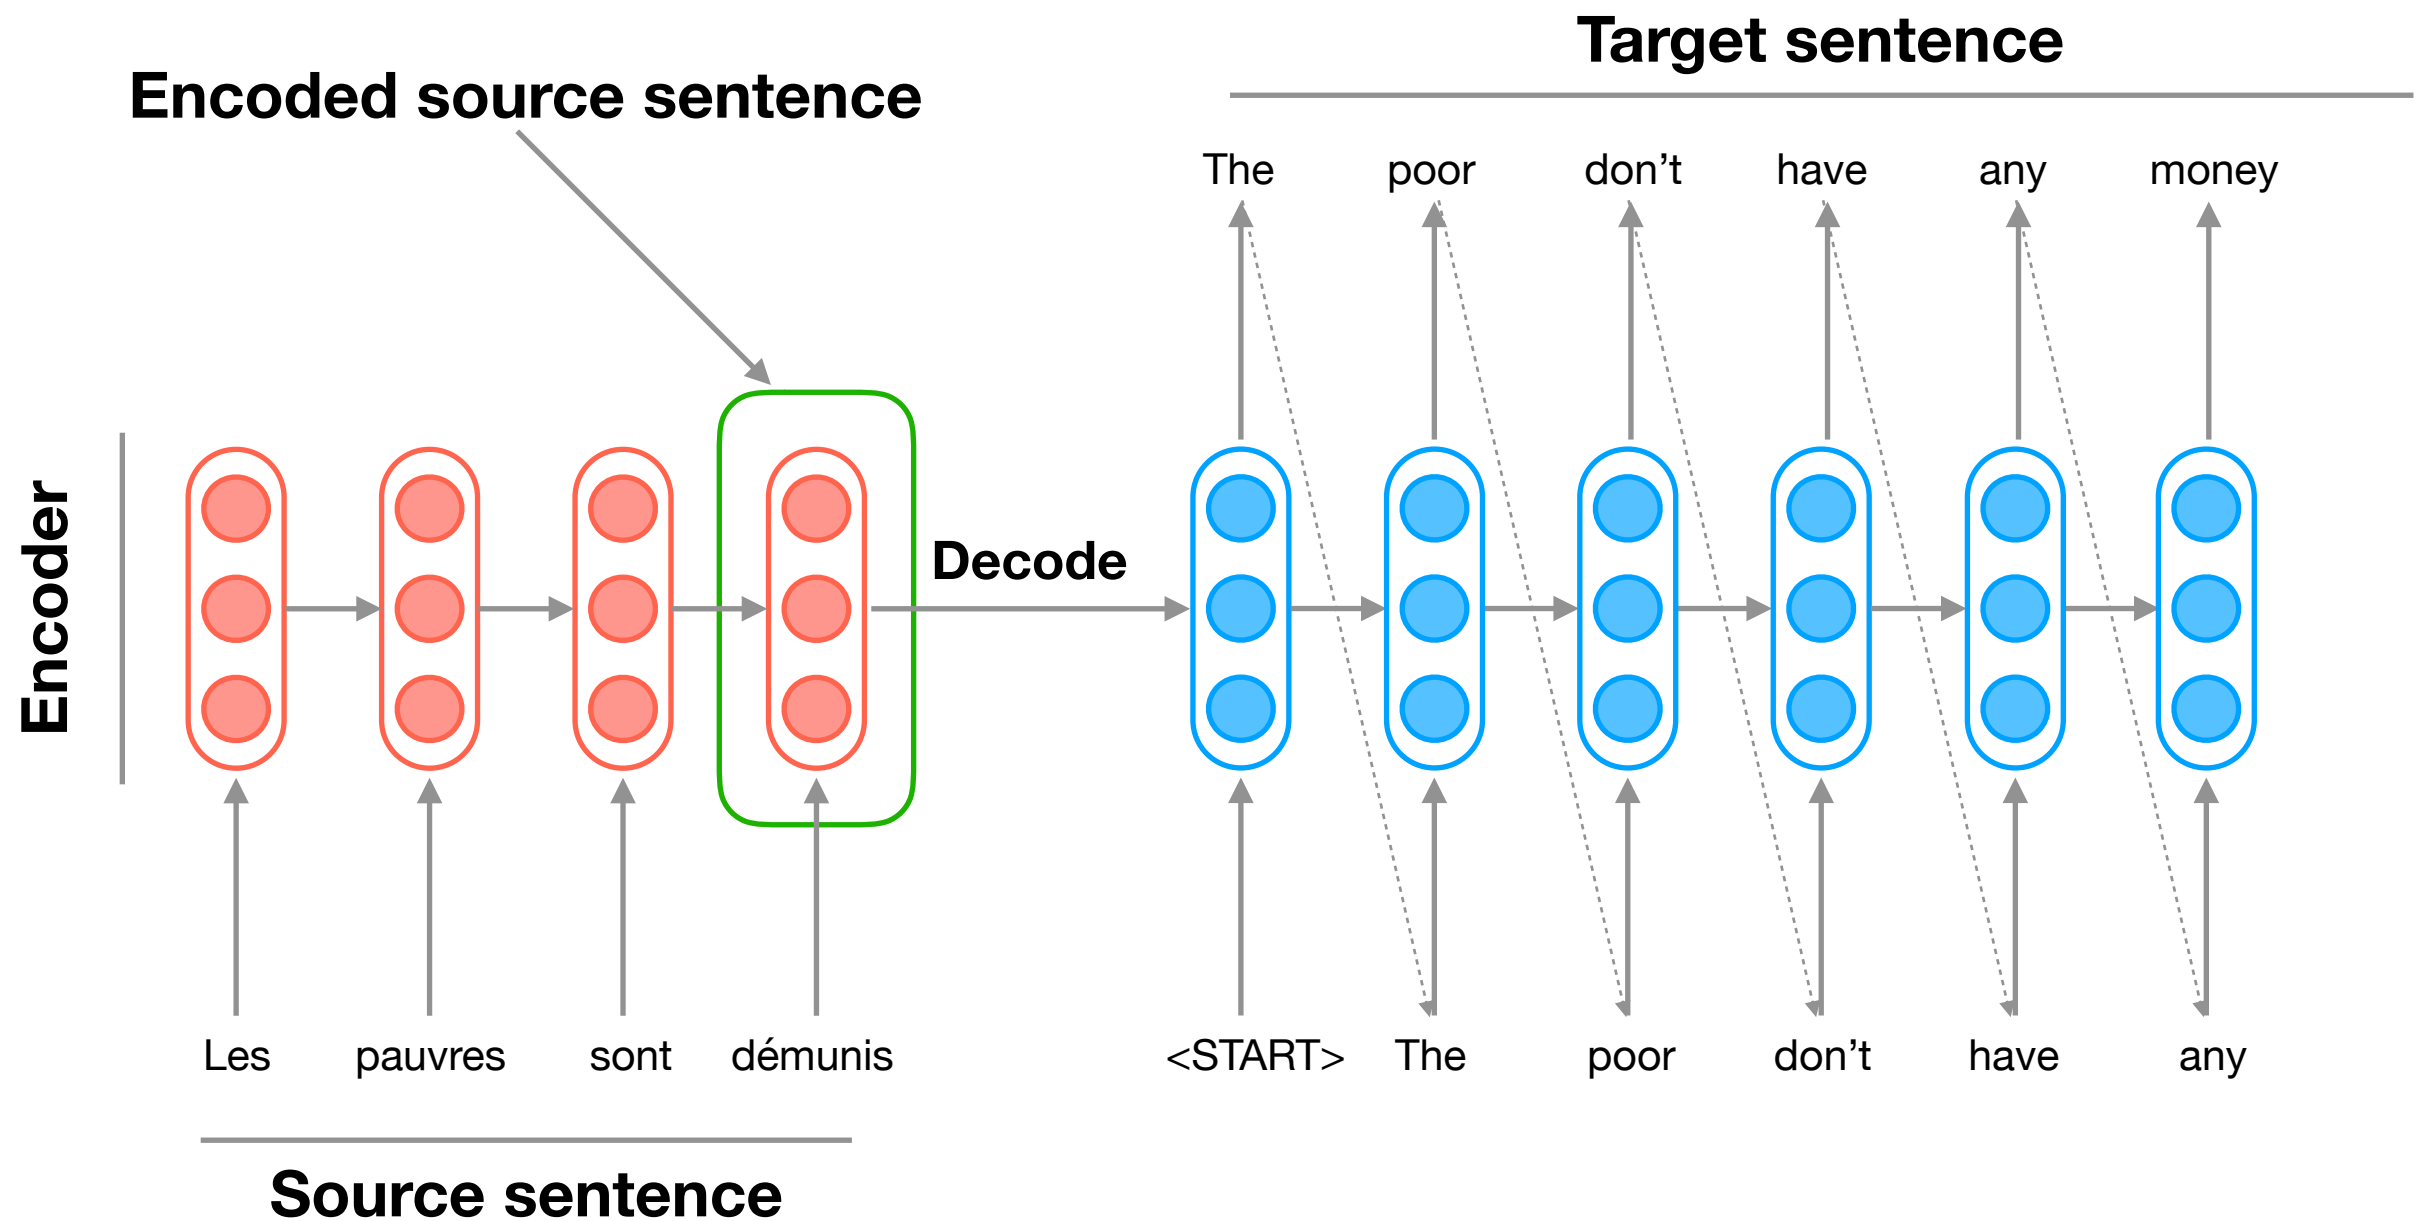
Sequence to sequence

Inference



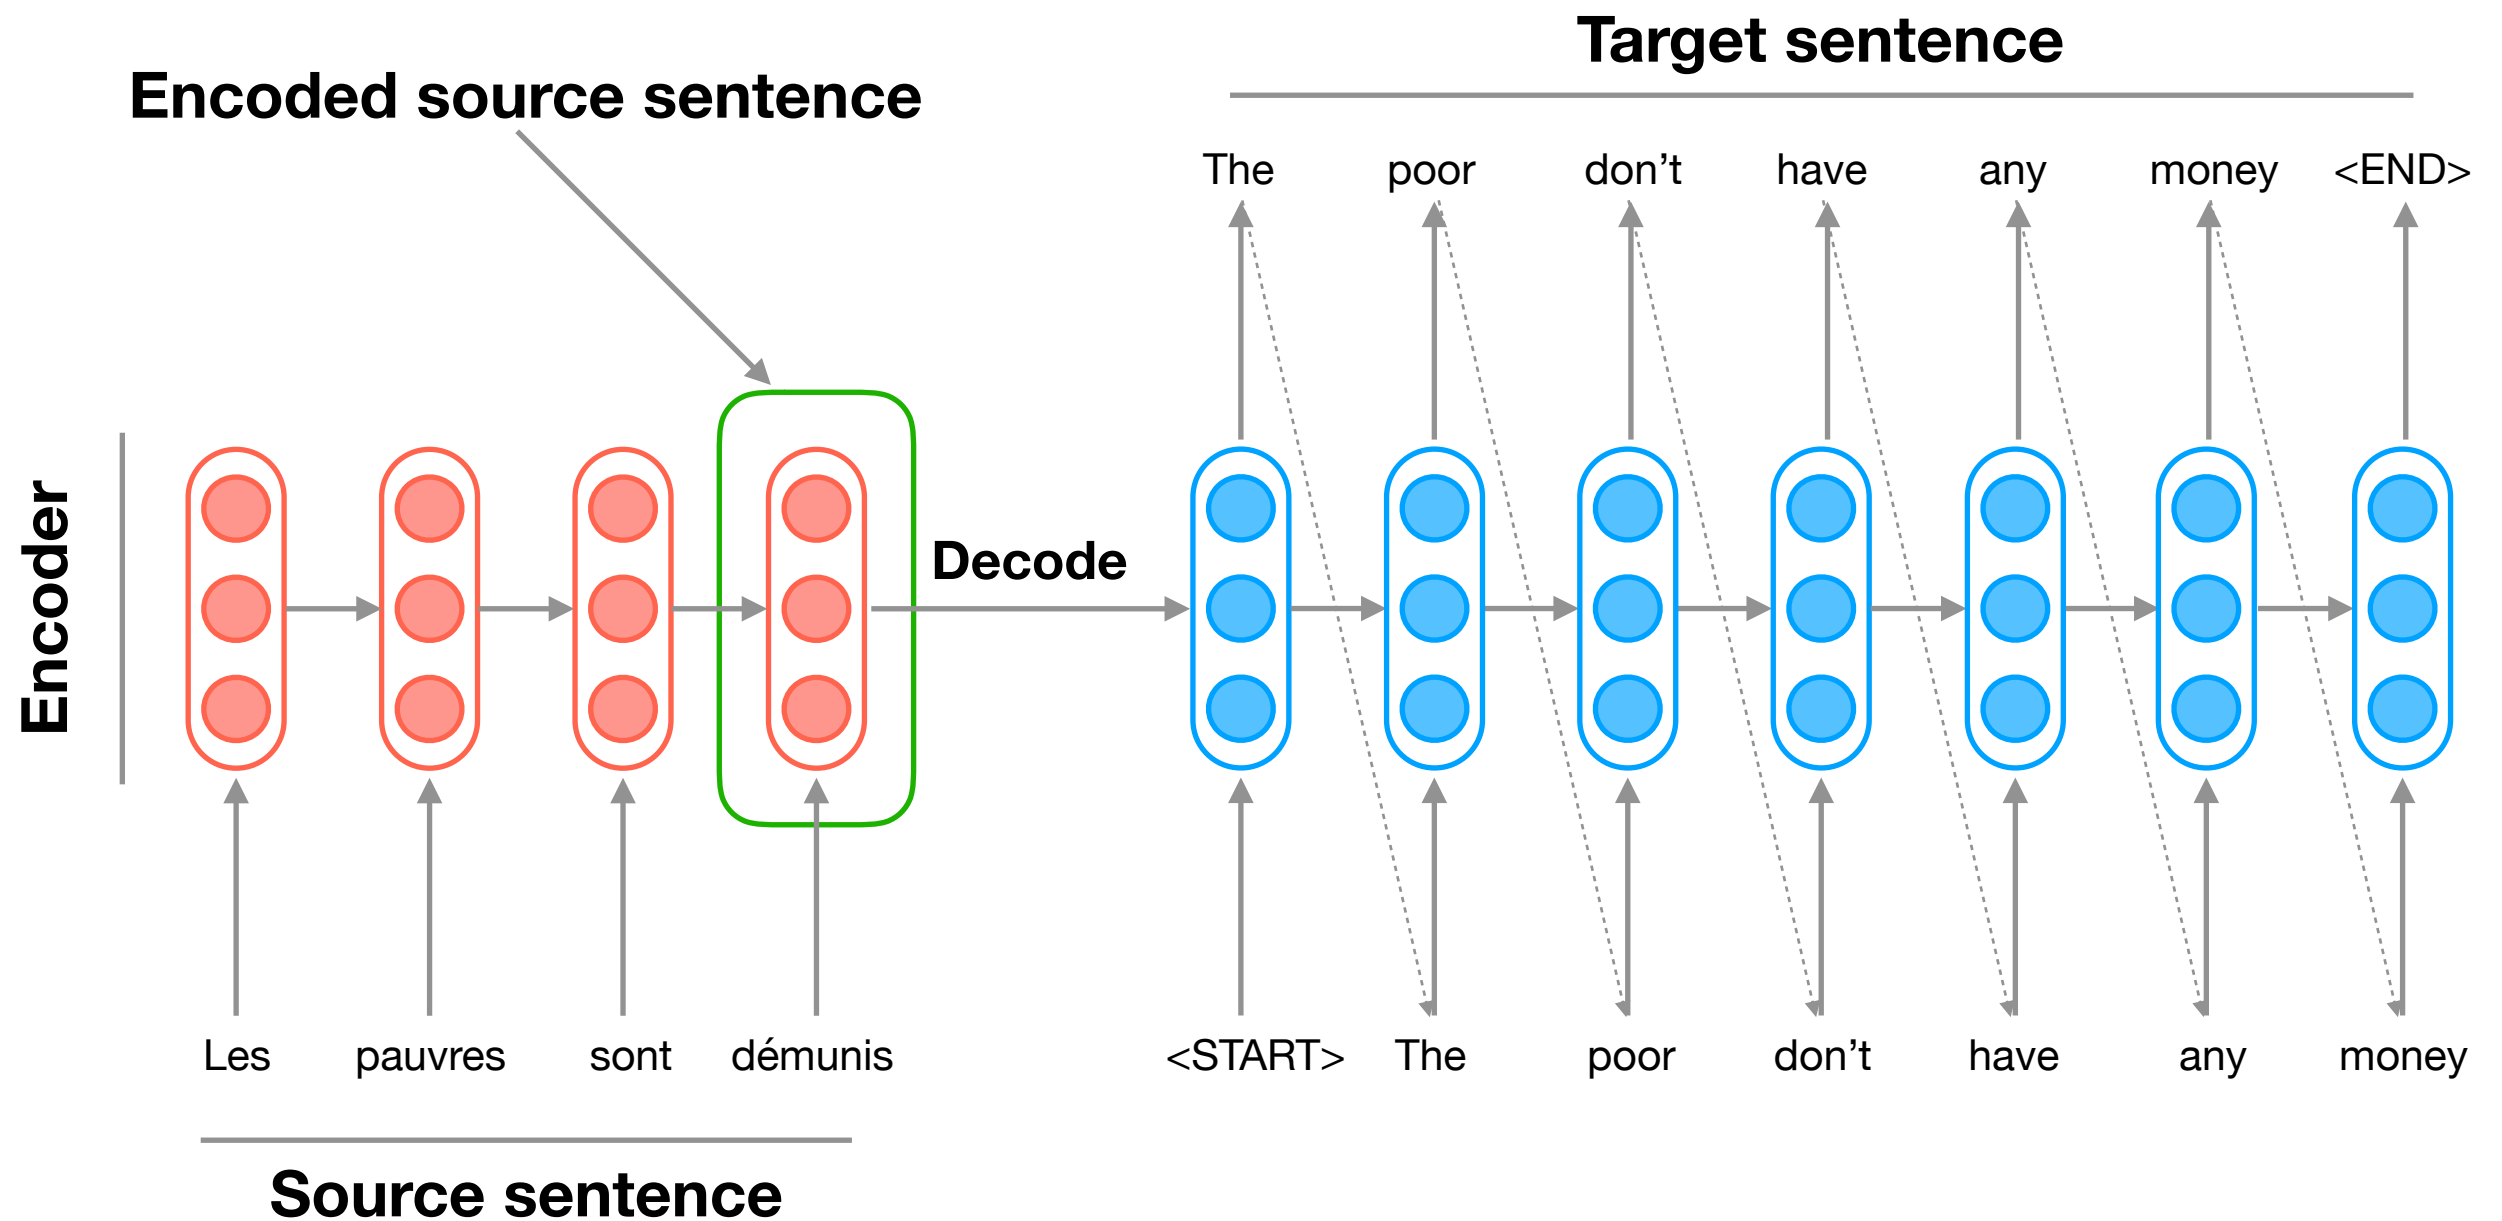
Sequence to sequence

Inference




Sequence to sequence

Inference



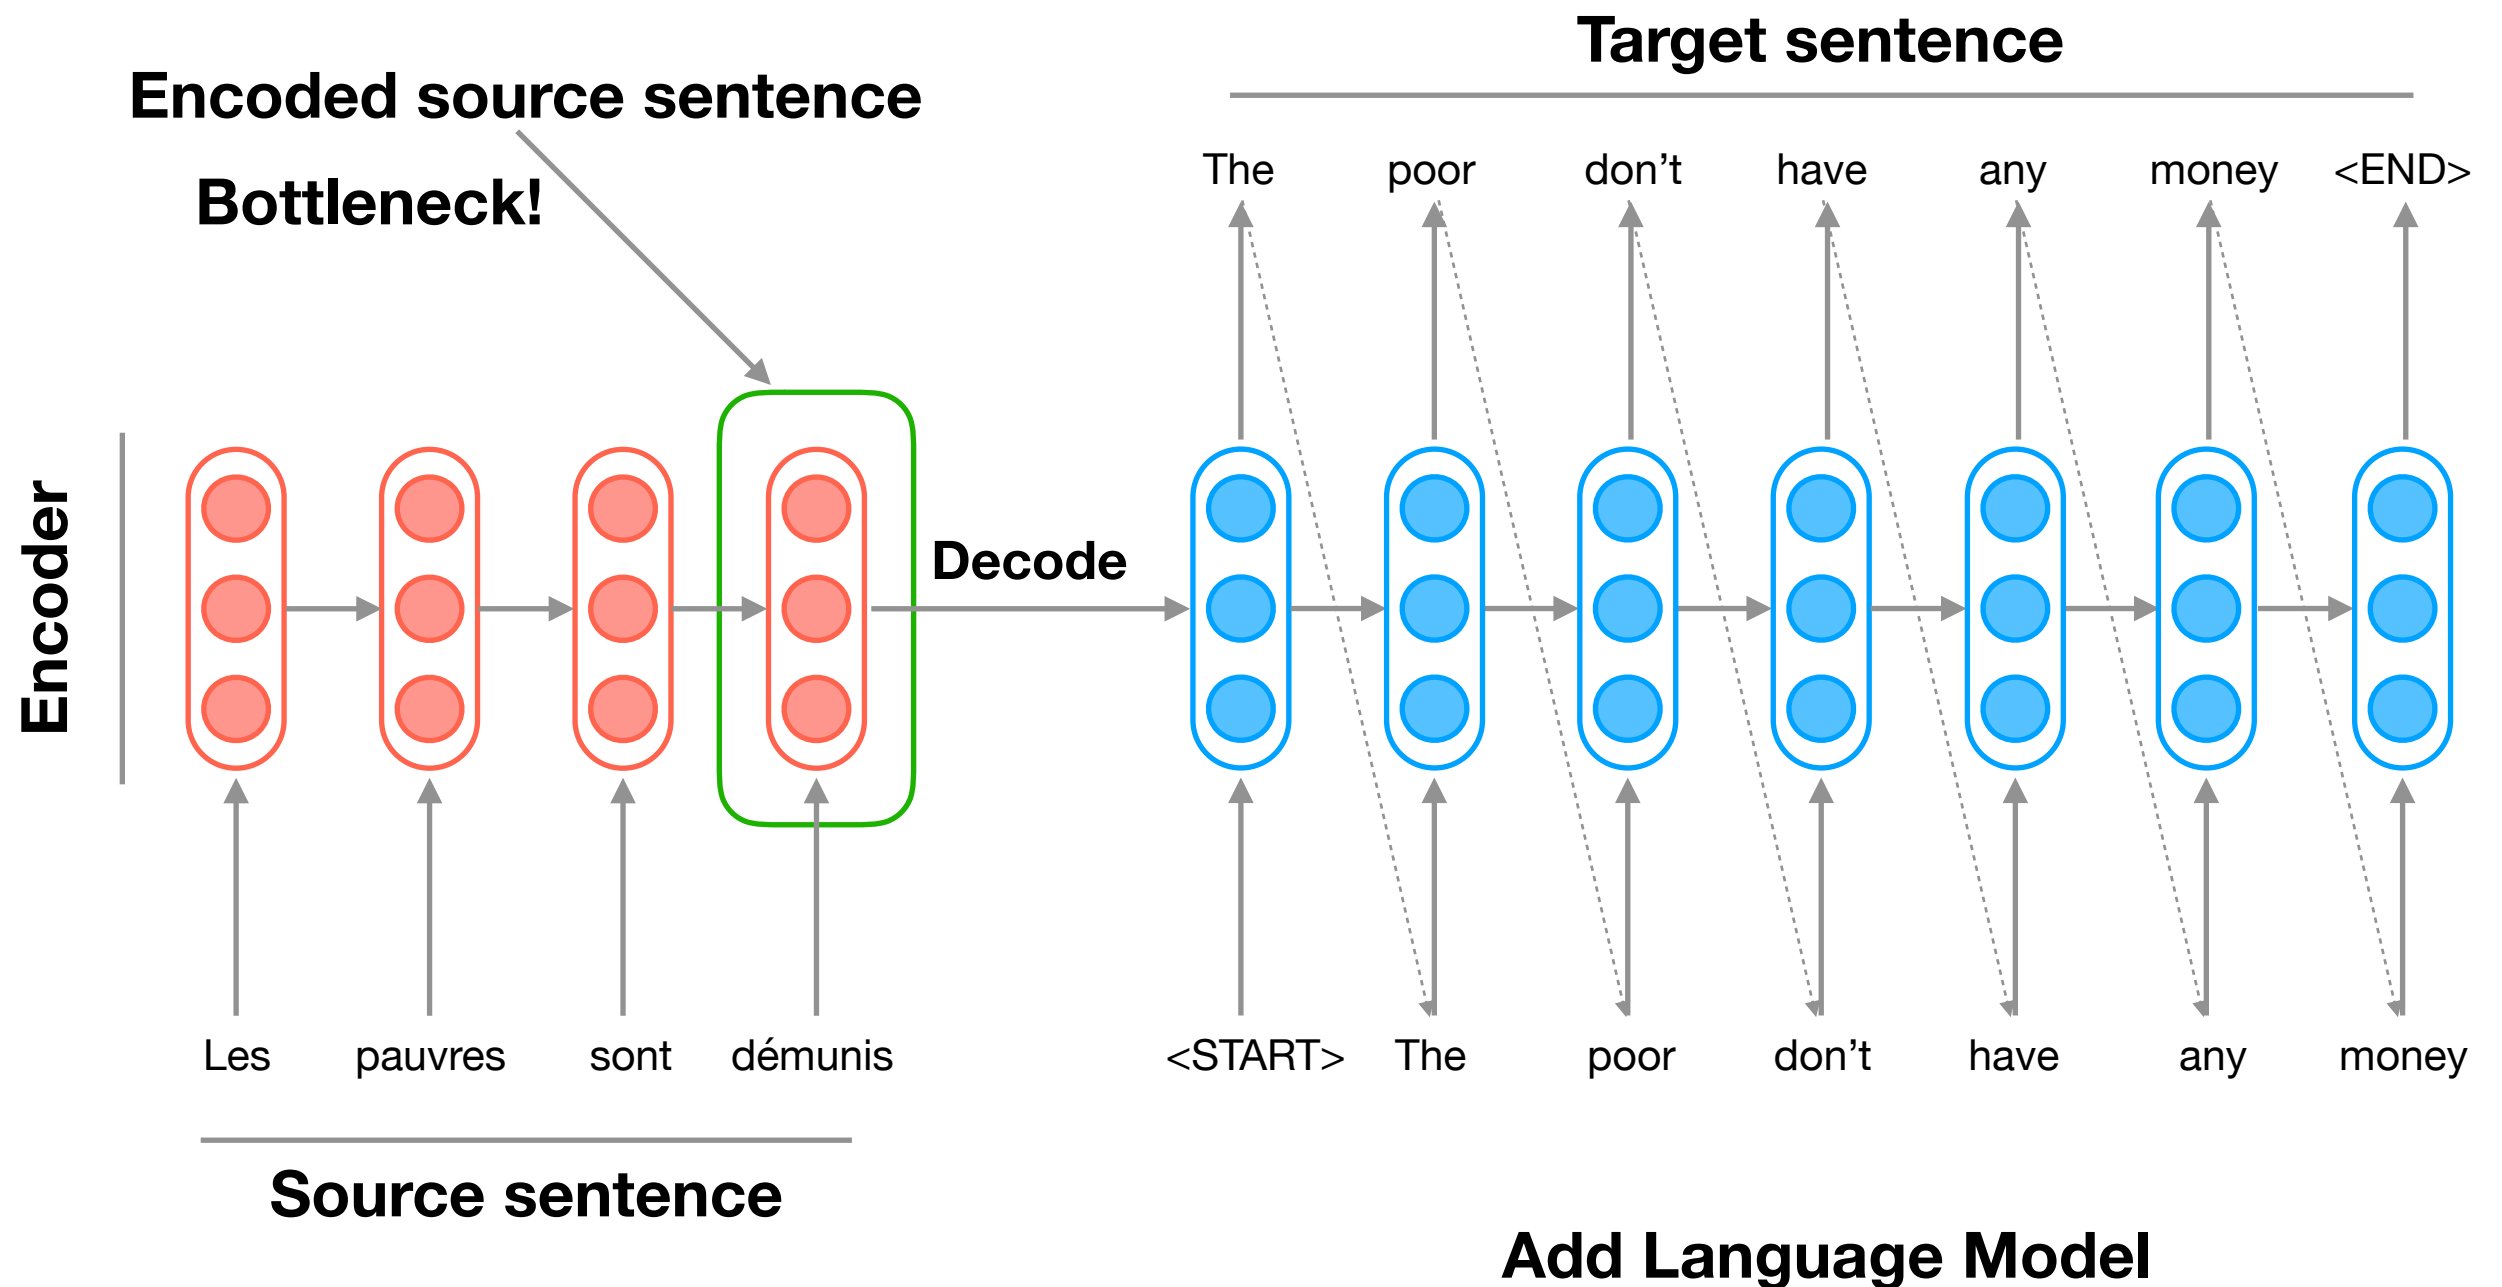
Sequence to sequence

Inference



Sequence to sequence

Inference



Contextualized Word Vectors

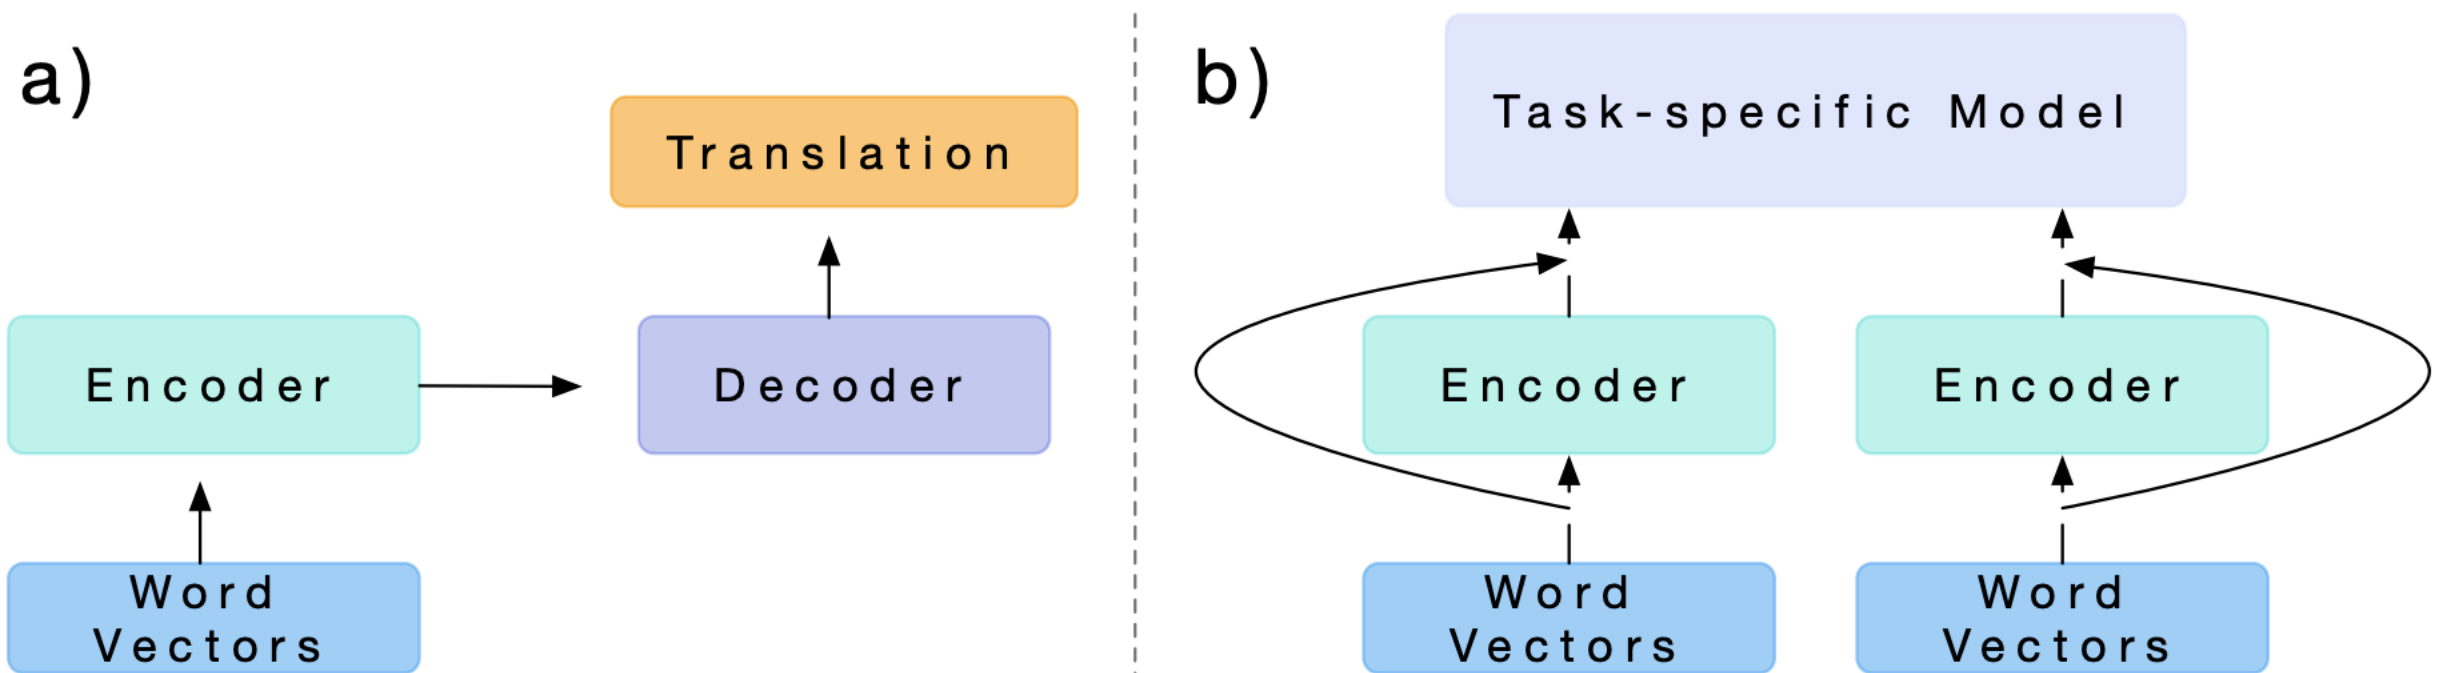


Figure 1: We a) train a two-layer, bidirectional LSTM as the encoder of an attentional sequence-to-sequence model for machine translation and b) use it to provide context for other NLP models.

Thanks for your Attention!

Boris Zubarev

 @bobazooba